

Poisson and Logistic Regression

Statistics 516, Homework 3 (Solutions)

Lip Cancer in Scotland

The data frame `epi.SClip` from the **epiR** package features data on the incidence of lip cancer in Scotland from 1975 to 1980. The following code processes the data for plotting and modeling, and shows the first few observations of a new data frame called `lipcancer`.¹

```
library(epiR)
data(epi.SClip) # necessary to "load" the data

library(dplyr)
lipcancer <- epi.SClip %>%
  mutate(district = factor(district, levels = rev(sort(unique(district))))) %>%
  mutate(percent = paste(prop.ag, "%", sep = "")) %>%
  mutate(percent = reorder(percent, prop.ag)) %>%
  select(district, cases, population, percent)
head(lipcancer)
```

	district	cases	population	percent
1	Caithness	11	83190	10%
2	Sutherland	5	37521	16%
3	Ross-Cromarty	15	129271	10%
4	Banff-Buchan	39	231337	16%
5	Nairn	3	29374	10%
6	Skye-Lochalsh	9	28324	16%

The data show for each of 56 districts the number of cases of lip cancer, the population (in person-years), and the percent of the population engaged in outdoor activity. Person-years is the sum of the number of years of exposure of all the people living in each district between 1975 to 1980.² The percent of people involved in outdoor activity (e.g., agriculture, fishing, forestry) is of interest because exposure to sunlight is a risk factor for lip cancer.³ The plot below shows the number of cases of lip cancer per person-year for each district, grouped by percent of the population engaged in outdoor activity. Note that the size of each point is proportional to the number of person-years for that district.


```
library(ggplot2)
p <- ggplot(lipcancer, aes(y = district, x = cases/population)) +
  theme_minimal() + geom_point(aes(size = population)) +
  facet_grid(percent ~ ., scales = "free_y", space = "free_y") +
```

¹Aside from some variable selection and renaming, this code sets the order of the levels of `district` so that they are alphabetical from top to bottom in the plot, and the values of `percent` so that they increase from top to bottom in the plot.

²Person-years is a common denominator in rates in epidemiology. It allows us to estimate the number of cases per person per year, controlling for both the number of people and how long they were observed. To compute person-years one needs to sum the years of observation of all the people observed. It also allows to account for changes in the size of the population. For example, suppose we observed a very small population for five years. Two of the people lived there all five years, but one person moved away after the third year. The total number of person-years would then be $5 + 5 + 3 = 13$.

³One must be very cautious about assuming that a relationship at the group level (e.g., district) implies a similar association at the individual level. This is known as the ecological fallacy. Districts with a larger percent of the population engaged in outdoor activities may differ in other ways from those with a smaller percentage that may, in part, be responsible for a higher incidence of lip cancer.

```
labs(y = NULL, x = "Cases per Person-Year", size = "Person-Years:") +
scale_x_continuous(labels = scales::label_number()) +
theme(axis.text.y = element_text(size = 7), legend.position = "top")
plot(p)
```



The objective here will be to model the relationship between the incidence rate of lip cancer and the percent of the population engaged in outdoor activity using Poisson regression.

1. Estimate a Poisson regression model for the rate of lip cancer, using the percent of the population engaged in outdoor activity as the only explanatory variable. Note that it will be treated here as a categorical explanatory variable, which will happen automatically since it is stored in the data frame as a character rather than a number. You will **not** be using `district` as an explanatory variable in your model.⁴ Be sure to include an offset variable to account for differences in the person-years across districts. Show the parameter estimates and their standard errors using the `summary` function so that I can verify that you estimated the model correctly.

Solution: We can estimate this model as follows.

⁴There may be considerable variation in the rate of lip cancer across districts, even among those with the same percent of the population engaged in outdoor activity. There are several ways that we might try to account for this. One would be to account for *over-dispersion* in the data due to variation between districts. Another approach would be to introduce a *random effect* for district. You might consider these approaches in future homework assignments.

```
m <- glm(cases ~ offset(log(population)) + percent,
  family = poisson, data = lipcancer)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-11.1471	0.1443	-77.229	0.000e+00
percent1%	0.2239	0.1805	1.240	2.149e-01
percent7%	1.2909	0.1696	7.610	2.732e-14
percent10%	1.2944	0.1708	7.579	3.481e-14
percent16%	1.5005	0.1691	8.875	6.983e-19
percent24%	2.0233	0.2378	8.508	1.764e-17

2. Estimate the expected number of cases of lip cancer per 100K (i.e., 100,000) person-years for each value of the percent of the explanatory variable. But be sure that you set the value of the offset variable to account for the fact that you are estimating the rate per 100K person-years and not per person-year.

Solution: I will show a couple of ways to do this.

```
trtools::contrast(m, tf = exp,
  a = list(percent = sort(unique(lipcancer$percent)), population = 100000),
  cnames = sort(unique(lipcancer$percent)))
```

	estimate	lower	upper
0%	1.442	1.086	1.913
1%	1.803	1.458	2.231
7%	5.242	4.402	6.242
10%	5.260	4.398	6.291
16%	6.465	5.440	7.682
24%	10.903	7.528	15.791

```
library(emmeans)
emmeans(m, ~percent, type = "response", offset = log(100000))
```

percent	rate	SE	df	asympt.LCL	asympt.UCL
0%	1.44	0.208	Inf	1.09	1.91
1%	1.80	0.196	Inf	1.46	2.23
7%	5.24	0.467	Inf	4.40	6.24
10%	5.26	0.480	Inf	4.40	6.29
16%	6.46	0.569	Inf	5.44	7.68
24%	10.90	2.061	Inf	7.53	15.79


Confidence level used: 0.95

Intervals are back-transformed from the log scale

Now suppose we want to show the estimated rate of lip cancer on the plot. This needs to be per person-year, since that is the scale used for the plot. Here is one way to do that.

```
d <- data.frame(percent = sort(unique(lipcancer$percent)),
  population = 1)
d$yhat <- predict(m, newdata = d, type = "response")
p <- ggplot(lipcancer, aes(y = district, x = cases/population)) +
  theme_minimal() + geom_point(aes(size = population)) +
  facet_grid(percent ~ ., scales = "free_y", space = "free_y") +
  labs(y = NULL, x = "Cases per Person-Year", size = "Person-Years:") +
  scale_x_continuous(labels = scales::label_number()) +
  theme(axis.text.y = element_text(size = 7), legend.position = "top") +
  geom_vline(aes(xintercept = yhat), data = d)
```

```
plot(p)
```



- Estimate five rate ratios to compare the rate of lip cancer at 1%, 7%, 10%, 16%, and 24% versus 0% of the population involved in outdoor activity. Write a sentence or two to interpret each estimated rate ratio in terms of the relationship between the percent of the population involved in outdoor activity and the rate of lip cancer.

Solution: I will show a couple of ways to estimate these rate ratios.

```
trtools::contrast(m, tf = exp,
  a = list(percent = c("1%", "7%", "10%", "16%", "24%"), population = 1),
  b = list(percent = "0%", population = 1),
  cnames = paste(c("1%", "7%", "10%", "16%", "24%"), "vs 0%"))
```

	estimate	lower	upper
1% vs 0%	1.251	0.8781	1.782
7% vs 0%	3.636	2.6076	5.070
10% vs 0%	3.649	2.6108	5.099
16% vs 0%	4.484	3.2193	6.246
24% vs 0%	7.563	4.7455	12.053

```
emmeans::contrast(emmeans(m, ~percent, offset = log(1), type = "response"),
  method = "trt.vs.ctrl", ref = 1, infer = TRUE, adjust = "none")
```

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
----------	-------	----	----	------------	------------	------	---------	---------

1% / 0%	1.25	0.226	Inf	0.878	1.78	1	1.240	0.2149
7% / 0%	3.64	0.617	Inf	2.608	5.07	1	7.610	<.0001
10% / 0%	3.65	0.623	Inf	2.611	5.10	1	7.579	<.0001
16% / 0%	4.48	0.758	Inf	3.219	6.25	1	8.875	<.0001
24% / 0%	7.56	1.798	Inf	4.745	12.05	1	8.508	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

Note that the value of the offset does not matter here since we are looking at ratios. We can interpret these as follows.

The rate of lip cancer is in a district with a percent outdoor activity of 1% is about 1.25 times that of a district with an outdoor activity of 0% (i.e., about 25% higher).

The rate of lip cancer is in a district with a percent outdoor activity of 7% is about 3.64 times that of a district with an outdoor activity of 0% (i.e., about 264% higher).

The rate of lip cancer is in a district with a percent outdoor activity of 10% is about 3.65 times that of a district with an outdoor activity of 0% (i.e., about 265% higher).

The rate of lip cancer is in a district with a percent outdoor activity of 16% is about 4.48 times that of a district with an outdoor activity of 0% (i.e., about 348% higher).

The rate of lip cancer is in a district with a percent outdoor activity of 24% is about 7.59 times that of a district with an outdoor activity of 0% (i.e., about 659% higher).

Alternatively we can “flip” the rate ratios in the following way.

```
trtools::contrast(m, tf = exp,
  a = list(percent = "0%", population = 1),
  b = list(percent = c("1%", "7%", "10%", "16%", "24%"), population = 1),
  cnames = paste("0% vs", c("1%", "7%", "10%", "16%", "24%")))
```

	estimate	lower	upper
0% vs 1%	0.7994	0.56115	1.1388
0% vs 7%	0.2750	0.19725	0.3835
0% vs 10%	0.2741	0.19611	0.3830
0% vs 16%	0.2230	0.16011	0.3106
0% vs 24%	0.1322	0.08296	0.2107

```
emmeans::contrast(emmeans(m, ~percent, offset = log(1), type = "response"),
  method = "trt.vs.ctrl1", ref = 1, infer = TRUE, adjust = "none", reverse = TRUE)
```

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
0% / 1%	0.799	0.1443	Inf	0.561	1.139	1	-1.240	0.2149
0% / 7%	0.275	0.0466	Inf	0.197	0.384	1	-7.610	<.0001
0% / 10%	0.274	0.0468	Inf	0.196	0.383	1	-7.579	<.0001
0% / 16%	0.223	0.0377	Inf	0.160	0.311	1	-8.875	<.0001
0% / 24%	0.132	0.0314	Inf	0.083	0.211	1	-8.508	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

We can interpret these rate ratios as follows.

The rate of lip cancer in a district with a percent of outdoor activity of 0% is about 0.8 times that of a district with a percent of outdoor activity of 1% (i.e., about 20% lower).

The rate of lip cancer in a district with a percent of outdoor activity of 0% is about 0.275 times that of a district with a percent of outdoor activity of 7% (i.e., about 72.5% lower).

The rate of lip cancer in a district with a percent of outdoor activity of 0% is about 0.274 times that of a district with a percent of outdoor activity of 10% (i.e., about 72.6% lower).

The rate of lip cancer in a district with a percent of outdoor activity of 0% is about 0.22 times that of a district with a percent of outdoor activity of 16% (i.e., about 78% lower).

The rate of lip cancer in a district with a percent of outdoor activity of 0% is about 0.13 times that of a district with a percent of outdoor activity of 24% (i.e., about 87% lower).

Treatment of Familial Adenomatous Polyposis

Familial adenomatous polyposis (FAP) is a rare genetic condition that causes polyps to form in the large intestine and rectum. These polyps are likely to later become cancerous. Data from a study that investigated the effectiveness of a non-steroidal anti-inflammatory drug for treating FAP are in the data frame `polyps` in the package `HSAUR3`.⁵

```
library(HSAUR3)
polyps
```

	number	treat	age
1	63	placebo	20
2	2	drug	16
3	28	placebo	18
4	17	drug	22
5	61	placebo	13
6	1	drug	23
7	7	placebo	34
8	15	placebo	50
9	44	placebo	19
10	25	drug	17
11	3	drug	23
12	28	placebo	22
13	10	placebo	30
14	40	placebo	27
15	33	drug	23
16	46	placebo	22
17	50	placebo	34
18	3	drug	23
19	1	drug	22
20	4	drug	42

The plots below show the raw data.

```
library(ggplot2)
library(cowplot)


p1 <- ggplot(polyps, aes(x = treat, y = number, fill = treat)) + theme_minimal() +
  geom_dotplot(binaxis = "y", stackdir = "center", show.legend = FALSE, binwidth = 1) +
  labs(x = "Treatment", y = "Number of Polyps") + ylim(0, 63)

p2 <- ggplot(polyps, aes(x = age, y = number, color = treat)) + theme_minimal() +
```

⁵Giardiello, F. M., Hamilton, S. R., Krush, A. J., Piantadosi, S., Hyland, L. M., Celano, P., Booker, S. V., Robinson, C. R., & Offerhaus, G. J. (1993). Treatment of colonic and rectal adenomas with sulindac in familial adenomatous polyposis. *New England Journal of Medicine*, 328(18), 1313–1316.

```
geom_point() + labs(x = "Age (years)", y = NULL, color = "Treatment") +
  theme(legend.position = c(0.8, 0.8)) + ylim(0, 63)

plot_grid(p1, p2, rel_widths = c(1, 3))
```



Twenty subjects with FAP were randomly assigned to either a treatment (drug) or control (placebo) group. After twelve months the polyps in each subject were counted. The focus here is on the statistical relationship between the number of polyps and the treatment condition. But subject age can be used as a covariate since the number of polyps may depend on age.⁶ Here you will use Poisson regression to make inferences about the treatment effect as well as the effect of age on the number of polyps for people with FAP.

1. Estimate a Poisson regression model with the number of polyps as the response variable and only treatment as the explanatory variable (do not include age yet). Report the parameter estimates and their standard errors using `summary` so that I can verify that you estimated the model correctly. Estimate the rate ratio for the effect of the treatment, and write a sentence or two to interpret this rate ratio in terms of the effect of the treatment on the number of polyps. Finally estimate the expected number of polyps for each treatment group.

Solution: We can estimate this model as follows.

```
m <- glm(number ~ treat, family = poisson, data = polyps)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	3.573	0.05051	70.75	0.000e+00
treatdrug	-1.282	0.11742	-10.92	9.446e-28

We can estimate the rate ratio as follows.

```
trtools::contrast(m, tf = exp,
  a = list(treat = "drug"), b = list(treat = "placebo"))
```

estimate	lower	upper
0.2775	0.2204	0.3493

⁶This model can be thought of a Poisson regression version of an analysis of covariance which uses a covariate to “control” for some variation in the response variable to improve inferences for the treatment variable in a randomized experiment.

Note that for this model we can also estimate it by just exponentiating the parameters.

```
exp(cbind(coef(m), confint(m)))
```

```

              2.5 %   97.5 %
(Intercept) 35.6364 32.2240 39.2814
treatdrug    0.2775  0.2192  0.3475

```

We can interpret this rate ratio as showing us that *the expected number of polyps for patients given the drug is about 0.28 times that of patients given the placebo (i.e., about 72% lower)*. We can also “flip” the rate ratio as follows.

```
trtools::contrast(m, tf = exp,
  a = list(treat = "placebo"), b = list(treat = "drug"))
```

```

estimate lower upper
3.604 2.863 4.536

```

This rate ratio can be interpreted by saying that *the expected number of polyps for patients given the placebo is about 3.6 times higher than patients given the drug (i.e., about 260% higher)*. Here is how to estimate these two rate ratios using the **emmeans** package.

```
pairs(emmeans(m, ~treat, type = "response"), infer = TRUE)
```

```

contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
placebo / drug  3.6 0.423 Inf      2.86      4.54    1  10.918 <.0001

```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~treat, type = "response"), infer = TRUE, reverse = TRUE)
```

```

contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
drug / placebo 0.277 0.0326 Inf      0.22      0.349    1 -10.918 <.0001

```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

Finally here is how we can estimate the expected number of polyps for each treatment condition.

```
trtools::contrast(m, tf = exp,
  a = list(treat = c("drug", "placebo")), cnames = c("drug", "placebo"))
```

```

      estimate lower upper
drug      9.889  8.034 12.17
placebo   35.636 32.278 39.34

```

```
emmeans(m, ~treat, type = "response")
```

```

treat  rate  SE  df asymp.LCL asymp.UCL
placebo 35.64 1.80 Inf      32.28      39.3
drug     9.89 1.05 Inf      8.03      12.2

```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

2. Estimate a Poisson regression model with the number of polyps as the response variable and both treatment and age as the explanatory variables. Do not include an interaction between treatment and

age.⁷ Estimate two rate ratios — one for the effect of the treatment and one for the effect of age.⁸ Write a sentence or two that interprets each rate ratio in terms of the relationship between the explanatory variable and the number of polyps. Also estimate the expected number of polyps for a 30-year old person who receives the treatment, and for someone of the same age that does not receive the treatment.

Solution: We can estimate the model as follows.

```
m <- glm(number ~ treat + age, family = poisson, data = polyps)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.52902	0.146872	30.84	8.487e-209
treatdrug	-1.35908	0.117643	-11.55	7.164e-31
age	-0.03883	0.005955	-6.52	7.016e-11

We can estimate the rate ratio for the treatment variable as follows.

```
trtools::contrast(m, tf = exp,
  a = list(treat = "drug", age = c(20,30,40)),
  b = list(treat = "placebo", age = c(20,30,40)),
  cnames = c("@20", "@30", "@40"))
```

	estimate	lower	upper
@20	0.2569	0.204	0.3235
@30	0.2569	0.204	0.3235
@40	0.2569	0.204	0.3235

Note that the age does not matter since there is no interaction between it and treatment, but I included several values to show this. This shows that *the expected number of polyps for a subject given the drug is about 0.26 times that of one of the same age but given the placebo (i.e., about 0.74% lower)*. Here is the “flipped” rate ratio.

```
trtools::contrast(m, tf = exp,
  a = list(treat = "placebo", age = c(20,30,40)),
  b = list(treat = "drug", age = c(20,30,40)),
  cnames = c("@20", "@30", "@40"))
```

	estimate	lower	upper
@20	3.893	3.091	4.902
@30	3.893	3.091	4.902
@40	3.893	3.091	4.902

So *the expected number of polyps for a patient given the placebo is about 3.9 times higher than one of the same age but given the drug (i.e., about 290% higher)*. We can estimate the rate ratio for the effect of as follows.

```
trtools::contrast(m, tf = exp,
  a = list(treat = c("drug", "placebo"), age = 30),
  b = list(treat = c("drug", "placebo"), age = 29),
  cnames = c("drug", "placebo"))
```

	estimate	lower	upper
drug	0.9619	0.9508	0.9732
placebo	0.9619	0.9508	0.9732

Because there is no interaction involving the treatment, the value of the treatment variable does not

⁷The interaction is negligible so we will ignore it.

⁸Note that since there is no interaction modeled between treatment and age the estimated rate ratio for one variable will not depend on the value of the other, so what value you use is arbitrary if you specify it using `contrast` or `emmeans`.

affect the rate ratio. This rate ratio can be interpreted as showing that *for every year of age the expected number of polyps decreases by a factor of about 0.96 (i.e., about 4%)*. The above rate ratios can be estimated using the **emmeans** package as follows.

```
pairs(emmeans(m, ~treat|age, at = list(age = c(20,30,40)),
  type = "response"), infer = TRUE)
```

```
age = 20:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
placebo / drug  3.89 0.458 Inf      3.09      4.9    1  11.553 <.0001
```

```
age = 30:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
placebo / drug  3.89 0.458 Inf      3.09      4.9    1  11.553 <.0001
```

```
age = 40:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
placebo / drug  3.89 0.458 Inf      3.09      4.9    1  11.553 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~treat, type = "response"), infer = TRUE)
```

```
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
placebo / drug  3.89 0.458 Inf      3.09      4.9    1  11.553 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~treat|age, at = list(age = c(20,30,40)),
  type = "response"), infer = TRUE, reverse = TRUE)
```

```
age = 20:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
drug / placebo 0.257 0.0302 Inf      0.204      0.324    1 -11.553 <.0001
```

```
age = 30:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
drug / placebo 0.257 0.0302 Inf      0.204      0.324    1 -11.553 <.0001
```

```
age = 40:
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
drug / placebo 0.257 0.0302 Inf      0.204      0.324    1 -11.553 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~treat, type = "response"), infer = TRUE, reverse = TRUE)
```

```
  contrast      ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
drug / placebo 0.257 0.0302 Inf      0.204      0.324    1 -11.553 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~age|treat, at = list(age = c(30,29)), type = "response"), infer = TRUE)
```

treat = placebo:

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
age30 / age29	0.962	0.00573	Inf	0.951	0.973	1	-6.520	<.0001

treat = drug:

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
age30 / age29	0.962	0.00573	Inf	0.951	0.973	1	-6.520	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
pairs(emmeans(m, ~age, at = list(age = c(30,29)), type = "response"), infer = TRUE)
```

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
age30 / age29	0.962	0.00573	Inf	0.951	0.973	1	-6.520	<.0001

Results are averaged over the levels of: treat

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

Finally here are some ways to estimate the expected number of polyps for a 30-year old patient in each treatment condition.

```
trtools::contrast(m, tf = exp,  
  a = list(treat = c("drug","placebo"), age = 30),  
  cnames = c("drug","placebo"))
```

	estimate	lower	upper
drug	7.426	5.912	9.329
placebo	28.908	25.455	32.829

```
emmeans(m, ~treat|age, at = list(age = 30), type = "response")
```

age = 30:

treat	rate	SE	df	asympt.LCL	asympt.UCL
placebo	28.91	1.876	Inf	25.45	32.83
drug	7.43	0.864	Inf	5.91	9.33

Confidence level used: 0.95


Intervals are back-transformed from the log scale

3. Using the model with both treatment and age as explanatory variables, plot two curves showing the estimated expected number of polyps as a function of treatment and age. Your plot should also include the raw data like the plot on the right-hand side above. You can use the code above that creates the object called p2 as the basis of your plot.

Solution: Here is the plot for the estimated model.

```
d <- expand.grid(treat = c("drug","placebo"), age = seq(13, 50, length = 100))  
d$yhat <- predict(m, newdata = d, type = "response")  
  
p <- ggplot(polyps, aes(x = age, y = number, color = treat)) + theme_minimal() +
```

```
geom_point() + labs(x = "Age (years)", y = NULL, color = "Treatment") +
  theme(legend.position = c(0.8, 0.8)) +
  geom_line(aes(y = yhat), data = d)
plot(p)
```



4. The model you estimated with both treatment and age as explanatory variables can be written as the nonlinear model

$$E(Y_i) = \exp(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2}),$$

where Y_i is the number of polyps, and x_{i1} and x_{i2} represent treatment and age (although which represents treatment and which represents age will depend on the order that you specified the explanatory variables in the model formula argument of `glm`). Estimate the model above using the `nls` function and show the parameter estimates and standard errors using the `summary` function. These should be similar but not equal to what you got using `glm` (note that you can use the estimates from `glm` as your starting values). The Poisson distribution assumes a variance structure such that $\text{Var}(Y_i) = E(Y_i)$, which suggests that if dealt with the heteroscedasticity using weights then we would use weights of $w_i = 1/E(Y_i)$ which can be approximated as $w_i = 1/\hat{y}_i$, where \hat{y}_i is the predicted value. Now estimate the model above again using `nls`, but use iteratively weighted least squares with the weights defined as above. Report the parameter estimates and their standard errors using `summary`. If you do this correctly the estimates should be the same as what you got from `glm`, but the standard errors will not necessarily be the same.⁹ **Note:** This problem is *extra credit* for students in Stat 436, but is *required* for students in Stat 516.

Solution: First I will estimate the model without using weights.

```
m <- nls(number ~ exp(b0 + b1 * (treat == "drug") + b2 * age),
  data = polyps, start = list(b0 = 4.5, b1 = -1.36, b2 = -0.04))
summary(m)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
b0	4.5674	0.37453	12.195	7.861e-10
b1	-1.3582	0.48022	-2.828	1.160e-02
b2	-0.0406	0.01707	-2.378	2.938e-02

⁹The reason that the standard errors are not the same is that Poisson regression assumes the variance structure $\text{Var}(Y_i) = E(Y_i)$ whereas using the weights as above assumes that $\text{Var}(Y_i) = \phi E(Y_i)$ for some unknown value of ϕ . That is, iteratively weighted least squares assumes that $\text{Var}(Y_i)$ is *proportional to* (rather than *equal to*) $E(Y_i)$, which we can write as $\text{Var}(Y_i) \propto E(Y_i)$. Interestingly it turns out that using iteratively weighted least squares with `nls` is effectively equivalent to using quasi-likelihood.

Now I will estimate it using iteratively weighted least squares.

```
polyps$w <- 1
for (i in 1:10) {
  m <- nls(number ~ exp(b0 + b1 * (treat == "drug") + b2 * age), data = polyps,
    start = list(b0 = 4.5, b1 = -1.36, b2 = -0.04), weights = w)
  polyps$w <- 1 / predict(m)
}
summary(m)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
b0	4.52902	0.48106	9.415	3.720e-08
b1	-1.35908	0.38533	-3.527	2.587e-03
b2	-0.03883	0.01951	-1.991	6.284e-02

Notice that the estimates are the same as when using `glm` to estimate a Poisson regression model, but the standard errors are different.

```
m <- glm(number ~ treat + age, family = poisson, data = polyps)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.52902	0.146872	30.84	8.487e-209
treatdrug	-1.35908	0.117643	-11.55	7.164e-31
age	-0.03883	0.005955	-6.52	7.016e-11

However, using iteratively weighted least squares here is equivalent to using a quasi-likelihood approach with `family = quasipoisson`.

```
m <- glm(number ~ treat + age, family = quasipoisson, data = polyps)
summary(m)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.52902	0.48106	9.415	3.720e-08
treatdrug	-1.35908	0.38533	-3.527	2.587e-03
age	-0.03883	0.01951	-1.991	6.284e-02

Note that both the estimates and the standard errors are the same as what was obtained when using iteratively weighted least squares.

Analgesic Potency

The data frame `analgesics` in the `trtools` package is from a study comparing four analgesics at varying doses.¹⁰ Here are the data in aggregated form.

```
library(trtools)
analgesics
```

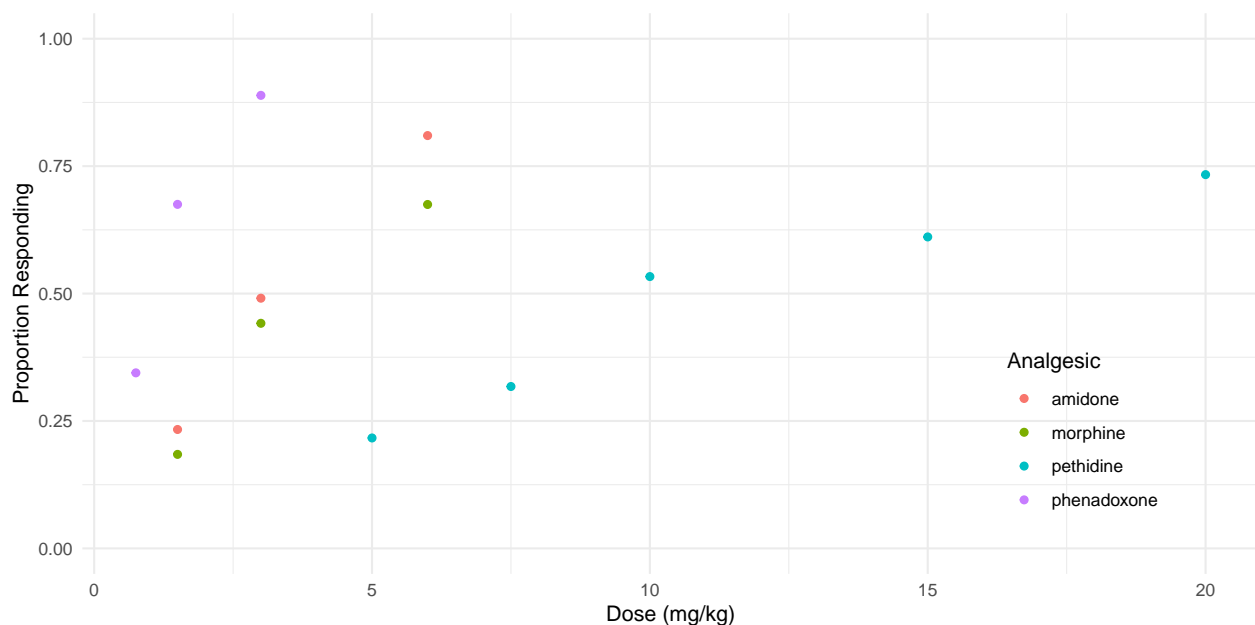
	analgesic	dose	responding	tested
1	morphine	1.50	19	103
2	morphine	3.00	53	120
3	morphine	6.00	83	123
4	amidone	1.50	14	60
5	amidone	3.00	54	110
6	amidone	6.00	81	100
7	pethidine	5.00	13	60

¹⁰Grewal, R. S. (1952). A method for testing analgesics in mice. *British Journal of Pharmacology and Chemotherapy*, 7, 433–437.

8	pethidine	7.50	27	85
9	pethidine	10.00	32	60
10	pethidine	15.00	55	90
11	pethidine	20.00	44	60
12	phenadoxone	0.75	31	90
13	phenadoxone	1.50	54	80
14	phenadoxone	3.00	80	90

Mice were randomly assigned one of four analgesics (morphine hydrochloride, amidone, phenadoxone, or pethidine hydrochloride) at one of nine doses. Each mouse was administered a number of electric shocks until a pain response (a squeak) was elicited. This was done before the administration of the analgesic (as a baseline) and after. A mouse was recorded as “responding” to an analgesic if at least four more shocks were required to elicit a pain response than before the analgesic was administered. The plot below shows the proportion of mice responding to the analgesic.

```
library(ggplot2)
p <- ggplot(analgesics, aes(x = dose, y = responding/tested, color = analgesic)) +
  theme_minimal() + geom_point() + theme(legend.position = c(0.85,0.25)) + ylim(0, 1) +
  labs(x = "Dose (mg/kg)", y = "Proportion Responding", color = "Analgesic")
plot(p)
```



The goal here is to model how the mice respond to the different analgesics using logistic regression.

1. Estimate a logistic regression model for the proportion of mice responding to the analgesics, using both the type of analgesic and the dose as explanatory variables. This model will include an interaction so that the odds ratio for dose depends on the type of analgesic used, but will also include a constraint that the probability of a response is the same at a dose of zero (since then the type of analgesic is irrelevant). To do this, specify the right-hand side of the model formula argument of the `glm` function as simply `analgesic:dose`. To verify that you estimated the model correctly, show the parameter estimates and their standard errors using `summary`.

Solution: We can estimate this model as follows.

```
m <- glm(cbind(responding, tested - responding) ~ analgesic:dose,
  family = binomial, data = analgesics)
summary(m)$coefficients
```

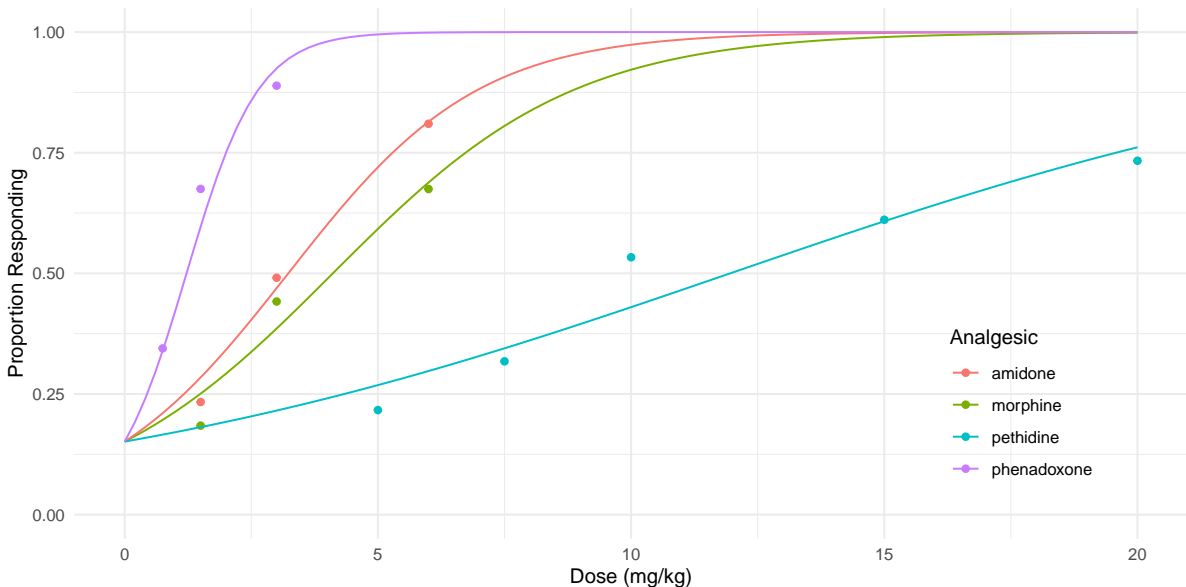
	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.7236	0.14566	-11.833	2.628e-32
analgesicamidone:dose	0.5336	0.04878	10.937	7.649e-28
analgesicmorphine:dose	0.4193	0.04255	9.853	6.636e-23
analgesicpethidine:dose	0.1442	0.01411	10.216	1.676e-24
analgesicphenadoxone:dose	1.4110	0.12569	11.226	3.046e-29

2. Using the model you estimated above, plot the estimated expected proportion of responses (which is also the estimated probability of a response) as a function of type of analgesic used and the dose. This plot should also include the observed proportions as shown above.

Solution: Here is a plot of the estimated model with the raw data.

```
d <- expand.grid(analgesic = unique(analgesics$analgesic), dose = seq(0, 20, length = 100))
d$yhat <- predict(m, newdata = d, type = "response")

p <- ggplot(analgesics, aes(x = dose, y = responding/tested, color = analgesic)) +
  theme_minimal() + geom_point() + theme(legend.position = c(0.85, 0.25)) + ylim(0, 1) +
  labs(x = "Dose (mg/kg)", y = "Proportion Responding", color = "Analgesic") +
  geom_line(aes(y = yhat), data = d)
plot(p)
```



3. Estimate the probability of a response for each analgesic at a dose of 0 mg/kg, and again at a dose of 5 mg/kg. Note that you should find that the estimates at 0 mg/kg are the same for the four analgesics.

Solution: These probabilities can be estimated as follows.

```
drug <- sort(unique(analgesics$analgesic))
trtools::contrast(m, tf = plogis,
  a = list(analgesic = drug, dose = 0),
  cnames = paste(drug, "at 0 mg/kg"))
```

	estimate	lower	upper
amidone at 0 mg/kg	0.1514	0.1182	0.1918
morphine at 0 mg/kg	0.1514	0.1182	0.1918
pethidine at 0 mg/kg	0.1514	0.1182	0.1918

```
phenadoxone at 0 mg/kg    0.1514 0.1182 0.1918
```

```
trtools::contrast(m, tf = plogis,
  a = list(analgesic = drug, dose = 5),
  cnames = paste(drug, "at 5 mg/kg"))
```

	estimate	lower	upper
amidone at 5 mg/kg	0.7200	0.6453	0.7842
morphine at 5 mg/kg	0.5921	0.5236	0.6573
pethidine at 5 mg/kg	0.2684	0.2308	0.3096
phenadoxone at 5 mg/kg	0.9952	0.9861	0.9983

Note that creating `drug` here lets me avoid having to type out all of the analgesics each time. Here is another way to do this.

```
emmeans(m, ~analgesic|dose, at = list(dose = c(0,5)), type = "response")
```

```
dose = 0:
  analgesic    prob      SE   df asymp.LCL asymp.UCL
amidone       0.151 0.01871 Inf     0.118    0.192
morphine      0.151 0.01871 Inf     0.118    0.192
pethidine     0.151 0.01871 Inf     0.118    0.192
phenadoxone   0.151 0.01871 Inf     0.118    0.192
```

```
dose = 5:
  analgesic    prob      SE   df asymp.LCL asymp.UCL
amidone       0.720 0.03558 Inf     0.645    0.784
morphine      0.592 0.03431 Inf     0.524    0.657
pethidine     0.268 0.02013 Inf     0.231    0.310
phenadoxone   0.995 0.00261 Inf     0.986    0.998
```

Confidence level used: 0.95

Intervals are back-transformed from the logit scale

4. Estimate the odds ratio for the effect of increasing dose by one unit for *each* of the four analgesics. For each odds ratio, write a sentence or two that interprets the odds ratio in terms of the relationship between dose and the response.

Solution: Here are a couple of ways to estimate these odds ratios.

```
drug <- sort(unique(analgesics$analgesic))
trtools::contrast(m, tf = exp,
  a = list(analgesic = drug, dose = 1),
  b = list(analgesic = drug, dose = 0),
  cnames = drug)
```

	estimate	lower	upper
amidone	1.705	1.550	1.876
morphine	1.521	1.399	1.653
pethidine	1.155	1.124	1.187
phenadoxone	4.100	3.205	5.245

```
pairs(emmeans(m, ~dose|analgesic, at = list(dose = c(1,0)), type = "response"), infer = TRUE)
```

```
analgesic = amidone:
  contrast      odds.ratio      SE   df asymp.LCL asymp.UCL null z.ratio p.value
dose1 / dose0      1.71 0.0832 Inf     1.55      1.88    1 10.937 <.0001
```



```

analgesic = morphine:
  contrast      odds.ratio      SE  df asymp.LCL asymp.UCL null z.ratio p.value
dose1 / dose0      1.52 0.0647 Inf      1.40      1.65      1  9.853 <.0001

```

```

analgesic = pethidine:
  contrast      odds.ratio      SE  df asymp.LCL asymp.UCL null z.ratio p.value
dose1 / dose0      1.16 0.0163 Inf      1.12      1.19      1 10.216 <.0001

```

```

analgesic = phenadoxone:
  contrast      odds.ratio      SE  df asymp.LCL asymp.UCL null z.ratio p.value
dose1 / dose0      4.10 0.5153 Inf      3.21      5.25      1 11.226 <.0001

```

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

We can interpret these four odds ratios as follows.

For every unit increase in dose, the odds of a response increases by a factor of about 1.71 (71%) when using amidone, 1.52 (52%) when using morphine hydrochloride, 1.16 (16%) when using pethidine hydrochloride, and 4.1 (310%) when using phenadoxone.

5. It can be seen from the output of `summary` for the model you estimated above that it can be written as

$$O_i = \begin{cases} e^{\beta_0} e^{\beta_1 d_i}, & \text{if the analgesic is amidone,} \\ e^{\beta_0} e^{\beta_2 d_i}, & \text{if the analgesic is morphine hydrochloride,} \\ e^{\beta_0} e^{\beta_3 d_i}, & \text{if the analgesic is pethidine,} \\ e^{\beta_0} e^{\beta_4 d_i}, & \text{if the analgesic is phenadoxone hydrochloride,} \end{cases}$$

where O_i is the odds of responding and d_i is the dose of the analgesic. The odds ratios for the effect of a one unit increase in dose for amidone, morphine hydrochloride, pethidine, and phenadoxone hydrochloride can be shown to be equal to e^{β_1} , e^{β_2} , e^{β_3} , and e^{β_4} , respectively. You estimated these odds ratios in the previous problem. Now consider the problem of testing if the odds ratio for, say, morphine hydrochloride is different from that for each of the other three analgesics. This would amount to testing each of the null hypotheses $\beta_1 - \beta_2 = 0$, $\beta_3 - \beta_2 = 0$, and $\beta_4 - \beta_2 = 0$. This can be done several ways. One is to use the `lincon` function from the `trtools` package. But it can also be done using `contrast` from the same package, or by using functions from the `emmeans` package. Test each of the three null hypotheses using a significance level of $\alpha = 0.05$. Be sure to state your conclusion. **Note:** This problem is *extra credit* for students in Stat 436, but is *required* for students in Stat 516.

Solution: Here is how we can do it using `lincon`.

```
lincon(m, a = c(0, 1, -1, 0, 0)) # b1 - b2 = 0
```

```

      estimate      se  lower  upper tvalue  df  pvalue
(0,1,-1,0,0),0  0.1143 0.04502 0.02605 0.2025  2.539 Inf 0.01113

```

```
lincon(m, a = c(0, 0, -1, 1, 0)) # b3 - b2 = 0
```

```

      estimate      se  lower  upper tvalue  df  pvalue
(0,0,-1,1,0),0 -0.2751 0.0364 -0.3465 -0.2038 -7.558 Inf 4.098e-14

```

```
lincon(m, a = c(0, 0, -1, 0, 1)) # b4 - b2 = 0
```

```

      estimate      se  lower  upper tvalue  df  pvalue
(0,0,-1,0,1),0  0.9917 0.1112 0.7737  1.21  8.916 Inf 4.822e-19

```

Note that all three null hypotheses are rejected at a significance level of 0.05. Using `contrast` is a bit trickier. Here I will just show how to do it to test $\beta_1 - \beta_2 = 0$ (i.e., comparing amidone and morphine).

```
trtools::contrast(m,
  a = list(analgesic = "amidone", dose = 1),
  b = list(analgesic = "amidone", dose = 0),
  u = list(analgesic = "morphine", dose = 1),
  v = list(analgesic = "morphine", dose = 0))
```

estimate	se	lower	upper	tvalue	df	pvalue
0.1143	0.04502	0.02605	0.2025	2.539	Inf	0.01113

Using the **emmeans** package we can compare all pairs of analgesics.

```
pairs(pairs(emmeans(m, ~dose|analgesic, at = list(dose = c(1,0))),
  by = NULL, adjust = "none"))
```

contrast	estimate	SE	df	z.ratio	p.value
(dose1 - dose0 amidone) - (dose1 - dose0 morphine)	0.114	0.0450	Inf	2.539	0.0111
(dose1 - dose0 amidone) - (dose1 - dose0 pethidine)	0.389	0.0430	Inf	9.064	<.0001
(dose1 - dose0 amidone) - (dose1 - dose0 phenadoxone)	-0.877	0.1121	Inf	-7.826	<.0001
(dose1 - dose0 morphine) - (dose1 - dose0 pethidine)	0.275	0.0364	Inf	7.558	<.0001
(dose1 - dose0 morphine) - (dose1 - dose0 phenadoxone)	-0.992	0.1112	Inf	-8.916	<.0001
(dose1 - dose0 pethidine) - (dose1 - dose0 phenadoxone)	-1.267	0.1193	Inf	-10.623	<.0001

Results are given on the log odds ratio (not the response) scale.

It is possible to target certain comparisons.

```
emmeans::contrast(pairs(emmeans(m, ~dose|analgesic, at = list(dose = c(1,0))),
  method = "trt.vs.ctrl", by = NULL, ref = 2, adjust = "none"))
```

contrast	estimate	SE	df	z.ratio	p.value
(dose1 - dose0 amidone) - (dose1 - dose0 morphine)	0.114	0.0450	Inf	2.539	0.0111
(dose1 - dose0 pethidine) - (dose1 - dose0 morphine)	-0.275	0.0364	Inf	-7.558	<.0001
(dose1 - dose0 phenadoxone) - (dose1 - dose0 morphine)	0.992	0.1112	Inf	8.916	<.0001

Results are given on the log odds ratio (not the response) scale.

Here is another approach using the **emtrends** function. First I will show all possible pairwise comparisons and then only those comparing against morphine.

```
pairs(emtrends(m, ~analgesic, var = "dose"), adjust = "none")
```

contrast	estimate	SE	df	z.ratio	p.value
amidone - morphine	0.114	0.0450	Inf	2.539	0.0111
amidone - pethidine	0.389	0.0430	Inf	9.064	<.0001
amidone - phenadoxone	-0.877	0.1121	Inf	-7.826	<.0001
morphine - pethidine	0.275	0.0364	Inf	7.558	<.0001
morphine - phenadoxone	-0.992	0.1112	Inf	-8.916	<.0001
pethidine - phenadoxone	-1.267	0.1193	Inf	-10.623	<.0001

```
emmeans::contrast(emtrends(m, ~analgesic, var = "dose"), method = "trt.vs.ctrl",
  ref = 2, adjust = "none")
```

contrast	estimate	SE	df	z.ratio	p.value
amidone - morphine	0.114	0.0450	Inf	2.539	0.0111
pethidine - morphine	-0.275	0.0364	Inf	-7.558	<.0001
phenadoxone - morphine	0.992	0.1112	Inf	8.916	<.0001

Here is another approach. Consider that if, for example, $\beta_1 - \beta_2 = 0$ then the expected response is the same for amidone and morphine hydrochloride *for any dose*. So we can actually test the null hypothesis

by testing the difference in the expected response between two analgesics for any given dose (other than zero where they are the same for this model).

```
trtools::contrast(m,
  a = list(analgesic = c("amidone", "pethidine", "phenadoxone"), dose = 5),
  b = list(analgesic = "morphine", dose = 5),
  cnames = c("amidone versus morphine", "pethidine versus morphine",
    "phenadoxone versus morphine"))
```

	estimate	se	lower	upper	tvalue	df	pvalue
amidone versus morphine	0.5715	0.2251	0.1302	1.013	2.539	Inf	1.113e-02
pethidine versus morphine	-1.3756	0.1820	-1.7323	-1.019	-7.558	Inf	4.098e-14
phenadoxone versus morphine	4.9586	0.5561	3.8686	6.049	8.916	Inf	4.822e-19

```
emmeans::contrast(emmeans(m, ~analgesic|dose, at = list(dose = 5)),
  method = "trt.vs.ctrl", ref = 2, adjust = "none")
```

```
dose = 5:
contrast      estimate    SE  df z.ratio p.value
amidone - morphine      0.571 0.225 Inf   2.539  0.0111
pethidine - morphine    -1.376 0.182 Inf  -7.558 <.0001
phenadoxone - morphine   4.959 0.556 Inf   8.916 <.0001
```

Results are given on the log odds ratio (not the response) scale.

Note that since we are just testing the differences between odds ratios there is no need to use `tf = exp` in `lincon` or `contrast`, or to use `type = "response"` with the `emmeans` package.

Here is another approach using a reparameterization of the model.

```
analgesics$analgesic <- relevel(factor(analgesics$analgesic), ref = "morphine")
m <- glm(cbind(responding, tested - responding) ~ dose + dose:analgesic,
  family = binomial, data = analgesics)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.7236	0.14566	-11.833	2.628e-32
dose	0.4193	0.04255	9.853	6.636e-23
dose:analgesicamidone	0.1143	0.04502	2.539	1.113e-02
dose:analgesicpethidine	-0.2751	0.03640	-7.558	4.098e-14
dose:analgesicphenadoxone	0.9917	0.11123	8.916	4.822e-19

This model can be written as

$$O_i = \begin{cases} e^{\beta_0} e^{(\beta_1 + \beta_2)d_i}, & \text{if the analgesic is amidone,} \\ e^{\beta_0} e^{\beta_1 d_i}, & \text{if the analgesic is morphine hydrochloride,} \\ e^{\beta_0} e^{(\beta_1 + \beta_3)d_i}, & \text{if the analgesic is pethidine,} \\ e^{\beta_0} e^{(\beta_1 + \beta_4)d_i}, & \text{if the analgesic is phenadoxone hydrochloride,} \end{cases}$$

so the hypotheses are now written as $\beta_2 = 0$, $\beta_3 = 0$, and $\beta_4 = 0$. These can be tested just using the output from `summary`.

Here is one more approach, but this time using a likelihood ratio test. Here I will demonstrate it for testing the null hypothesis $\beta_1 - \beta_2 = 0$. To do this we need to specify a null model that is implied by the null hypothesis. In that model there is no distinction between amidone or morphine hydrochloride. To do this we can create a new variable by combining two levels.

```
library(forcats)
analgesics <- analgesics %>%
  mutate(adrug = fct_collapse(analgesic,
    amidone_or_morphine = c("amidone", "morphine"),
  ))
analgesics
```

	analgesic	dose	responding	tested	adrug
1	morphine	1.50	19	103	amidone_or_morphine
2	morphine	3.00	53	120	amidone_or_morphine
3	morphine	6.00	83	123	amidone_or_morphine
4	amidone	1.50	14	60	amidone_or_morphine
5	amidone	3.00	54	110	amidone_or_morphine
6	amidone	6.00	81	100	amidone_or_morphine
7	pethidine	5.00	13	60	pethidine
8	pethidine	7.50	27	85	pethidine
9	pethidine	10.00	32	60	pethidine
10	pethidine	15.00	55	90	pethidine
11	pethidine	20.00	44	60	pethidine
12	phenadoxone	0.75	31	90	phenadoxone
13	phenadoxone	1.50	54	80	phenadoxone
14	phenadoxone	3.00	80	90	phenadoxone

```
m <- glm(cbind(responding, tested - responding) ~ analgesic:dose,
  family = binomial, data = analgesics)
m.null <- glm(cbind(responding, tested - responding) ~ adrug:dose,
  family = binomial, data = analgesics)
anova(m.null, m, test = "LRT")
```

Analysis of Deviance Table

```
Model 1: cbind(responding, tested - responding) ~ adrug:dose
Model 2: cbind(responding, tested - responding) ~ analgesic:dose
  Resid. Df Resid. Dev Df Deviance Pr(>Chi)
1         10         19.4
2          9         12.8  1     6.59    0.01 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Race and the Death Penalty

The data frame `deathpenalty` in the package `catdata` tabulates judgments of the death penalty in cases of murder in Florida between 1976 and 1987.¹¹ The race of the defendant and victim were also recorded. The data are in a form that is not very convenient for modeling.

```
library(catdata)
data(deathpenalty)
deathpenalty
```

	DeathPenalty	VictimRace	DefendantRace	Freq
1	0	0	0	139
2	1	0	0	4
3	0	1	0	37

¹¹Radelet, M. & Pierce, G. L. (1991). Choosing those who will die: Race and the death penalty in Florida, *Florida Law Review*, 43(1), 1–34.

4	1	1	0	11
5	0	0	1	16
6	1	0	1	0
7	0	1	1	414
8	1	1	1	53

Here values of 0 and 1 correspond to *no* and *yes*, respectively, for the variable `DeathPenalty`, and *black* and *white*, respectively, for the variables `VictimRace` and `DefendantRace`. It is useful to make these variables factors with clearly labeled values. To avoid confusion I will create a new data frame called `dpenalty`.

```
library(dplyr)
dpenalty <- deathpenalty %>%
  mutate(DeathPenalty = factor(DeathPenalty, levels = c(0,1), labels = c("no","yes"))) %>%
  mutate(VictimRace = factor(VictimRace, levels = c(0,1), labels = c("black","white"))) %>%
  mutate(DefendantRace = factor(DefendantRace, levels = c(0,1), labels = c("black", "white")))
dpenalty
```

	DeathPenalty	VictimRace	DefendantRace	Freq
1	no	black	black	139
2	yes	black	black	4
3	no	white	black	37
4	yes	white	black	11
5	no	black	white	16
6	yes	black	white	0
7	no	white	white	414
8	yes	white	white	53

Consider using logistic regression to model these data with `DeathPenalty` as the response variable. It is possible to model the data in this form, and I will show you how that can be done in the solutions. But I will put the data in a form that is more familiar to you when using logistic regression.

```
library(tidyr)
dpenalty <- dpenalty %>%
  pivot_wider(names_from = DeathPenalty, values_from = Freq)
dpenalty
```

```
# A tibble: 4 x 4
  VictimRace DefendantRace    no    yes
  <fct>      <fct>      <int> <int>
1 black     black       139     4
2 white     black        37    11
3 black     white        16     0
4 white     white       414    53
```

Here the data are in the familiar aggregated form where for each combination of `VictimRace` and `DefendantRace` we have the number of cases where the death penalty was and was not decided. These data can be used to demonstrate what is sometimes called Simpson's Paradox which occurs when the association between two variables reverses when we collapse across a third variable. Here you will use logistic regression with the data frame `dpenalty` to explore this phenomenon.

1. First consider using only paying attention to the defendant's race. We can compute the proportion of cases where the death penalty was imposed as follows.

```
dpenalty %>% group_by(DefendantRace) %>%
  summarize(no = sum(no), yes = sum(yes)) %>%
  mutate(proportion = yes / (no + yes))
```

```
# A tibble: 2 x 4
```

	DefendantRace	no	yes	proportion
	<fct>	<int>	<int>	<dbl>
1	black	176	15	0.0785
2	white	430	53	0.110

From this we can see that a slightly larger proportion of white defendants were given the death penalty in comparison to black defendants. Estimate a logistic regression model with the proportion of judgments resulting in the death penalty as your response variable, but using *only the defendant's race* as an explanatory variable. Report the parameter estimates and their standard errors using `summary` so I can verify that you estimated this model correctly. Estimate (a) the probability of a death penalty decision for each race of the defendant, and (b) the odds ratio that summarizes the relationship between the defendant's race and the death penalty decision. Summarize the odds ratio in a sentence or two that describes the relationship between the race of the defendant and the death penalty decision.

Solution: This model can be estimated as follows.

```
m <- glm(cbind(yes, no) ~ DefendantRace, family = binomial, data = dpenalty)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.4624	0.2690	-9.155	5.444e-20
DefendantRacewhite	0.3689	0.3058	1.206	2.277e-01

The probability of the death penalty conditional on the defendant's race can be estimated several ways.

```
trtools::contrast(m, tf = plogis,
  a = list(DefendantRace = c("black", "white")))
```

	estimate	lower	upper
black	0.07853	0.0479	0.1262
white	0.10973	0.0848	0.1409

```
emmeans(m, ~DefendantRace, type = "response")
```

	DefendantRace	prob	SE	df	asympt.LCL	asympt.UCL
black	black	0.0785	0.0195	Inf	0.0479	0.126
white	white	0.1097	0.0142	Inf	0.0848	0.141

Confidence level used: 0.95

Intervals are back-transformed from the logit scale

Note that for this particular model these estimates are equivalent to the sample proportions. Here is the estimated odds ratio.

```
trtools::contrast(m, tf = exp,
  a = list(DefendantRace = "black"),
  b = list(DefendantRace = "white"))
```

	estimate	lower	upper
black / white	0.6915	0.3797	1.259

```
pairs(emmeans(m, ~DefendantRace, type = "response"), infer = TRUE)
```

	contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
black / white	black / white	0.692	0.211	Inf	0.38	1.26	1	-1.206	0.2277

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

Thus we estimate that *the odds of the death penalty being given to a black defendant is 0.69 times that of a white defendant (i.e., 31% lower)*. Alternatively we can “flip” the odds ratio as follows.

```
trtools::contrast(m, tf = exp,
  a = list(DefendantRace = "white"),
  b = list(DefendantRace = "black"))

estimate lower upper
1.446 0.7941 2.634

pairs(emmeans(m, ~DefendantRace, type = "response"), infer = TRUE, reverse = TRUE)

contrast odds.ratio SE df asymp.LCL asymp.UCL null z.ratio p.value
white / black 1.45 0.442 Inf 0.794 2.63 1 1.206 0.2277
```

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

This says that we estimate that *the odds of the death penalty being given to a white defendant is 1.45 times larger (i.e., 45% higher) than for a black defendant*. Note that we can also get this odds ratio by just exponentiating the parameters.

```
exp(cbind(coef(m), confint(m)))

                2.5 % 97.5 %
(Intercept)      0.08523 0.04818 0.1393
DefendantRacewhite 1.44620 0.81342 2.7198
```

2. Now consider all conditioning on the race of the *victim*. We can compute the proportion of cases where the death penalty was imposed as follows.

```
dpenalty %>% mutate(proportion = yes / (no + yes)) %>% arrange(VictimRace)

# A tibble: 4 x 5
  VictimRace DefendantRace   no   yes proportion
  <fct>      <fct>      <int> <int>      <dbl>
1 black     black      139     4     0.0280
2 black     white       16     0     0
3 white     black       37    11     0.229
4 white     white      414    53     0.113
```

Note that now a larger proportion of black defendants were given the death penalty when the victim is black, but also when the victim is white. This apparent reversal is due to two things. First, the death penalty is given more often when the *victim* is white, and secondly there is an association between the race of the defendant and that of the victim which can be seen in the original data. In a majority of the cases the race of the victim and the defendant are the same as can be seen below.

```
dpenalty %>% mutate(total = no + yes)

# A tibble: 4 x 5
  VictimRace DefendantRace   no   yes total
  <fct>      <fct>      <int> <int> <int>
1 black     black      139     4   143
2 white     black       37    11    48
3 black     white       16     0    16
4 white     white      414    53   467
```

To account for the race of the victim as well as the defendant, estimate a logistic like that you estimated in the previous problem, but now using *both* the race of the defendant *and* the race of the victim as

explanatory variables, but do not include an interaction.¹² Report the parameter estimates and their standard errors using `summary` so that I can verify that you estimated this model correctly. Estimate (a) the probability of a death penalty decision for each combination of levels of victim's race and defendant's race, (b) the odds ratio that summarizes the relationship between the defendant's race and the death penalty decision, and (c) the odds ratio that summarizes the relationship between the victim's race and the death penalty decision. Summarize each odds ratio in a sentence or two that describes the relationship between the explanatory variable and the death penalty decision.

Solution: Here is how to estimate this model.

```
m <- glm(cbind(yes, no) ~ DefendantRace + VictimRace, family = binomial, data = dpenalty)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.5961	0.5069	-7.094	1.302e-12
DefendantRacewhite	-0.8678	0.3671	-2.364	1.807e-02
VictimRacewhite	2.4044	0.6006	4.003	6.247e-05

We can estimate the probabilities as follows.

```
trtools::contrast(m, tf = plogis,
  a = list(DefendantRace = c("black", "white", "black", "white"),
    VictimRace = c("black", "black", "white", "white")),
  cnames = c("black D, black V", "white D, black V", "black D, white V", "white D, white V"))
```

	estimate	lower	upper
black D, black V	0.02670	0.010054	0.06897
white D, black V	0.01139	0.003433	0.03708
black D, white V	0.23296	0.135367	0.37075
white D, white V	0.11310	0.087439	0.14509

```
emmeans(m, ~DefendantRace*VictimRace, type = "response")
```

DefendantRace	VictimRace	prob	SE	df	asympt.LCL	asympt.UCL
black	black	0.0267	0.01317	Inf	0.01005	0.0690
white	black	0.0114	0.00693	Inf	0.00343	0.0371
black	white	0.2330	0.06041	Inf	0.13537	0.3707
white	white	0.1131	0.01463	Inf	0.08744	0.1451

Confidence level used: 0.95

Intervals are back-transformed from the logit scale

Here are the odds ratios for the effect of defendant's race, when controlling for the victim's race.

```
trtools::contrast(m, tf = exp,
  a = list(DefendantRace = "black", VictimRace = c("black", "white")),
  b = list(DefendantRace = "white", VictimRace = c("black", "white")),
  cnames = c("black victim", "white victim"))
```

	estimate	lower	upper
black victim	2.382	1.16	4.89
white victim	2.382	1.16	4.89

```
pairs(emmeans(m, ~DefendantRace|VictimRace, type = "response"))
```

VictimRace = black:

contrast	odds.ratio	SE	df	null	z.ratio	p.value
black / white	2.38	0.874	Inf	1	2.364	0.0181

¹²The interaction is negligible so for simplicity we will ignore it.


```
VictimRace = white:
  contrast      odds.ratio    SE  df null z.ratio p.value
black / white      2.38 0.874 Inf    1   2.364  0.0181
```

Tests are performed on the log odds ratio scale

```
trtools::contrast(m, tf = exp,
  a = list(DefendantRace = "white", VictimRace = c("black","white")),
  b = list(DefendantRace = "black", VictimRace = c("black","white")),
  cnames = c("black victim", "white victim"))
```

```
          estimate lower upper
black victim  0.4199 0.2045 0.8621
white victim  0.4199 0.2045 0.8621
```

```
pairs(emmeans(m, ~DefendantRace|VictimRace, type = "response"), reverse = TRUE)
```

```
VictimRace = black:
  contrast      odds.ratio    SE  df null z.ratio p.value
white / black      0.42 0.154 Inf    1  -2.364  0.0181
```

```
VictimRace = white:
  contrast      odds.ratio    SE  df null z.ratio p.value
white / black      0.42 0.154 Inf    1  -2.364  0.0181
```

Tests are performed on the log odds ratio scale

Thus the odds of the death penalty for a black defendant is about 2.38 times higher (i.e., 138% larger) than for a white defendant when controlling for the race of the victim. Alternatively, we might say that the odds of the death penalty for a white defendant is about 0.42 times that of a black defendant (i.e., 58% smaller) when controlling for the race of the victim.

Here are the odds ratios for the effect of the victim's race, when controlling for the defendant's race.

```
trtools::contrast(m, tf = exp,
  a = list(VictimRace = "black", DefendantRace = c("black","white")),
  b = list(VictimRace = "white", DefendantRace = c("black","white")),
  cnames = c("black victim", "white victim"))
```

```
          estimate lower upper
black victim  0.09032 0.02783 0.2931
white victim  0.09032 0.02783 0.2931
```

```
pairs(emmeans(m, ~VictimRace|DefendantRace, type = "response"))
```

```
DefendantRace = black:
  contrast      odds.ratio    SE  df null z.ratio p.value
black / white      0.0903 0.0542 Inf    1  -4.003  0.0001
```

```
DefendantRace = white:
  contrast      odds.ratio    SE  df null z.ratio p.value
black / white      0.0903 0.0542 Inf    1  -4.003  0.0001
```

Tests are performed on the log odds ratio scale

```
trtools::contrast(m, tf = exp,
  a = list(VictimRace = "white", DefendantRace = c("black","white")),
```

```
b = list(VictimRace = "black", DefendantRace = c("black","white")),
cnames = c("black victim", "white victim"))

      estimate lower upper
black victim    11.07 3.412 35.93
white victim    11.07 3.412 35.93

pairs(emmeans(m, ~VictimRace|DefendantRace, type = "response"), reverse = TRUE)
```

```
DefendantRace = black:
contrast      odds.ratio    SE   df null z.ratio p.value
white / black      11.1 6.65 Inf    1   4.003  0.0001
```

```
DefendantRace = white:
contrast      odds.ratio    SE   df null z.ratio p.value
white / black      11.1 6.65 Inf    1   4.003  0.0001
```

Tests are performed on the log odds ratio scale

These odds ratios show that *the odds of the death penalty when the victim is black is about 0.09 times that when the victim is white (i.e., about 91% smaller) when controlling for the race of the defendant.* Or we could say that *the odds of the death penalty when the victim is white is about 11.1 times that when the victim is black (i.e., 1010% larger) when controlling for the race of the defendant.*

Note that the expression “controlling for” means that we are estimating the odds ratio for one variable when the other variable does not change. We can do this in the second model, because the other variable is also an explanatory variable which we can hold constant when defining an odds ratio. But this was not the case in the first model where we could hold constant the race of the victim.

Here I will show you how to specify the model with the data in their original form with one frequency per row. Here is how those data looked.

```
dpenalty <- dpenalty %>% pivot_longer(c(no,yes),
  names_to = "DeathPenalty", values_to = "Freq")
dpenalty
```

```
# A tibble: 8 x 4
  VictimRace DefendantRace DeathPenalty  Freq
  <fct>      <fct>          <chr>      <int>
1 black     black         no         139
2 black     black         yes          4
3 white     black         no          37
4 white     black         yes          11
5 black     white          no          16
6 black     white         yes           0
7 white     white          no         414
8 white     white         yes          53
```

To handle the data in this form we can use weights to indicate the frequency of each observation. Note that the output of `summary` is the same as before.

```
m <- glm(DeathPenalty == "yes" ~ DefendantRace + VictimRace,
  family = binomial, weights = Freq, data = dpenalty)
summary(m)$coefficients
```

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)   -3.5961     0.5069  -7.095 1.296e-12
DefendantRacewhite -0.8678     0.3671  -2.364 1.807e-02
```

VictimRacewhite	2.4044	0.6006	4.004	6.240e-05
-----------------	--------	--------	-------	-----------