

# Categorical Response Variables and the Incidental Parameters Problem

Statistics 516, Homework 5 (Solutions)

## Lip Cancer in Scotland: The Random Effect of District


Recall the model for the lip cancer data from Scotland used in the third and fourth homework assignments. Here again is the code to process and plot the raw data showing the observed rate of lip cancer per person-year by percent of the population engaged in outdoor activity.

```
library(epiR)
library(dplyr)
library(ggplot2)

data(epi.SClip)

lipcancer <- epi.SClip %>%
  mutate(district = factor(district, levels = rev(sort(unique(district))))) %>%
  mutate(percent = paste(prop.ag, "%", sep = "")) %>%
  mutate(percent = reorder(percent, prop.ag)) %>%
  select(district, cases, population, percent)

p <- ggplot(lipcancer, aes(y = district, x = cases/population)) +
  theme_minimal() + geom_point(aes(size = population), shape = 21) +
  facet_grid(percent ~ ., scales = "free_y", space = "free_y") +
  labs(y = NULL, x = "Cases per Person-Year", size = "Person-Years:") +
  scale_x_continuous(labels = scales::label_number()) +
  theme(axis.text.y = element_text(size = 7), legend.position = "top")
plot(p)
```



In the previous homework assignment you noted that there may be some over-dispersion which may be due to variation in the lip cancer rate over districts other than that accounted for by the percent of the population engaged in outdoor activity. To account for that over-dispersion you used a quasi-likelihood approach to “adjust” inferences to the over-dispersion. Another approach that could be used is to specify the effect of district as a *random effect*.<sup>1</sup>

1. Estimate a Poisson regression model like you did in the third homework assignment, but this time specifying a random “main effect” for district. Report the model parameter estimates and standard errors using `summary` or `lincon` so that I can verify that you estimated the model correctly.

**Solution:** Here is how to estimate this mixed effects model.

```
library(lme4)
m <- glmer(cases ~ percent + offset(log(population)) + (1 | district),
  family = poisson, data = lipcancer)
summary(m)
```

```
Generalized linear mixed model fit by maximum likelihood (Laplace Approximation) ['glmerMod']
Family: poisson ( log )
Formula: cases ~ percent + offset(log(population)) + (1 | district)
Data: lipcancer
```

<sup>1</sup>In most cases when random effects are specified there are two or more observations per level of the factor(s) that define the random effect(s). Specifying the random effects can then help account for a lack of independence of observations from the same experimental/observational unit. But in some cases like this one a random effect can induce over-dispersion without dependencies among the observations because there is only one observation per unit.

```

      AIC      BIC   logLik deviance df.resid
355.7    369.9   -170.9    341.7      49

Scaled residuals:
    Min       1Q   Median       3Q      Max
-1.5491 -0.4031 -0.0435  0.4107  1.4272

Random effects:
 Groups   Name      Variance Std.Dev.
district (Intercept) 0.339    0.582
Number of obs: 56, groups: district, 56

Fixed effects:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -11.0362     0.3251  -33.95 < 2e-16 ***
percent1%    -0.0687     0.3932   -0.17  0.86127
percent7%     1.1837     0.3723    3.18  0.00148 **
percent10%    1.1775     0.3821    3.08  0.00206 **
percent16%    1.3654     0.3915    3.49  0.00049 ***
percent24%    1.8539     0.4770    3.89  0.00010 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:
          (Intr) prcn1% prcn7% prc10% prc16%
percent1% -0.825
percent7% -0.873  0.722
percent10% -0.850  0.704  0.743
percent16% -0.830  0.688  0.725  0.707
percent24% -0.681  0.564  0.595  0.580  0.566

trtools::lincon(m)

```

- Using contrast or the **emmeans** package, estimate the expected number of cases of lip cancer per 100K person-years for each value of the percent of the population spent in outdoor activities.

**Solution:** Here are the estimated rates of cases of lip cancer (per 100K person-years).

```

library(emmeans)
emmeans(m, ~percent, type = "response", offset = log(100000))

percent rate   SE df asymp.LCL asymp.UCL
0%      1.61 0.524 Inf    0.852    3.05
1%      1.50 0.334 Inf    0.973    2.32
7%      5.26 0.956 Inf    3.685    7.51
10%     5.23 1.051 Inf    3.526    7.75
16%     6.31 1.379 Inf    4.111    9.69

```

24%	10.28	3.591	Inf	5.187	20.39
-----	-------	-------	-----	-------	-------

Confidence level used: 0.95

Intervals are back-transformed from the log scale

3. In the third homework assignment I asked you to estimate and interpret rate ratios that compared the rate of lip cancer at 1%, 7%, 10%, 16%, and 24% versus 0% of the population involved in outdoor activity. This time estimate and interpret rate ratios that compare *consecutive* values of the percent of the population involved in outdoor activity (i.e., 0% versus 1%, 1% versus 7%, 7% versus 10%, and so on).

**Solution:** I will show you a couple of ways to do this. One is to use `contrast`.

```
trtools::contrast(m, tf = exp,
  a = list(percent = c("1%", "7%", "10%", "16%", "24%"), population = 1),
  b = list(percent = c("0%", "1%", "7%", "10%", "16%"), population = 1),
  cnames = c("1% vs 0%", "7% vs 1%", "10% vs 7%", "16% vs 10%", "24% vs 16%"))
```

	estimate	lower	upper
1% vs 0%	0.9336	0.4319	2.018
7% vs 1%	3.4987	1.9964	6.132
10% vs 7%	0.9938	0.5846	1.689
16% vs 10%	1.2068	0.6750	2.157
24% vs 16%	1.6298	0.7275	3.651

With a little programming effort we can avoid so much tedious typing.

```
pa <- sort(unique(lipcancer$percent))[2:6]
pb <- sort(unique(lipcancer$percent))[1:5]
trtools::contrast(m, tf = exp,
  a = list(percent = pa, population = 1),
  b = list(percent = pb, population = 1),
  cnames = paste(pa, "vs", pb))
```

	estimate	lower	upper
1% vs 0%	0.9336	0.4319	2.018
7% vs 1%	3.4987	1.9964	6.132
10% vs 7%	0.9938	0.5846	1.689
16% vs 10%	1.2068	0.6750	2.157
24% vs 16%	1.6298	0.7275	3.651

The `contrast` function in the **emmeans** package can also be used to do consecutive paired comparisons.

```
emmeans::contrast(emmeans(m, ~percent, offset = log(1), type = "response"),
  method = "consec", infer = TRUE, adjust = "none")
```

contrast	ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
1% / 0%	0.934	0.367	Inf	0.432	2.02	1	-0.175	0.8613
7% / 1%	3.499	1.002	Inf	1.996	6.13	1	4.375	<.0001
10% / 7%	0.994	0.269	Inf	0.585	1.69	1	-0.023	0.9818
16% / 10%	1.207	0.358	Inf	0.675	2.16	1	0.634	0.5261
24% / 16%	1.630	0.671	Inf	0.728	3.65	1	1.187	0.2353

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

Here is shortcut that gives both the estimated rates and consecutive pairwise comparisons.

```
emmeans(m, consec ~ percent, offset = log(100000),
  type = "response", infer = TRUE, adjust = "none")
```

```
$emmeans
percent rate SE df asymp.LCL asymp.UCL null z.ratio p.value
0%      1.61 0.524 Inf 0.852 3.05 1 1.466 0.1426
1%      1.50 0.334 Inf 0.973 2.32 1 1.838 0.0661
7%      5.26 0.956 Inf 3.685 7.51 1 9.142 <.0001
10%     5.23 1.051 Inf 3.526 7.75 1 8.229 <.0001
16%     6.31 1.379 Inf 4.111 9.69 1 8.426 <.0001
24%    10.28 3.591 Inf 5.187 20.39 1 6.675 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
$contrasts
contrast ratio SE df asymp.LCL asymp.UCL null z.ratio p.value
1% / 0% 0.934 0.367 Inf 0.432 2.02 1 -0.175 0.8613
7% / 1% 3.499 1.002 Inf 1.996 6.13 1 4.375 <.0001
10% / 7% 0.994 0.269 Inf 0.585 1.69 1 -0.023 0.9818
16% / 10% 1.207 0.358 Inf 0.675 2.16 1 0.634 0.5261
24% / 16% 1.630 0.671 Inf 0.728 3.65 1 1.187 0.2353
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

Note that the value of the offset does not affect the rate ratios, but it does affect the estimated rates as I had to specify it here to give the estimated number of cases of lip cancer per 100K person-years. Here is how we could interpret these rate ratios.

*The estimated rate of cases of lip cancer is about 6.6% less at 1% percent of the population engaged in outdoor activity than it is at 0%.*

*The estimated rate of cases of lip cancer is about 3.5 times higher (250% higher) at 7% percent of the population engaged in outdoor activity than it is at 1%.*

*The estimated rate of cases of lip cancer is about 1% less at 10% percent of the population engaged in outdoor activity than it is at 7%.*

*The estimated rate of cases of lip cancer is about 1.21 times higher (21% higher) at 16% percent of the population engaged in outdoor activity than it is at 10%.*

*The estimated rate of cases of lip cancer is about 1.63 times higher (63% higher) at 24% percent of the population engaged in outdoor activity than it is at 16%.*

It may be worth noting that only one of these rate ratios is statistically significant at conventional significance levels. We can also “flip” the rate ratios to put the lower percent in the numerator of the ratio.

```
emmeans(m, consec ~ percent, offset = log(100000),
  type = "response", infer = TRUE, adjust = "none", reverse = TRUE)$contrasts
```

```
contrast ratio SE df asymp.LCL asymp.UCL null z.ratio p.value
0% / 1% 1.071 0.4212 Inf 0.496 2.315 1 0.175 0.8613
1% / 7% 0.286 0.0818 Inf 0.163 0.501 1 -4.375 <.0001
7% / 10% 1.006 0.2724 Inf 0.592 1.710 1 0.023 0.9818
```

10% / 16%	0.829	0.2456	Inf	0.464	1.481	1	-0.634	0.5261
16% / 24%	0.614	0.2525	Inf	0.274	1.375	1	-1.187	0.2353

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

We can interpret these rate ratios as follows.

*The estimated rate of cases of lip cancer is about 1.07 times higher (7% higher) at 0% percent of the population engaged in outdoor activity than it is at 1%.*

*The estimated rate of cases of lip cancer is about 73% lower at 1% percent of the population engaged in outdoor activity than it is at 7%.*

*The estimated rate of cases of lip cancer is about 1.006 times higher (0.6% higher) at 7% percent of the population engaged in outdoor activity than it is at 10%.*

*The estimated rate of cases of lip cancer is about 17% lower at 10% percent of the population engaged in outdoor activity than it is at 16%.*

*The estimated rate of cases of lip cancer is about 39% lower at 16% percent of the population engaged in outdoor activity than it is at 24%.*

## Comparing Methods of Measuring Blood Oxygen Saturation

The data frame `ox` in the **MethComp** package is from a study comparing two methods of measuring blood oxygen saturation: chemical analysis or using a pulse oximeter.<sup>2</sup> Measurements were taken using both methods for each of 61 children. Multiple readings (replicates) were taken using each method in quick succession. For most children there were three replicates per method for a total of six measurements per child. Here are first few observations, sorted by the child (`item`).

```
library(dplyr)
library(tibble) # to use the remove_rownames function
library(MethComp)
data(ox)
ox <- ox %>% arrange(item) %>% remove_rownames()
head(ox, 12)
```

	meth	item	repl	y
1	CO	1	1	78.0
2	CO	1	2	76.4
3	CO	1	3	77.2
4	pulse	1	1	71.0
5	pulse	1	2	72.0
6	pulse	1	3	73.0
7	CO	2	1	68.7
8	CO	2	2	67.6
9	CO	2	3	68.3
10	pulse	2	1	68.0
11	pulse	2	2	67.0
12	pulse	2	3	68.0

Note that here `meth` is the method of measuring oxygen saturation (`CO` is chemical analysis and `pulse` is pulse oximeter), `item` identifies the child, `repl` is the replicate, and `y` is the measurement of oxygen saturation. The plot below shows the raw data.

<sup>2</sup>Source: Carstensen, B. (2010). *Comparing clinical measurement methods: A practical guide*. Wiley.

```
library(ggplot2)
p <- ggplot(ox, aes(y = factor(item), x = y, color = meth)) +
  theme_minimal() + geom_point(alpha = 0.5) + scale_y_discrete(limits = rev) +
  labs(y = "Child Identifier", x = "Oxygen Saturation (Percent)", color = "Method")
plot(p)
```




The plot shows some evidence of an “effect” for the child — i.e., children may naturally vary with respect to their average blood oxygen saturation. This can induce a lack of independence of the observations among those from the same child. We can also see this in a scatter plot of the average measurement for each child by method.<sup>3</sup>

<sup>3</sup>Using `coord_fixed()` forces a one-to-one aspect ratio of the axes, which is appropriate here since they are on the same scale.

```
library(tidyr) # to use the pivot_wider function
oxaverages <- ox %>% group_by(item, meth) %>%
  summarize(y = mean(y)) %>% pivot_wider(names_from = meth, values_from = y)

p <- ggplot(oxaverages, aes(x = CO, y = pulse)) + theme_minimal() +
  geom_abline(intercept = 0, slope = 1) +
  geom_point(alpha = 0.5) + coord_fixed() +
  labs(x = "Chemical Analysis", y = "Pulse Oximeter")

plot(p)
```



The lack of independence is suggested by the strong correlation between the average measurements. Here you will consider some different ways to account for this (or not).

1. Estimate the following four linear models with the oxygen saturation measurement as the response variable and the method of measurement as an explanatory variable: a model that ignores the effect of child (using the `lm` function), a marginal model estimated using generalized estimating equations that attempts to account for the lack of independence of observations from the same child (using the `geeglm` function), a fixed effects model with child used as a factor explanatory variable (using the `lm` function), and a mixed effects model with a random “main effect” for child (using the `lmer` function). For each model report the parameter estimates and standard errors using either the `summary` or `lincon` function so that I can verify that you estimated the model correctly. When using `geeglm` be sure to use the *sorted* data frame (see the code above) and an exchangeable correlation structure.

**Solution:** First here is the model with no effect for child.

```
m1 <- lm(y ~ meth, data = ox)
trtools::lincon(m1)
```



	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	75.658	0.9002	73.888	77.42864	84.046	352	4.824e-235
methpulse	-2.477	1.2731	-4.981	0.02639	-1.946	352	5.245e-02

Here is the marginal model (note that the normal distribution and identity link function are the defaults).

```
library(geepack)
m2 <- geeglm(y ~ meth, family = gaussian(link = identity),
  id = item, corstr = "exchangeable", data = ox)
trtools::lincon(m2)
```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	75.650	1.5175	72.665	78.634	49.85	352	1.381e-161
methpulse	-2.477	0.6385	-3.733	-1.222	-3.88	352	1.248e-04

Here is the fixed effects model with a fixed effect for child. Note that I have trimmed the output a bit since there are 60 indicator variables for 61 levels of item.

```
m3 <- lm(y ~ meth + factor(item), data = ox)
trtools::lincon(m3)[1:5,]
```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	75.839	2.0553	71.7937	79.884	36.900	292	5.841e-112
methpulse	-2.477	0.5307	-3.5218	-1.433	-4.668	292	4.632e-06
factor(item)2	-6.667	2.8823	-12.3393	-0.994	-2.313	292	2.142e-02
factor(item)3	5.350	2.8823	-0.3227	11.023	1.856	292	6.444e-02
factor(item)4	-10.500	2.8823	-16.1727	-4.827	-3.643	292	3.188e-04

And here is the mixed effects model.

```
m4 <- lmer(y ~ meth + (1 | item), data = ox)
trtools::lincon(m4)
```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	75.650	1.4546	72.799	78.501	52.009	Inf	0.000e+00
methpulse	-2.477	0.5307	-3.517	-1.437	-4.669	Inf	3.033e-06

- For each of the four models you estimated above, estimate (a) the expected blood oxygen saturation observed for each method, and (b) the difference in the expected blood oxygen saturation between the two methods. This can be done using `lincon`, `contrast`, or with the **emmeans** package, but be sure that you provide estimates and standard errors.

**Solution:** I will use the **emmeans** package. I am going to use a shortcut that provides the answers to both (a) and (b).

```
emmeans(m1, pairwise ~ meth, infer = TRUE) # ignore
```

```
$emmeans
meth emmean SE df lower.CL upper.CL t.ratio p.value
CO      75.7 0.9 352      73.9      77.4  84.050 <.0001
pulse   73.2 0.9 352      71.4      75.0  81.290 <.0001
```

Confidence level used: 0.95

```
$contrasts
contrast estimate SE df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 1.27 352  -0.0264      4.98   1.946  0.0524
```

Confidence level used: 0.95

```
emmeans(m2, pairwise ~ meth, infer = TRUE) # marginal
```

```
$emmeans
meth emmean SE df asymp.LCL asymp.UCL z.ratio p.value
CO      75.7 1.52 Inf      72.7      78.6 49.850 <.0001
pulse   73.2 1.39 Inf      70.5      75.9 52.740 <.0001
```

Covariance estimate used: vbeta  
Confidence level used: 0.95

```
$contrasts
contrast estimate SE df asymp.LCL asymp.UCL z.ratio p.value
CO - pulse      2.48 0.638 Inf      1.23      3.73 3.880 0.0001
```

Confidence level used: 0.95

```
emmeans(m3, pairwise ~ meth, infer = TRUE) # fixed effect
```

```
$emmeans
meth emmean SE df lower.CL upper.CL t.ratio p.value
CO      75.7 0.378 292      74.9      76.4 200.080 <.0001
pulse   73.2 0.378 292      72.4      73.9 193.530 <.0001
```

Results are averaged over the levels of: item  
Confidence level used: 0.95

```
$contrasts
contrast estimate SE df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 0.531 292      1.43      3.52 4.668 <.0001
```

Results are averaged over the levels of: item  
Confidence level used: 0.95

```
emmeans(m4, pairwise ~ meth, infer = TRUE) # random effect
```

```
$emmeans
meth emmean SE df lower.CL upper.CL t.ratio p.value
CO      75.7 1.45 64.2      72.7      78.6 52.010 <.0001
pulse   73.2 1.45 64.2      70.3      76.1 50.310 <.0001
```

Degrees-of-freedom method: kenward-roger  
Confidence level used: 0.95

```
$contrasts
contrast estimate SE df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 0.531 292      1.43      3.52 4.669 <.0001
```

Degrees-of-freedom method: kenward-roger  
Confidence level used: 0.95

Note that for the fixed effects model the estimates of the expected blood oxygen saturation for each method are *averaged* over the children. They can vary from child to child as seen below which shows the estimates for the first three children.

```
emmeans(m3, ~ meth | item, at = list(item = 1:3))
```

```

item = 1:
  meth  emmean    SE  df lower.CL upper.CL
CO      75.8  2.06 292    71.8    79.9
pulse   73.4  2.06 292    69.3    77.4

```

```

item = 2:
  meth  emmean    SE  df lower.CL upper.CL
CO      69.2  2.06 292    65.1    73.2
pulse   66.7  2.06 292    62.6    70.7

```

```

item = 3:
  meth  emmean    SE  df lower.CL upper.CL
CO      81.2  2.06 292    77.1    85.2
pulse   78.7  2.06 292    74.7    82.8

```

Confidence level used: 0.95

The estimate of the differences in these expectations, however, do not depend on the child.

```

emmeans(m3, pairwise ~ meth | item,
  at = list(item = 1:3), infer = TRUE)$contrast

```

```

item = 1:
  contrast    estimate    SE  df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 0.531 292    1.43    3.52   4.668 <.0001

```

```

item = 2:
  contrast    estimate    SE  df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 0.531 292    1.43    3.52   4.668 <.0001

```

```

item = 3:
  contrast    estimate    SE  df lower.CL upper.CL t.ratio p.value
CO - pulse      2.48 0.531 292    1.43    3.52   4.668 <.0001

```

Confidence level used: 0.95

3. Compare the parameter estimates and standard errors you obtained for the quantities you estimated in the previous problem across the four models. Discuss briefly what did and did not change (much).

**Solution:** The estimates of  $\beta_0$  and  $\beta_1$  are very similar across the four method, but there are some differences in the standard errors. This is also true when estimating the expected response for each method and the difference in the expected response between these two methods. The standard errors from the first approach which ignores the effect of child, which tend to be smaller for the expected response but larger for the difference, should not be trusted. The standard errors for the difference in the expected response between the two methods are more similar across the other approaches. The standard errors for the expected response are relatively similar between the marginal and mixed effects approaches. The standard error for the expected response for the fixed effects approach depend heavily on if we are making inferences for a given child or the average response across all children in the study.

4. Estimate a linear mixed effects model that specifies an *interaction* between the child and the method of measurement as a second random effect. Report the parameter estimates and standard errors, and estimate the quantities that you estimated in the second problem now using this model.

**Solution:** Here is the estimated mixed effects model with the interaction.

```

m5 <- lmer(y ~ meth + (meth | item), data = ox)
trtools::lincon(m5)

```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	75.649	1.5285	72.653	78.645	49.493	Inf	0.0000000
methpulse	-2.476	0.6382	-3.727	-1.225	-3.879	Inf	0.0001047

Note that (meth | item) is the same thing as (meth | item). The random “main effect” is implied. Here are the estimated quantities.

```
emmmeans(m5, pairwise ~ meth, infer = TRUE)
```

```
$emmmeans
  meth emmean  SE df lower.CL upper.CL t.ratio p.value
CO      75.7 1.53 60     72.6     78.7  49.490 <.0001
pulse   73.2 1.40 60     70.4     76.0  52.300 <.0001
```

Degrees-of-freedom method: kenward-roger  
Confidence level used: 0.95

```
$contrasts
  contrast estimate SE df lower.CL upper.CL t.ratio p.value
CO - pulse    2.48 0.638 59.5      1.2     3.75  3.878 0.0003
```

Degrees-of-freedom method: kenward-roger  
Confidence level used: 0.95

## Weight Gain in Rats Exposed to Thiouracil and Thyroxin


In the first homework assignment you encountered data from a study on the effects of thiouracil and thyroxin on growth of rats. Here are the first few rows of the data frame.

```
library(ALA)
head(rat,10)
```

	id	treatment	week	weight
1	1	control	0	57
28	1	control	1	86
55	1	control	2	114
82	1	control	3	139
109	1	control	4	172
2	2	control	0	60
29	2	control	1	93
56	2	control	2	123
83	2	control	3	146
110	2	control	4	177

Note that the variable id identifies the rat and that there are multiple observations for each rat. The plot below shows the raw data with line segments joining observations from the same rat.

```
library(ALA)
library(ggplot2)
p <- ggplot(rat, aes(x = week, y = weight)) + theme_minimal() +
  geom_line(aes(group = id), alpha = 0.25) + geom_point(alpha = 0.5) +
  labs(x = "Week", y = "Weight (g)") + facet_wrap(~treatment)
plot(p)
```



In the first homework assignment we ignored the fact that there are multiple observations from each rat, but this may be important to take into account. We can think of the rat as being a factor with as many levels as there are rats (27), and that this categorical variable may have an “effect” in the sense that rats can vary in terms of their size, but also may grow at different rates (similar to what we observed for the Sitka spruce tree experiment from lecture). Here you will model the data using a random effects approach to account for the effect of rat.

1. In the seventh part of the *Weight Gain in Rats Exposed to Thiouracil and Thyroxine* problem on the first homework assignment you were asked to estimate a model with the model formula `weight ~ treatment:week`. This model can be written as

$$E(Y_i) = \begin{cases} \beta_0 + \beta_1 w_i, & \text{if the treatment for the } i\text{-th observation is control,} \\ \beta_0 + \beta_2 w_i, & \text{if the treatment for the } i\text{-th observation is thiouracil,} \\ \beta_0 + \beta_3 w_i, & \text{if the treatment for the } i\text{-th observation is thyroxine,} \end{cases}$$

where  $Y_i$  and  $w_i$  are the  $i$ -th observations of weight and week, respectively. Note that this model allows for differences in the growth rate depending on the treatment, but assumes that the expected weight at zero weeks is the same across treatments since at that point no drug had been administered. Now consider a mixed effects model with a random “main effect” for rat which can be written as

$$E(Y_{ij}) = \begin{cases} \beta_0 + \beta_1 w_{ij} + \delta_i, & \text{if the treatment for the } i\text{-th rat is control,} \\ \beta_0 + \beta_2 w_{ij} + \delta_i, & \text{if the treatment for the } i\text{-th rat is thiouracil,} \\ \beta_0 + \beta_3 w_{ij} + \delta_i, & \text{if the treatment for the } i\text{-th rat is thyroxine,} \end{cases}$$

where now  $Y_{ij}$  and  $w_{ij}$  are the weight and week for the  $j$ -th observation of the  $i$ -th rat. Here  $\delta_i$  is a random parameter that represents the random effect of rat. Another model that allows for differences in the growth rate of rats by including an “interaction” between rat and week could be written as

$$E(Y_{ij}) = \begin{cases} \beta_0 + \beta_1 w_{ij} + \delta_i + \gamma_i w_{ij}, & \text{if the treatment for the } i\text{-th rat is control,} \\ \beta_0 + \beta_2 w_{ij} + \delta_i + \gamma_i w_{ij}, & \text{if the treatment for the } i\text{-th rat is thiouracil,} \\ \beta_0 + \beta_3 w_{ij} + \delta_i + \gamma_i w_{ij}, & \text{if the treatment for the } i\text{-th rat is thyroxine,} \end{cases}$$

where  $\gamma_i$  is a second random parameter for the interaction between rat and week. Estimate these two mixed effects models, and show the results using either `summary` or `lincon` so that I can verify that you estimated the model correctly. Note that you will still want to include the same fixed effects structure

as the original model with the term `treatment:week`, but extend the model to include one or two random effects.

**Solution:** The first model can be estimated as follows.

```
m1 <- lmer(weight ~ treatment:week + (1 | id), data = rat)
trtools::lincon(m1)
```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	54.46	1.9310	50.67	58.24	28.20	Inf	5.521e-175
treatmentcontrol:week	26.32	0.6653	25.02	27.62	39.56	Inf	0.000e+00
treatmentthiouracil:week	17.38	0.6653	16.07	18.68	26.12	Inf	2.276e-150
treatmentthyroxin:week	26.90	0.7849	25.37	28.44	34.28	Inf	1.823e-257

The second model can be estimated as follows.

```
m2 <- lmer(weight ~ treatment:week + (week | id), data = rat)
trtools::lincon(m2)
```

	estimate	se	lower	upper	tvalue	df	pvalue
(Intercept)	54.46	1.317	51.88	57.04	41.35	Inf	0.000e+00
treatmentcontrol:week	26.25	1.236	23.82	28.67	21.24	Inf	4.005e-100
treatmentthiouracil:week	17.53	1.236	15.11	19.95	14.19	Inf	1.128e-45
treatmentthyroxin:week	26.79	1.471	23.91	29.68	18.21	Inf	4.413e-74

- Using the model you estimated with two random effects (i.e., the model with a random effect for the interaction between rat and week), estimate the expected weight of a rat at the fourth week for each treatment condition. Also estimate the three pairwise differences among these three expected weights. You can use `lincon`, `contrast`, or the **emmeans** package for this. Be sure to include standard errors and confidence intervals.

**Solution:** Here is how we can do this with one statement using the **emmeans** package.

```
emmeans(m2, pairwise ~ treatment, at = list(week = 4), infer = TRUE, adjust = "none")
```

```
$emmeans
  treatment emmean  SE  df lower.CL upper.CL t.ratio p.value
control      159 5.04 24.4      149      170  31.660 <.0001
thiouracil    125 5.04 24.4      114      135  24.730 <.0001
thyroxin      162 6.04 24.3      149      174  26.780 <.0001
```

Degrees-of-freedom method: kenward-roger

Confidence level used: 0.95

```
$contrasts
  contrast      estimate  SE  df lower.CL upper.CL t.ratio p.value
control - thiouracil    34.87 7.18 24      20.0      49.7   4.854 0.0001
control - thyroxin     -2.18 7.92 24     -18.5      14.2  -0.276 0.7852
thiouracil - thyroxin  -37.06 7.92 24     -53.4     -20.7  -4.680 0.0001
```

Degrees-of-freedom method: kenward-roger

Confidence level used: 0.95

An interesting thing about this model is that the tests of these pairwise differences are equivalent to testing the differences in the rates of change in weight between the treatment conditions. This is because the model is constrained so that the expected weight is the same for each treatment at week zero. For example, the difference in the expected weights between the control condition and the thiouracil treatment condition is

$$\beta_0 + \beta_1 4 - (\beta_0 + \beta_2 4) = (\beta_1 - \beta_2) 4.$$

Note that we omit the terms  $\delta_i$  and  $\gamma_i w_{ij}$  because  $\delta_i$  and  $\gamma_i$  both are modeled as having mean values of zero and so do not affect the expected weight for a given treatment condition and week if we are not conditioning on a particular rat. We can see the equivalence of these tests by testing the differences between the rates of change.

```
pairs(pairs(emmeans(m2, ~week|treatment, at = list(week = c(4,3))),
  by = NULL, adjust = "none"))
```

contrast	estimate	SE	df	t.ratio	p.value
(week4 - week3 control) - (week4 - week3 thiouracil)	8.718	1.80	24	4.854	0.0001
(week4 - week3 control) - (week4 - week3 thyroxin)	-0.546	1.98	24	-0.276	0.7852
(week4 - week3 thiouracil) - (week4 - week3 thyroxin)	-9.264	1.98	24	-4.680	0.0001

Degrees-of-freedom method: kenward-roger

## Seatbelt Use and Degree of Injury in Auto Accidents

The following data shows the number of auto accidents resulting in various levels of injury by victim gender, accident location, and whether not the victim was wearing a seat belt.<sup>4</sup>

```
seatbelt <- data.frame(gender = rep(c("female","male"), each = 4),
  location = rep(rep(c("urban","rural"), each = 2), 2), seatbelt = rep(c("no","yes"), 4),
  I1 = c(7287, 11587, 3246, 6134, 10381, 10969, 6123, 6693),
  I2 = c(175, 126, 73, 94, 136, 83, 141, 74),
  I3 = c(720, 577, 710, 564, 566, 259, 710, 353),
  I4 = c(91, 48, 159, 92, 96, 37, 188, 74),
  I5 = c(10, 8, 31, 17, 14, 1, 45, 12))
seatbelt
```

	gender	location	seatbelt	I1	I2	I3	I4	I5
1	female	urban	no	7287	175	720	91	10
2	female	urban	yes	11587	126	577	48	8
3	female	rural	no	3246	73	710	159	31
4	female	rural	yes	6134	94	564	92	17
5	male	urban	no	10381	136	566	96	14
6	male	urban	yes	10969	83	259	37	1
7	male	rural	no	6123	141	710	188	45
8	male	rural	yes	6693	74	353	74	12

The variables I1 through I5 denote the level of injury: *non injured* (I1), *injured but not transported by emergency medical services* (I2), *injured and transported by emergency medical services but not hospitalized* (I3), *injured and hospitalized but did not die* (I4), and *injured and died* (I5). The injury levels are ordered by severity. Note that the data are in aggregated form showing the number of accidents at each injury level for each combination of gender, location, and seat belt use. Use these data in the following.

1. Use the `vglm` function from the **VGAM** package to estimate a *sequential* regression model similar to a model we considered for the `pneumo` data in lecture on April 14. Use the injury level as the response variable, and use location and seat belt use as the explanatory variables, but not gender.<sup>5</sup> For this model do not specify any interactions. Note that the model should be specified with the injury levels ordered as listed above from I1 to I5. Use a logit link function. Do not use the `parallel` option for this model. Report the parameter estimates using `summary` so that I can verify that you specified the model correctly.

<sup>4</sup>Source: Agresti, A. (2013). *Categorical data analysis* (3rd ed). Hoboken, NJ: Wiley.

<sup>5</sup>Including location may be important as the severity of accidents may vary between urban and rural locations. Gender could be included as a (poor) proxy for weight, which could be related to severity of accidents — particular for those not wearing seat belts. In that case it would be important to include an interaction between gender (or, ideally, weight) and seat belt use. But for the purpose of this exercise we are going to keep the model very simple.

**Solution:** Here is the estimated model.

```
library(VGAM)
m <- vglm(cbind(I1,I2,I3,I4,I5) ~ seatbelt + location,
  family = cratio(link = "logitlink"), data = seatbelt)
summary(m)
```

Call:

```
vglm(formula = cbind(I1, I2, I3, I4, I5) ~ seatbelt + location,
  family = cratio(link = "logitlink"), data = seatbelt)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept):1	-1.5326	0.0219	-70.02	< 2e-16 ***
(Intercept):2	2.1098	0.0626	33.72	< 2e-16 ***
(Intercept):3	-1.2063	0.0525	-22.99	< 2e-16 ***
(Intercept):4	-1.5210	0.1230	-12.36	< 2e-16 ***
seatbeltyes:1	-0.7487	0.0273	-27.41	< 2e-16 ***
seatbeltyes:2	-0.1606	0.0735	-2.19	0.029 *
seatbeltyes:3	-0.3514	0.0776	-4.53	6.0e-06 ***
seatbeltyes:4	-0.2177	0.2054	-1.06	0.289
locationurban:1	-0.7303	0.0268	-27.25	< 2e-16 ***
locationurban:2	-0.5052	0.0728	-6.94	3.9e-12 ***
locationurban:3	-0.6096	0.0762	-8.00	1.3e-15 ***
locationurban:4	-0.5251	0.2131	-2.46	0.014 *

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Names of linear predictors: logitlink(P[Y>1|Y>=1]), logitlink(P[Y>2|Y>=2]),  
logitlink(P[Y>3|Y>=3]), logitlink(P[Y>4|Y>=4])

Residual deviance: 458.2 on 20 degrees of freedom

Log-likelihood: -329.6 on 20 degrees of freedom

Number of Fisher scoring iterations: 4

No Hauck-Donner effect found in any of the estimates

- The sequential model you specified in the previous problem is effectively equivalent to four logistic regression models: one for whether an individual was injured (I2, I3, I4, or I5) or not injured (I1), another for whether an injured person was not transported (I2) or was transported (I3, I4, or I5), another for whether a transported person was hospitalized (I4 or I5) or was not hospitalized (I3), and finally a model for whether those that required hospitalization died (I5) or did not die (I4). The model you estimated in the previous problem can be used to compute odds ratios for the effect of using a seat belt (or not) with respect to each of these four situations. Report those four odds ratios and write a sentence for each odds ratio that explains how it is interpreted in terms of the statistical effect of wearing a seat belt (or not).

**Solution:** We can estimate the same model as four logistic regression models. Note that this is not necessary to obtain the odds ratios as they can also be found by exponentiating the parameter estimates from the model estimated earlier. But it is illuminating to see how we can estimate the model above as four separate logistic regression models.



```
m1 <- glm(cbind(I2+I3+I4+I5,I1) ~ seatbelt + location,
  family = binomial, data = seatbelt)
summary(m1)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.5326	0.02189	-70.02	0.000e+00
seatbeltyes	-0.7487	0.02731	-27.41	2.005e-165
locationurban	-0.7303	0.02680	-27.25	1.591e-163

```
m2 <- glm(cbind(I3+I4+I5,I2) ~ seatbelt + location,
  family = binomial, data = seatbelt)
summary(m2)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	2.1098	0.06227	33.880	1.298e-251
seatbeltyes	-0.1606	0.07341	-2.188	2.869e-02
locationurban	-0.5052	0.07277	-6.942	3.865e-12

```
m3 <- glm(cbind(I4+I5,I3) ~ seatbelt + location,
  family = binomial, data = seatbelt)
summary(m3)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.2063	0.05203	-23.184	6.648e-119
seatbeltyes	-0.3514	0.07776	-4.519	6.221e-06
locationurban	-0.6096	0.07627	-7.993	1.317e-15

```
m4 <- glm(cbind(I5,I4) ~ seatbelt + location,
  family = binomial, data = seatbelt)
summary(m4)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-1.5210	0.1222	-12.447	1.452e-35
seatbeltyes	-0.2177	0.2060	-1.057	2.906e-01
locationurban	-0.5251	0.2133	-2.462	1.383e-02

Notice that the parameter estimates and the standard errors from these four models are the same as those obtained above when using the `vglm` function. Here are the estimated odds ratios for seatbelt wearing. Note that since there is no interaction specified we do not need to condition on location.

```
pairs(emmeans(m1, ~seatbelt, type = "response"), infer = TRUE)
```

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
no / yes	2.11	0.0578	Inf	2	2.23	1	27.411	<.0001

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m2, ~seatbelt, type = "response"), infer = TRUE)
```

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
no / yes	1.17	0.0862	Inf	1.02	1.36	1	2.188	0.0287

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m3, ~seatbelt, type = "response"), infer = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
no / yes	1.42	0.11	Inf	1.22	1.65	1	4.519	<.0001

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m4, ~seatbelt, type = "response"), infer = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
no / yes	1.24	0.256	Inf	0.83	1.86	1	1.057	0.2906

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

We can interpret these odds ratios as follows.

*The odds of an injury are about 2.11 times higher (110% higher) when not wearing a seatbelt.*

*The odds that an injured person would need to be transported is about 1.17 times higher (17% higher) when not wearing a seatbelt.*

\*The odds that a transported person would require hospitalization is about 1.42 times higher (42% higher) when not wearing a seatbelt.

*The odds that a hospitalized person would die is about 1.24 times higher (24% higher) when not wearing a seatbelt.*

Note that all but the last odds ratios are statistically significant at convention significance levels. We can also flip the odds ratios.

```
pairs(emmeans(m1, ~seatbelt, type = "response"), infer = TRUE, reverse = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
yes / no	0.473	0.0129	Inf	0.448	0.499	1	-27.411	<.0001

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m2, ~seatbelt, type = "response"), infer = TRUE, reverse = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
yes / no	0.852	0.0625	Inf	0.738	0.983	1	-2.188	0.0287

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m3, ~seatbelt, type = "response"), infer = TRUE, reverse = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
----------	------------	----	----	----------	----------	------	---------	---------

yes / no	0.704	0.0547	Inf	0.604	0.82	1	-4.519	<.0001
----------	-------	--------	-----	-------	------	---	--------	--------

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m4, ~seatbelt, type = "response"), infer = TRUE, reverse = TRUE)
```

contrast	odds.ratio	SE	df	asyp.LCL	asyp.UCL	null	z.ratio	p.value
yes / no	0.804	0.166	Inf	0.537	1.2	1	-1.057	0.2906

Results are averaged over the levels of: location

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

We can interpret these odds ratios as follows.

*The odds of an injury is reduced by a factor of about 0.47 (53% less) when wearing a seat belt.*

*The odds that an injured person would need to be transported is reduced by a factor of about 0.85 (15% less) when wearing a seat belt.*

*The odds that a transported person would need to be hospitalized decreases by a factor of about 0.7 (30% less) when wearing a seat belt.*

*The odds that a hospitalized person would die decreases by a factor of about 0.8 (20% less) when wearing a seat belt.*