

Friday, Feb 2

Marginal Means

A *marginal mean* is effectively an average of expected responses. The **emmeans** package is particularly useful for making inferences about marginal means.

```
library(trtools)
library(emmeans)
```

Warning: The **emmeans** package contains a function called **contrast** which is not the same as the function of the same name in the **trtools** package, resulting in a namespace conflict if both packages are loaded. If you have both packages loaded in a given session, use **trtools::contrast** and **emmeans::contrast** to refer to a given function.

Example: Consider again the data from the platyfish study.

```
m <- lm(Percentage ~ Pair, data = Sleuth3::case0602)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|-----------|
| (Intercept) | 56.406 | 3.864 | 14.5965 | 5.208e-24 |
| PairPair2 | 4.479 | 5.657 | 0.7919 | 4.308e-01 |
| PairPair3 | 6.023 | 5.384 | 1.1187 | 2.667e-01 |
| PairPair4 | 10.594 | 5.657 | 1.8727 | 6.485e-02 |
| PairPair5 | 7.805 | 6.441 | 1.2118 | 2.292e-01 |
| PairPair6 | 6.929 | 5.657 | 1.2250 | 2.243e-01 |

We see that there are indicator variables for male pairs 2-6. The model can be written as

$$E(Y_i) = \begin{cases} \beta_0, & \text{if the } i\text{-th observation was from the first male pair,} \\ \beta_0 + \beta_1, & \text{if the } i\text{-th observation was from the second male pair,} \\ \beta_0 + \beta_2, & \text{if the } i\text{-th observation was from the third male pair,} \\ \beta_0 + \beta_3, & \text{if the } i\text{-th observation was from the fourth male pair,} \\ \beta_0 + \beta_4, & \text{if the } i\text{-th observation was from the fifth male pair,} \\ \beta_0 + \beta_5, & \text{if the } i\text{-th observation was from the sixth male pair.} \end{cases}$$

We can use **contrast** to estimate the expected response for each pair.

```
contrast(m, a = list(Pair = paste("Pair", 1:6, sep = "")),
  cnames = paste("Pair", 1:6, sep = ""))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-------|----------|-------|-------|-------|--------|----|-----------|
| Pair1 | 56.41 | 3.864 | 48.71 | 64.10 | 14.60 | 78 | 5.208e-24 |
| Pair2 | 60.89 | 4.131 | 52.66 | 69.11 | 14.74 | 78 | 2.990e-24 |
| Pair3 | 62.43 | 3.749 | 54.97 | 69.89 | 16.65 | 78 | 2.114e-27 |
| Pair4 | 67.00 | 4.131 | 58.78 | 75.22 | 16.22 | 78 | 1.052e-26 |
| Pair5 | 64.21 | 5.152 | 53.95 | 74.47 | 12.46 | 78 | 3.039e-20 |
| Pair6 | 63.34 | 4.131 | 55.11 | 71.56 | 15.33 | 78 | 3.006e-25 |

Note how I used a shortcut to specify the pairs.

```
paste("Pair", 1:6, sep = "")
```

```
[1] "Pair1" "Pair2" "Pair3" "Pair4" "Pair5" "Pair6"
```

This can also be done using the `emmeans` function from the package `emmeans`.

```
library(emmeans)
emmeans(m, ~ Pair)
```

| Pair | emmean | SE | df | lower.CL | upper.CL |
|-------|--------|------|----|----------|----------|
| Pair1 | 56.4 | 3.86 | 78 | 48.7 | 64.1 |
| Pair2 | 60.9 | 4.13 | 78 | 52.7 | 69.1 |
| Pair3 | 62.4 | 3.75 | 78 | 55.0 | 69.9 |
| Pair4 | 67.0 | 4.13 | 78 | 58.8 | 75.2 |
| Pair5 | 64.2 | 5.15 | 78 | 54.0 | 74.5 |
| Pair6 | 63.3 | 4.13 | 78 | 55.1 | 71.6 |

Confidence level used: 0.95

Denote the six expected responses (one for each pair) as

$$\begin{aligned}\mu_1 &= \beta_0, \\ \mu_2 &= \beta_0 + \beta_1, \\ \mu_3 &= \beta_0 + \beta_2, \\ \mu_4 &= \beta_0 + \beta_3, \\ \mu_5 &= \beta_0 + \beta_4, \\ \mu_6 &= \beta_0 + \beta_5.\end{aligned}$$

One marginal mean would be the average expected response across the pairs. This could be written as

$$\mu = \frac{\mu_1 + \mu_2 + \mu_3 + \mu_4 + \mu_5 + \mu_6}{6} = \beta_0 + \frac{1}{6}\beta_1 + \frac{1}{6}\beta_2 + \frac{1}{6}\beta_3 + \frac{1}{6}\beta_4 + \frac{1}{6}\beta_5.$$

We can estimate this quantity with `lincon`.

```
lincon(m, a = c(1,1/6,1/6,1/6,1/6,1/6))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|---------------------------|----------|-------|-------|-------|--------|----|-----------|
| (1,1/6,1/6,1/6,1/6,1/6),0 | 62.38 | 1.722 | 58.95 | 65.81 | 36.23 | 78 | 1.501e-50 |

We can also use `emmeans`.

```
emmeans(m, ~ 1)
```

| 1 | emmean | SE | df | lower.CL | upper.CL |
|---------|--------|------|----|----------|----------|
| overall | 62.4 | 1.72 | 78 | 59 | 65.8 |

Results are averaged over the levels of: Pair

Confidence level used: 0.95

Note that we can use the confidence interval to test the null hypothesis that $\mu = 50$. For a test statistic and p-value for this test we could write this as

$$\mu = 50 \Leftrightarrow \beta_0 + \frac{1}{6}\beta_1 + \frac{1}{6}\beta_2 + \frac{1}{6}\beta_3 + \frac{1}{6}\beta_4 + \frac{1}{6}\beta_5 = 50 \Leftrightarrow \beta_0 + \frac{1}{6}\beta_1 + \frac{1}{6}\beta_2 + \frac{1}{6}\beta_3 + \frac{1}{6}\beta_4 + \frac{1}{6}\beta_5 - 50 = 0.$$

Here is how we can do that with `lincon`.

```
lincon(m, a = c(1,1/6,1/6,1/6,1/6,1/6), b = -50)
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-----------------------------|----------|-------|-------|-------|--------|----|-----------|
| (1,1/6,1/6,1/6,1/6,1/6),-50 | 12.38 | 1.722 | 8.95 | 15.81 | 7.189 | 78 | 3.439e-10 |

Here is how we do it with `emmeans`.

```
emmeans(m, ~ 1, offset = -50, infer = TRUE)
```

| | 1 | emmean | SE | df | lower.CL | upper.CL | t.ratio | p.value |
|---------|---|--------|------|----|----------|----------|---------|---------|
| overall | | 12.4 | 1.72 | 78 | 8.95 | 15.8 | 7.189 | <.0001 |

Results are averaged over the levels of: Pair

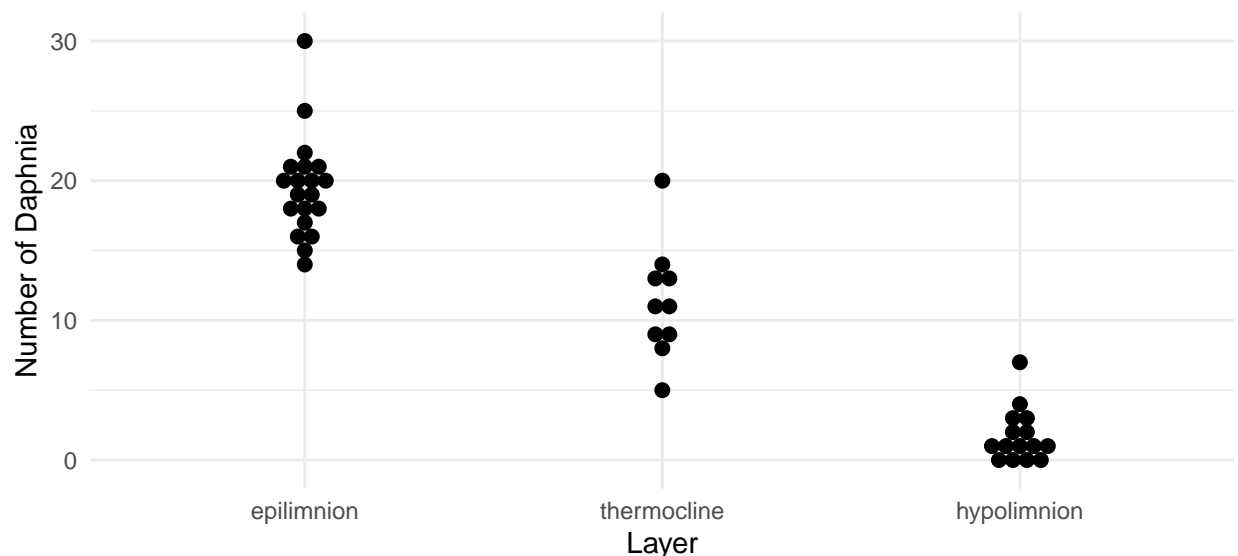
Confidence level used: 0.95

By *not* listing an explanatory variable on the right-hand side of `~`, we are asking that `emmeans` average over that explanatory variable. Also note that the argument `infer = TRUE` makes the `emmeans` function provide both confidence intervals as well as tests.

Note: If we just want to know whether or not we would reject the null hypothesis that $\mu = 50$ we can also just look at the confidence interval for μ .

Example: Consider the following data from a survey of water fleas.

```
library(ggplot2)
p <- ggplot(daphniastrat, aes(x = layer, y = count)) +
  geom_dotplot(binaxis = "y", stackdir = "center") +
  labs(x = "Layer", y = "Number of Daphnia") + theme_minimal()
plot(p)
```



We might model these data using the following linear model.

```
m <- lm(count ~ layer, data = daphniastrat)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|------------------|----------|------------|---------|-----------|
| (Intercept) | 19.50 | 0.7271 | 26.820 | 4.727e-28 |
| layerthermocline | -8.20 | 1.2593 | -6.512 | 7.293e-08 |
| layerhypolimnion | -17.77 | 1.1106 | -15.997 | 1.784e-19 |

So our model can be written as

$$E(Y_i) = \begin{cases} \beta_0, & \text{if the } i\text{-th observation is from the epilimnion layer,} \\ \beta_0 + \beta_1, & \text{if the } i\text{-th observation is from the thermocline layer,} \\ \beta_0 + \beta_2, & \text{if the } i\text{-th observation is from the hypolimnion layer.} \end{cases}$$

Let μ_e , μ_t , and μ_h denote the expected number of daphnia per liter for the epilimnion, thermocline, and hypolimnion layers, respectively (i.e., the density of daphnia in each layer). So

$$\mu_e = \beta_0, \mu_t = \beta_0 + \beta_1, \mu_h = \beta_0 + \beta_2.$$

It is known that the volumes of the epilimnion, thermocline, and hypolimnion layers are 100, 200, and 400 kL, respectively. The density for the entire lake is then

$$\mu = \frac{100}{700}\mu_e + \frac{200}{700}\mu_t + \frac{400}{700}\mu_h = \beta_0 + \frac{2}{7}\beta_1 + \frac{4}{7}\beta_2.$$

We can estimate this with `lincon` or `emmeans` using the `weights` option.

```
lincon(m, a = c(1, 2/7, 4/7))
```

```
      estimate      se lower upper tvalue df    pvalue
(1,2/7,4/7),0  7.005 0.572  5.85 8.159  12.25 42 1.907e-15
```

```
emmeans(m, ~ 1, weights = c(1/7, 2/7, 4/7))
```

```
  1      emmean    SE df lower.CL upper.CL
overall      7 0.572 42    5.85    8.16
```

Results are averaged over the levels of: layer

Confidence level used: 0.95

Note that when using `emmeans` it is important to put the weights in the correct order. We can verify the order using `level` (if the variable is a factor) or by using `weights = "slow.levels"`.

```
levels(daphniastrat$layer)
```

```
[1] "epilimnion" "thermocline" "hypolimnion"
```

```
emmeans(m, ~ 1, weights = "show.levels")
```

`emmeans` are obtained by averaging over these factor combinations

```
      layer
1 epilimnion
2 thermocline
3 hypolimnion
```

We can estimate the expected number of daphnia per liter for each layer.

```
emmeans(m, ~ layer)
```

```
      layer      emmean    SE df lower.CL upper.CL
epilimnion  19.50 0.727 42    18.033    20.97
thermocline  11.30 1.028 42     9.225    13.38
hypolimnion   1.73 0.840 42     0.039     3.43
```

Confidence level used: 0.95

```
trtools::contrast(m, a = list(layer = c("epilimnion", "thermocline", "hypolimnion")),
  cnames = c("epilimnion", "thermocline", "hypolimnion"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-------------|----------|--------|----------|--------|--------|----|-----------|
| epilimnion | 19.500 | 0.7271 | 18.03274 | 20.967 | 26.820 | 42 | 4.727e-28 |
| thermocline | 11.300 | 1.0282 | 9.22498 | 13.375 | 10.990 | 42 | 6.221e-14 |
| hypolimnion | 1.733 | 0.8395 | 0.03909 | 3.428 | 2.065 | 42 | 4.517e-02 |

We can also do inferences concerning the differences between pairs of layers.

```
pairs(emmeans(m, ~ layer), adjust = "none")
```

| contrast | estimate | SE | df | t.ratio | p.value |
|---------------------------|----------|------|----|---------|---------|
| epilimnion - thermocline | 8.20 | 1.26 | 42 | 6.512 | <.0001 |
| epilimnion - hypolimnion | 17.77 | 1.11 | 42 | 15.997 | <.0001 |
| thermocline - hypolimnion | 9.57 | 1.33 | 42 | 7.207 | <.0001 |

```
trtools::contrast(m,
  a = list(layer = c("epilimnion", "epilimnion", "thermocline")),
  b = list(layer = c("thermocline", "hypolimnion", "hypolimnion")),
  cnames = c("E-T", "E-H", "T-H"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-----|----------|-------|--------|-------|--------|----|-----------|
| E-T | 8.200 | 1.259 | 5.659 | 10.74 | 6.512 | 42 | 7.293e-08 |
| E-H | 17.767 | 1.111 | 15.525 | 20.01 | 15.997 | 42 | 1.784e-19 |
| T-H | 9.567 | 1.327 | 6.888 | 12.25 | 7.207 | 42 | 7.363e-09 |

The `adjust = "none"` option for `pairs` specifies that no adjustment be made to confidence intervals or tests for the family-wise Type I error rate.¹

Something to note when using the `weights` argument with the `emmeans` function is that the weights that are used must sum to one, and if they do not they will be normalized so that they do. For example, the following provide the same result.

```
emmeans(m, ~ 1, weights = c(1/7, 2/7, 4/7))
```

| 1 | emmean | SE | df | lower.CL | upper.CL |
|---------|--------|-------|----|----------|----------|
| overall | 7 | 0.572 | 42 | 5.85 | 8.16 |

Results are averaged over the levels of: layer

Confidence level used: 0.95

```
emmeans(m, ~ 1, weights = c(1, 2, 4)) # original weights multiplied by 7
```

| 1 | emmean | SE | df | lower.CL | upper.CL |
|---------|--------|-------|----|----------|----------|
| overall | 7 | 0.572 | 42 | 5.85 | 8.16 |

Results are averaged over the levels of: layer

Confidence level used: 0.95

If you want to use weights that do not sum to one, you can use the `contrast` function from the `emmeans` package (different from the function of the same name from `trtools`).

¹The family-wise Type I error rate is the probability of making *at least one* Type I error. If it is desired that the family-wise Type I error rate be no greater than α (default is 0.05), then some adjustment can be made. This adjustment is seen in the p-values and confidence intervals. The most general method is to use `adjust = "mvt"`. Some special cases are more widely known such as Tukey (`adjust = "tukey"`) and Bonferroni (`adjust = "bonferroni"`), but the adjustment based on the multivariate *t*-distribution (`adjust = "mvt"`) is the most general and accurate. Note that an adjustment will produce “simultaneous” confidence intervals. A method of producing simultaneous confidence intervals has the property that the probability that *all* of the confidence intervals will contain the quantities being estimated is equal to the specified confidence level (95% by default). The multivariate *t*-distribution adjustment is perhaps not as well known, so a reference that you can cite is Edwards, D. & Berry, J. T. (1987). The efficiency of simulation-based multiple comparisons. *Biometrics*, 43(4), 913–928.

```
emmeans(m, ~1, weights = c(1/7, 2/7, 4/7))
```

| | 1 | emmean | SE | df | lower.CL | upper.CL |
|---------|---|--------|-------|----|----------|----------|
| overall | | 7 | 0.572 | 42 | 5.85 | 8.16 |

Results are averaged over the levels of: layer

Confidence level used: 0.95

```
emmeans::contrast(emmeans(m, ~layer), method = list(layer = c(1/7, 2/7, 4/7)), infer = TRUE)
```

| | contrast | estimate | SE | df | lower.CL | upper.CL | t.ratio | p.value |
|-------|----------|----------|-------|----|----------|----------|---------|---------|
| layer | | 7 | 0.572 | 42 | 5.85 | 8.16 | 12.245 | <.0001 |

Confidence level used: 0.95

But suppose we wanted to estimate the *number* of daphnia in the lake (τ). It can be shown that this is

$$\tau = 700000\mu = 700000 \left(\frac{1}{7}\mu_e + \frac{2}{7}\mu_t + \frac{4}{7}\mu_h \right) = 100000\mu_e + 200000\mu_t + 400000\mu_h.$$

Note that there are 700kL in the lake, which is 700000L (which is the scale used for the observations). This can be estimated as follows.

```
emmeans::contrast(emmeans(m, ~layer),
  method = list(layer = 700000 * c(1/7, 2/7, 4/7)), infer = TRUE)
```

| | contrast | estimate | SE | df | lower.CL | upper.CL | t.ratio | p.value |
|-------|----------|----------|-------|----|----------|----------|---------|---------|
| layer | | 4903333 | 4e+05 | 42 | 4095230 | 5711437 | 12.245 | <.0001 |

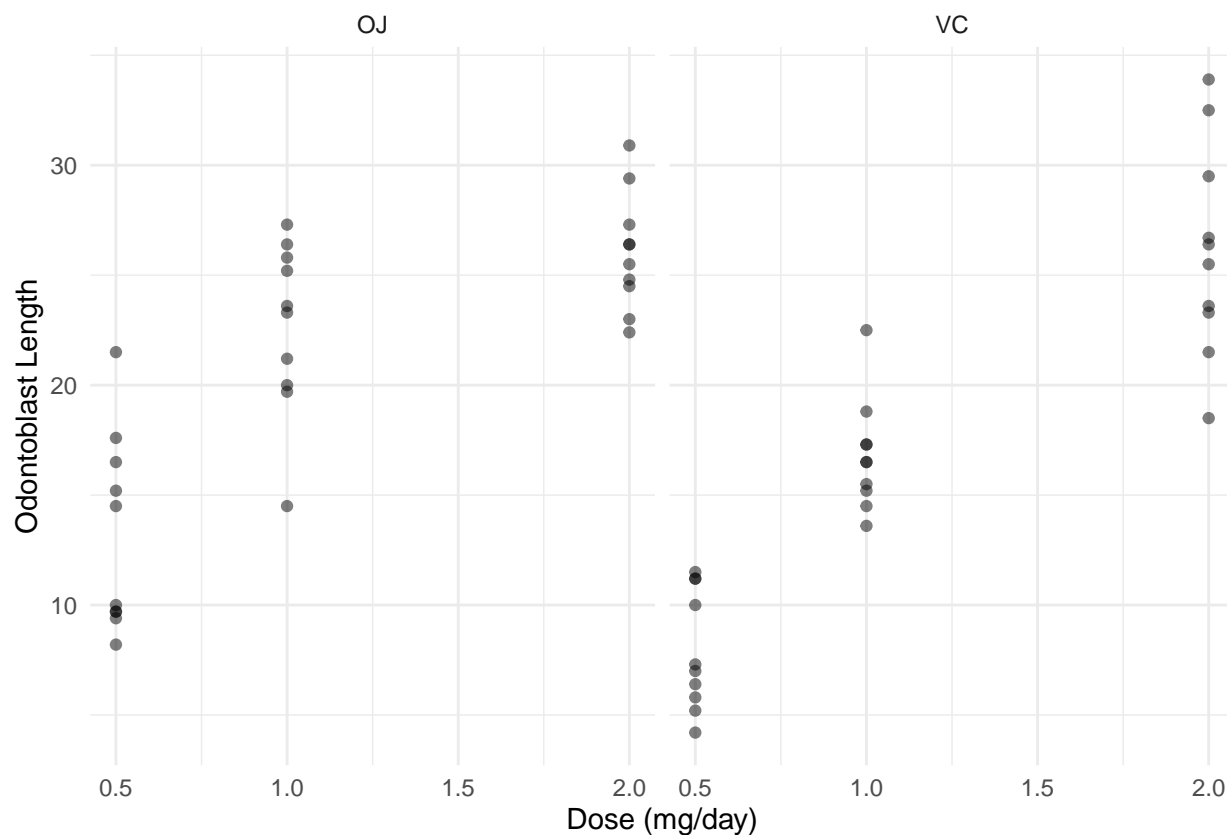
Confidence level used: 0.95

Another approach is to use `lincon` but your weights/coefficients will be different since they are applied to β_0 , β_1 , and β_2 .

Marginal Means and “Main Effects”

Consider data from a randomized experiment with guinea pigs administered one of three doses of vitamin C (0.5, 1, or 2 mg/day) via one of two supplement methods: orange juice (OJ) or ascorbic acid (VC).

```
p <- ggplot(ToothGrowth, aes(x = dose, y = len)) +
  geom_point(alpha = 0.5) + facet_wrap(~supp) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length") + theme_minimal()
plot(p)
```



Here we are going to model dose as a categorical variable so we need to coerce it to a factor. Perhaps the safest approach is to create a new variable.

```
ToothGrowth$dosef <- factor(ToothGrowth$dose)
```

Note: Whether a variable is a numeric, a factor, or something else can be seen use `str` (for “structure”).

```
str(ToothGrowth)
```

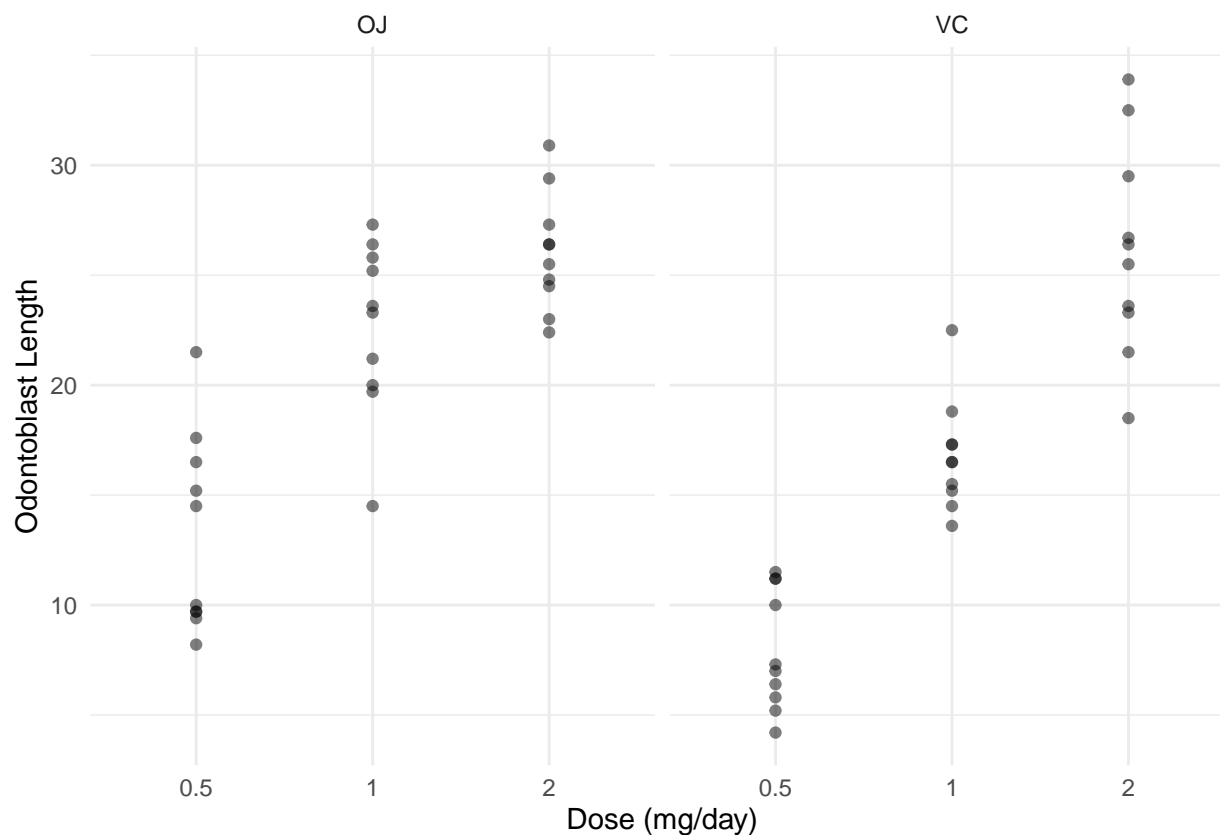
```
'data.frame':  60 obs. of  4 variables:
 $ len  : num  4.2 11.5 7.3 5.8 6.4 10 11.2 11.2 5.2 7 ...
 $ supp : Factor w/ 2 levels "OJ","VC": 2 2 2 2 2 2 2 2 2 2 ...
 $ dose : num  0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 ...
 $ dosef: Factor w/ 3 levels "0.5","1","2": 1 1 1 1 1 1 1 1 1 1 ...
```

Notice that `ggplot` responds differently.

```
summary(ToothGrowth)
```

| len | supp | dose | dosef |
|--------------|-------|--------------|--------|
| Min. : 4.2 | OJ:30 | Min. :0.50 | 0.5:20 |
| 1st Qu.:13.1 | VC:30 | 1st Qu.:0.50 | 1 :20 |
| Median :19.2 | | Median :1.00 | 2 :20 |
| Mean :18.8 | | Mean :1.17 | |
| 3rd Qu.:25.3 | | 3rd Qu.:2.00 | |
| Max. :33.9 | | Max. :2.00 | |

```
p <- ggplot(ToothGrowth, aes(x = dosef, y = len)) +
  geom_point(alpha = 0.5) + facet_wrap(~supp) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length") + theme_minimal()
plot(p)
```



Now consider the following linear model.

```
m <- lm(len ~ dosef + supp + dosef:supp, data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|---------------|----------|------------|---------|-----------|
| (Intercept) | 13.23 | 1.148 | 11.5208 | 3.603e-16 |
| dosef1 | 9.47 | 1.624 | 5.8312 | 3.176e-07 |
| dosef2 | 12.83 | 1.624 | 7.9002 | 1.430e-10 |
| suppVC | -5.25 | 1.624 | -3.2327 | 2.092e-03 |
| dosef1:suppVC | -0.68 | 2.297 | -0.2961 | 7.683e-01 |
| dosef2:suppVC | 5.33 | 2.297 | 2.3207 | 2.411e-02 |

The model is

$$E(Y_i) = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \beta_3 x_{i3} + \beta_4 x_{i4} + \beta_5 x_{i5},$$

where

$$x_{i1} = \begin{cases} 1, & \text{if dose is 1 mg/day,} \\ 0, & \text{otherwise,} \end{cases}$$

$$x_{i2} = \begin{cases} 1, & \text{if dose is 2 mg/day,} \\ 0, & \text{otherwise,} \end{cases}$$

$$x_{i3} = \begin{cases} 1, & \text{if supplement type is VC} \\ 0, & \text{otherwise,} \end{cases}$$

$$x_{i4} = x_{i1}x_{i3} = \begin{cases} 1, & \text{if dose is 1 mg/day and supplement type is VC,} \\ 0, & \text{otherwise,} \end{cases}$$

$$x_{i5} = x_{i2}x_{i3} = \begin{cases} 1, & \text{if dose is 2 mg/day and supplement type is VC,} \\ 0, & \text{otherwise.} \end{cases}$$

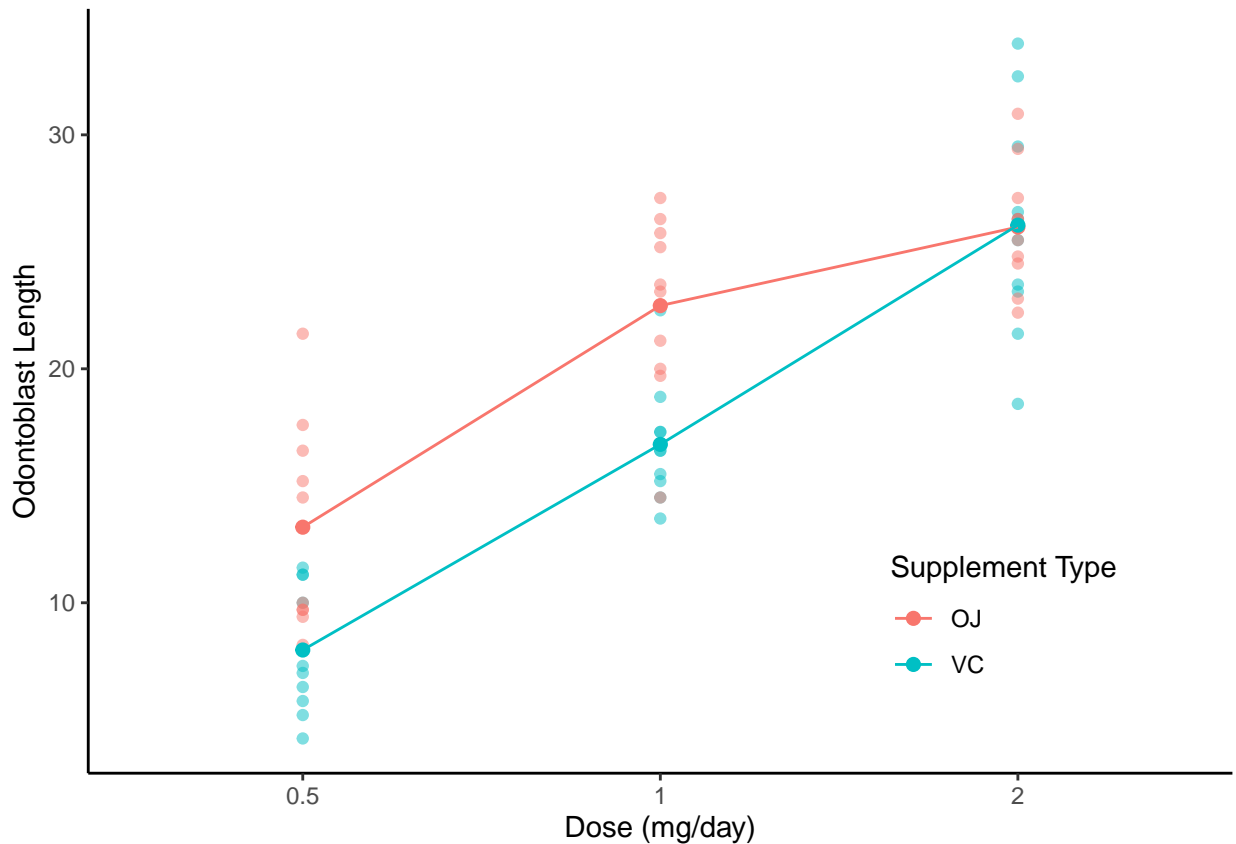
We can write this model case-wise.

$$E(Y_i) = \begin{cases} \beta_0, & \text{if dose is 0.5 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_1, & \text{if dose is 1 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_2, & \text{if dose is 2 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_3, & \text{if dose is 0.5 mg/day and supplement type is VC,} \\ \beta_0 + \beta_1 + \beta_3 + \beta_4, & \text{if dose is 1 mg/day and supplement type is VC,} \\ \beta_0 + \beta_2 + \beta_3 + \beta_5, & \text{if dose is 2 mg/day and supplement type is VC.} \end{cases}$$

And we can visualize it.

```
d <- expand.grid(dosef = levels(ToothGrowth$dosef), supp = levels(ToothGrowth$supp))
d$yhat <- predict(m, newdata = d)

p <- ggplot(ToothGrowth, aes(x = dosef, y = len, color = supp)) +
  geom_point(alpha = 0.5) + theme_classic() +
  theme(legend.position = c(0.8, 0.2)) +
  geom_point(aes(y = yhat), size = 2, data = d) +
  geom_line(aes(y = yhat, group = supp), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length", color = "Supplement Type")
plot(p)
```



We might want to compare the two supplement types at each dose level. These are sometimes called simple effects. We can do this using `contrast`.

```
trtools::contrast(m,
  a = list(supp = "OJ", dosef = c("0.5", "1", "2")),
  b = list(supp = "VC", dosef = c("0.5", "1", "2")),
  cnames = c("OJ-VC @ 0.5 mg/day", "OJ-VC @ 1.0 mg/day", "OJ-VC @ 2.0 mg/day"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|--------------------|----------|-------|--------|-------|----------|----|-----------|
| OJ-VC @ 0.5 mg/day | 5.25 | 1.624 | 1.994 | 8.506 | 3.23273 | 54 | 0.0020925 |
| OJ-VC @ 1.0 mg/day | 5.93 | 1.624 | 2.674 | 9.186 | 3.65144 | 54 | 0.0005897 |
| OJ-VC @ 2.0 mg/day | -0.08 | 1.624 | -3.336 | 3.176 | -0.04926 | 54 | 0.9608934 |

We can also do this use `emmeans`. First note that we can use `emmeans` to also estimate the expected response for each level of `supp` within each level of `dosef`.

```
emmeans(m, ~ supp | dosef)
```

```
dosef = 0.5:
  supp emmean   SE df lower.CL upper.CL
OJ    13.23 1.15 54    10.93    15.5
VC     7.98 1.15 54     5.68    10.3
```

```
dosef = 1:
  supp emmean   SE df lower.CL upper.CL
OJ    22.70 1.15 54    20.40    25.0
VC    16.77 1.15 54    14.47    19.1
```

```
dosef = 2:
  supp emmean   SE df lower.CL upper.CL
OJ    26.06 1.15 54    23.76    28.4
VC    26.14 1.15 54    23.84    28.4
```

Confidence level used: 0.95

We can use `pairs` to make inferences about the differences between levels of `supp` within each level of `dosef`.

```
pairs(emmeans(m, ~ supp | dosef), adjust = "none", infer = TRUE)
```

```
dosef = 0.5:
  contrast estimate   SE df lower.CL upper.CL t.ratio p.value
OJ - VC      5.25 1.62 54     1.99     8.51   3.233 0.0021
```

```
dosef = 1:
  contrast estimate   SE df lower.CL upper.CL t.ratio p.value
OJ - VC      5.93 1.62 54     2.67     9.19   3.651 0.0006
```

```
dosef = 2:
  contrast estimate   SE df lower.CL upper.CL t.ratio p.value
OJ - VC     -0.08 1.62 54    -3.34     3.18  -0.049 0.9609
```

Confidence level used: 0.95

The “main effect” of supplement method concerns μ_{OJ} and μ_{VC} , defined as

$$\mu_{OJ} = \frac{\mu_{OJ,0.5} + \mu_{OJ,1.0} + \mu_{OJ,2.0}}{3}, \quad \mu_{VC} = \frac{\mu_{VC,0.5} + \mu_{VC,1.0} + \mu_{VC,2.0}}{3}.$$

So μ_{OJ} and μ_{VC} are the *marginal means* for `supp`, averaging over levels of `dose`.

```
emmeans(m, ~ supp)
```

| supp | emmean | SE | df | lower.CL | upper.CL |
|------|--------|-------|----|----------|----------|
| OJ | 20.7 | 0.663 | 54 | 19.3 | 22.0 |
| VC | 17.0 | 0.663 | 54 | 15.6 | 18.3 |

Results are averaged over the levels of: dosef

Confidence level used: 0.95

Inference for the main effect $\mu_{OJ} - \mu_{VC}$ can then be obtained as follows.

```
pairs(emmeans(m, ~ supp), infer = TRUE)
```

| contrast | estimate | SE | df | lower.CL | upper.CL | t.ratio | p.value |
|----------|----------|-------|----|----------|----------|---------|---------|
| OJ - VC | 3.7 | 0.938 | 54 | 1.82 | 5.58 | 3.946 | 0.0002 |

Results are averaged over the levels of: dosef

Confidence level used: 0.95

This main effect is simply a single linear function of $\beta_0, \beta_1, \dots, \beta_5$. From the case-wise representation of the model,

$$E(Y_i) = \begin{cases} \beta_0, & \text{if dose is 0.5 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_1, & \text{if dose is 1 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_2, & \text{if dose is 2 mg/day and supplement type is OJ,} \\ \beta_0 + \beta_3, & \text{if dose is 0.5 mg/day and supplement type is VC,} \\ \beta_0 + \beta_1 + \beta_3 + \beta_4, & \text{if dose is 1 mg/day and supplement type is VC,} \\ \beta_0 + \beta_2 + \beta_3 + \beta_5, & \text{if dose is 2 mg/day and supplement type is VC,} \end{cases}$$

we have that $\mu_{OJ,0.5} = \beta_0$, $\mu_{OJ,1.0} = \beta_0 + \beta_1$, $\mu_{OJ,2.0} = \beta_0 + \beta_2$, $\mu_{VC,0.5} = \beta_0 + \beta_3$, $\mu_{VC,1.0} = \beta_0 + \beta_1 + \beta_3 + \beta_4$, $\mu_{VC,2.0} = \beta_0 + \beta_2 + \beta_3 + \beta_5$. So, we can write this as

$$\mu_{OJ} - \mu_{VC} = \frac{\mu_{OJ,0.5} + \mu_{OJ,1.0} + \mu_{OJ,2.0}}{3} - \frac{\mu_{VC,0.5} + \mu_{VC,1.0} + \mu_{VC,2.0}}{3} = -\beta_3 - \frac{1}{3}\beta_4 - \frac{1}{3}\beta_5.$$

```
lincon(m, a = c(0,0,0,-1,-1/3,-1/3))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|------------------------|----------|--------|-------|-------|--------|----|-----------|
| (0,0,0,-1,-1/3,-1/3),0 | 3.7 | 0.9376 | 1.82 | 5.58 | 3.946 | 54 | 0.0002312 |

Clearly using `emmeans` is much less tedious!

The “main effect” of dose concerns differences among the marginal means of dose defined as $\mu_{0.5}$, μ_1 and μ_2 where

$$\mu_{0.5} = \frac{\mu_{OJ,0.5} + \mu_{VC,0.5}}{2}, \quad \mu_1 = \frac{\mu_{OJ,1} + \mu_{VC,1}}{2}, \quad \mu_2 = \frac{\mu_{OJ,2} + \mu_{VC,2}}{2}.$$

```
emmeans(m, ~ dosef)
```

| dosef | emmean | SE | df | lower.CL | upper.CL |
|-------|--------|-------|----|----------|----------|
| 0.5 | 10.6 | 0.812 | 54 | 8.98 | 12.2 |
| 1 | 19.7 | 0.812 | 54 | 18.11 | 21.4 |
| 2 | 26.1 | 0.812 | 54 | 24.47 | 27.7 |

Results are averaged over the levels of: supp

Confidence level used: 0.95

```
pairs(emmeans(m, ~ dosef), adjust = "none")
```

| contrast | estimate | SE | df | t.ratio | p.value |
|-------------------|----------|------|----|---------|---------|
| dosef0.5 - dosef1 | -9.13 | 1.15 | 54 | -7.951 | <.0001 |
| dosef0.5 - dosef2 | -15.49 | 1.15 | 54 | -13.493 | <.0001 |
| dosef1 - dosef2 | -6.37 | 1.15 | 54 | -5.543 | <.0001 |

Results are averaged over the levels of: supp

```
pairs(emmeans(m, ~ dosef), reverse = TRUE, adjust = "none")
```

| contrast | estimate | SE | df | t.ratio | p.value |
|-------------------|----------|------|----|---------|---------|
| dosef1 - dosef0.5 | 9.13 | 1.15 | 54 | 7.951 | <.0001 |
| dosef2 - dosef0.5 | 15.49 | 1.15 | 54 | 13.493 | <.0001 |
| dosef2 - dosef1 | 6.37 | 1.15 | 54 | 5.543 | <.0001 |

Results are averaged over the levels of: supp

In ANOVA tables the test of the “main effect” is the (joint) null hypothesis that all pairwise differences are zero. For the variable dose the null hypothesis is $\mu_{0.5} = \mu_1 = \mu_2$. This can be done using the `test` function.

```
test(pairs(emmeans(m, ~ dosef)), joint = TRUE)
```

| df1 | df2 | F.ratio | p.value | note |
|-----|-----|---------|---------|------|
| 2 | 54 | 92.000 | <.0001 | d |

d: df1 reduced due to linear dependence

This is the traditional main effect that is sometimes reported in an “ANOVA table” such as that produced by `Anova` from the `car` package.

```
library(car)
m <- lm(len ~ dosef + supp + dosef:supp, data = ToothGrowth,
  contrast = list(dosef = contr.sum, supp = contr.sum))
Anova(m, type = 3)
```

Anova Table (Type III tests)

Response: len

| | Sum Sq | Df | F value | Pr(>F) |
|-------------|--------|----|---------|-------------|
| (Intercept) | 21236 | 1 | 1610.39 | < 2e-16 *** |
| dosef | 2426 | 2 | 92.00 | < 2e-16 *** |
| supp | 205 | 1 | 15.57 | 0.00023 *** |
| dosef:supp | 108 | 2 | 4.11 | 0.02186 * |
| Residuals | 712 | 54 | | |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

The option `contrast = list(dosef = contr.sum, supp = contr.sum)` is necessary here for the `Anova` function to do the correct calculations.²

The test of the main effect of supplement method was given by

```
pairs(emmeans(m, ~ supp), infer = TRUE)
```

| contrast | estimate | SE | df | lower.CL | upper.CL | t.ratio | p.value |
|----------|----------|-------|----|----------|----------|---------|---------|
| OJ - VC | 3.7 | 0.938 | 54 | 1.82 | 5.58 | 3.946 | 0.0002 |

²I am demonstrating here what is sometimes called inferences based on Type III sums of squares. Another common approach is to use what is called Type II sums of squares. This can be done with the `Anova` function with `type = 2`. For inferences based on Type II sums of squares with the functions from the `emmeans` package an extra step is needed (email me for an example if you really want to know how to do it).

Results are averaged over the levels of: dosef
Confidence level used: 0.95

We do not need a joint test here since there are only two marginal means, but here it is anyway.

```
test(pairs(emmeans(m, ~ supp)), joint = TRUE)
```

```
df1 df2 F.ratio p.value
1 54 15.572 0.0002
```

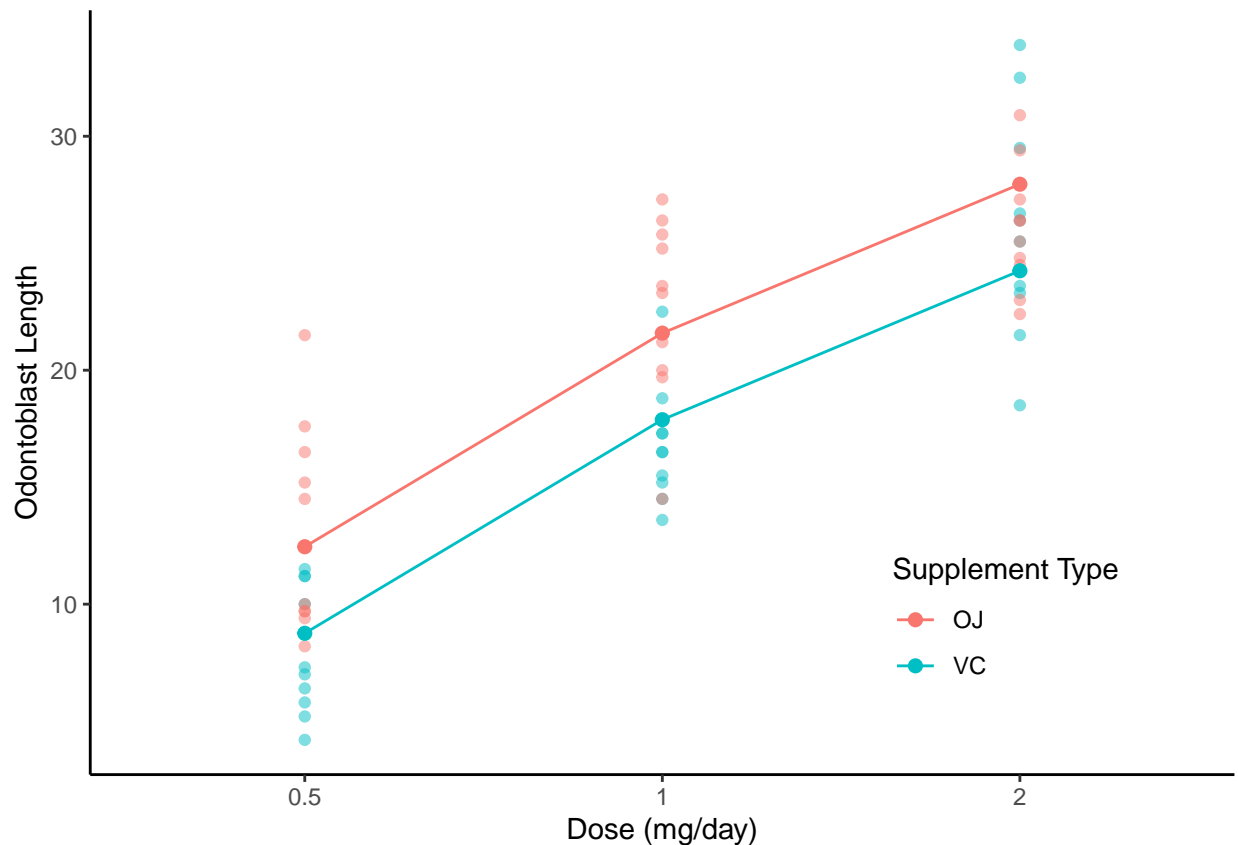
Now consider a model without the interaction.

```
m <- lm(len ~ dosef + supp, data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|-----------|
| (Intercept) | 12.45 | 0.9883 | 12.603 | 5.490e-18 |
| dosef1 | 9.13 | 1.2104 | 7.543 | 4.383e-10 |
| dosef2 | 15.50 | 1.2104 | 12.802 | 2.852e-18 |
| suppVC | -3.70 | 0.9883 | -3.744 | 4.293e-04 |

```
d <- expand.grid(dosef = levels(ToothGrowth$dosef), supp = levels(ToothGrowth$supp))
d$yhat <- predict(m, newdata = d)
```

```
p <- ggplot(ToothGrowth, aes(x = dosef, y = len, color = supp)) +
  geom_point(alpha = 0.5) + theme_classic() +
  theme(legend.position = c(0.8,0.2)) +
  geom_point(aes(y = yhat), size = 2, data = d) +
  geom_line(aes(y = yhat, group = supp), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length", color = "Supplement Type")
plot(p)
```



Note that without the interaction the “simple effects” and “main effects” are equivalent. Here are the simple and main effects for supplement.

```
pairs(emmeans(m, ~ supp | dosef)) # simple
```

dosef = 0.5:

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|-------|----|---------|---------|
| OJ - VC | 3.7 | 0.988 | 56 | 3.744 | 0.0004 |

dosef = 1:

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|-------|----|---------|---------|
| OJ - VC | 3.7 | 0.988 | 56 | 3.744 | 0.0004 |

dosef = 2:

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|-------|----|---------|---------|
| OJ - VC | 3.7 | 0.988 | 56 | 3.744 | 0.0004 |

```
pairs(emmeans(m, ~ supp)) # main
```

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|-------|----|---------|---------|
| OJ - VC | 3.7 | 0.988 | 56 | 3.744 | 0.0004 |

Results are averaged over the levels of: dosef

And here are the simple and main effects for dose.

```
pairs(emmeans(m, ~ dosef | supp), adjust = "none") # simple
```

supp = OJ:

| contrast | estimate | SE | df | t.ratio | p.value |
|----------|----------|----|----|---------|---------|
|----------|----------|----|----|---------|---------|

| | | | | | |
|-------------------|--------|------|----|---------|--------|
| dosef0.5 - dosef1 | -9.13 | 1.21 | 56 | -7.543 | <.0001 |
| dosef0.5 - dosef2 | -15.49 | 1.21 | 56 | -12.802 | <.0001 |
| dosef1 - dosef2 | -6.37 | 1.21 | 56 | -5.259 | <.0001 |

supp = VC:

| contrast | estimate | SE | df | t.ratio | p.value |
|-------------------|----------|------|----|---------|---------|
| dosef0.5 - dosef1 | -9.13 | 1.21 | 56 | -7.543 | <.0001 |
| dosef0.5 - dosef2 | -15.49 | 1.21 | 56 | -12.802 | <.0001 |
| dosef1 - dosef2 | -6.37 | 1.21 | 56 | -5.259 | <.0001 |

```
pairs(emmeans(m, ~ dosef), adjust = "none") # main
```

| contrast | estimate | SE | df | t.ratio | p.value |
|-------------------|----------|------|----|---------|---------|
| dosef0.5 - dosef1 | -9.13 | 1.21 | 56 | -7.543 | <.0001 |
| dosef0.5 - dosef2 | -15.49 | 1.21 | 56 | -12.802 | <.0001 |
| dosef1 - dosef2 | -6.37 | 1.21 | 56 | -5.259 | <.0001 |

Results are averaged over the levels of: supp

The joint test of the overall main effect for dose can be obtained as follows.

```
test(pairs(emmeans(m, ~ dosef)), joint = TRUE)
```

| df1 | df2 | F.ratio | p.value | note |
|-----|-----|---------|---------|------|
| 2 | 56 | 82.810 | <.0001 | d |

d: df1 reduced due to linear dependence