

Monday, Feb 7

Modeling Nonlinearity

Four approaches to modeling a nonlinear relationship between the expected response and a quantitative explanatory variable.

1. polynomials
2. transformations
3. splines
4. nonlinear regression

The first three can be done with *linear models*.

Polynomial Regression

If we have a single explanatory variable x_i , then a *polynomial regression model* of degree k is

$$E(Y_i) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \cdots + \beta_k x_i^k.$$

Note that this *is* a linear *model* since we can write it as

$$E(Y_i) = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik},$$

where $x_{i1} = x_i, x_{i2} = x_i^2, \dots, x_{ik} = x_i^k$.

Example: Consider again the `ToothGrowth` data but with dose treated as a quantitative explanatory variable, and ignoring supplement type for now. Note the use of the “inhibit” function `I` here.

```
m <- lm(len ~ dose + I(dose^2), data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|----------|------------|---------|-----------|
| (Intercept) | -2.49 | 3.178 | -0.7836 | 4.365e-01 |
| dose | 30.15 | 6.147 | 4.9052 | 8.148e-06 |
| I(dose^2) | -7.93 | 2.366 | -3.3514 | 1.432e-03 |


This model is

$$E(L_i) = \beta_0 + \beta_1 d_i + \beta_2 d_i^2,$$

where d_i is dose.

```
d <- expand.grid(dose = seq(0.25, 2.25, length = 100))
d$yhat <- predict(m, newdata = d)

p <- ggplot(ToothGrowth, aes(x = dose, y = len)) +
  geom_point(aes(color = supp), alpha = 0.5) +
  geom_line(aes(y = yhat), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") +
  theme_minimal() + theme(legend.position = c(0.8,0.2))
plot(p)
```



Note that the following are equivalent ways to specify this model.

```
# create a new variable for squared dose
ToothGrowth$dose2 <- ToothGrowth$dose^2
m <- lm(len ~ dose + dose2, data = ToothGrowth)

# specify squared dose in the model formula using the "inhibit" function
m <- lm(len ~ dose + I(dose^2), data = ToothGrowth)

# use the poly function to create the extra term
m <- lm(len ~ poly(dose, degree = 2), data = ToothGrowth)
```

I recommend not using the first approach of creating a new variable only because it is easier to have the transformation “built in” to the model when applying other functions to the model object like `predict` or `contrast`.

Note: Using `poly` without the option `raw = TRUE` will produce “orthogonal polynomials” which is a re-parameterization of the model. This approach is sometimes recommended due to numerical instability of “raw” polynomials, but in many cases this is not an issue. But the `poly` function is sometimes convenient, especially for polynomials of higher degree.

Clearly in such a model the rate of change in expected length is *not* necessarily constant.

```
library(trtools)
contrast(m, a = list(dose = 1), b = list(dose = 0.5)) # 0.5 to 1

estimate    se lower upper tvalue df    pvalue
9.13 1.341 6.444 11.82 6.806 57 6.697e-09

contrast(m, a = list(dose = 1.5), b = list(dose = 1)) # 1 to 1.5

estimate    se lower upper tvalue df    pvalue
5.165 0.4472 4.27 6.06 11.55 57 1.47e-16
```

This can also be seen mathematically by writing the model as

$$E(L_i) = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 = \beta_0 + \underbrace{(\beta_1 + \beta_2 x_i)}_{\delta_i} x_i = \beta_0 + \delta_i x_i,$$

so that the rate of change in length per unit increase in dose *depends on dose* (if $\beta_2 \neq 0$). In a sense, dose is “interacting with itself” — i.e., the “effect” of a one unit increase in dose depends on the dose.

We can have the polynomial depend on (i.e, interact with) supplement type.

```
m <- lm(len ~ dose + I(dose^2) + supp + dose:supp + I(dose^2):supp, data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|------------------|----------|------------|---------|-----------|
| (Intercept) | -1.433 | 3.847 | -0.3726 | 7.109e-01 |
| dose | 34.520 | 7.442 | 4.6384 | 2.272e-05 |
| I(dose^2) | -10.387 | 2.864 | -3.6260 | 6.383e-04 |
| suppVC | -2.113 | 5.440 | -0.3885 | 6.992e-01 |
| dose:suppVC | -8.730 | 10.525 | -0.8295 | 4.105e-01 |
| I(dose^2):suppVC | 4.913 | 4.051 | 1.2129 | 2.305e-01 |

Note that we could also have written

```
m <- lm(len ~ poly(dose, 2)*supp, data = ToothGrowth)
```

In a model formula argument, $a*b$ expands to $a + b + a:b$.

This model can be written as

$$E(L_i) = \begin{cases} \beta_0 + \beta_1 d_i + \beta_2 d_i^2, & \text{if supplement type is OJ,} \\ \beta_0 + \beta_3 + (\beta_1 + \beta_4) d_i + (\beta_2 + \beta_5) d_i^2, & \text{if supplement type is VC,} \end{cases}$$


where d_i is dose, or alternatively as

$$E(L_i) = \begin{cases} \beta_0 + \beta_1 d_i + \beta_2 d_i^2, & \text{if supplement type is OJ,} \\ \gamma_0 + \gamma_1 d_i + \gamma_2 d_i^2, & \text{if supplement type is VC,} \end{cases}$$

where $\gamma_0 = \beta_0 + \beta_3$, $\gamma_1 = \beta_1 + \beta_4$, and $\gamma_2 = \beta_2 + \beta_5$. There is a distinct polynomial of degree two for each supplement type.

```
d <- expand.grid(supp = c("OJ", "VC"), dose = seq(0.25, 2.25, length = 100))
d$yhat <- predict(m, newdata = d)

p <- ggplot(ToothGrowth, aes(x = dose, y = len, color = supp)) +
  geom_point(alpha = 0.5) + geom_line(aes(y = yhat), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") + theme_minimal() +
  theme(legend.position = c(0.8, 0.2))
plot(p)
```



Polynomials are, in principle, quite general. But in many cases we would like to have a monotonic relationship, and/or have a model exhibit an asymptote. Finally, the parameters of a polynomial model are not easily to interpret.

Logarithmic Transformations

Applying a *logarithmic transformation* to an explanatory variable may be useful for explanatory variables that tend to have “diminishing returns” with respect to the expected response.


Example: Consider a linear model for expected length but now with $\log_2(\text{dose})$ as the explanatory variable.

```
m <- lm(len ~ log2(dose) + supp + log2(dose):supp, data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------------|----------|------------|---------|-----------|
| (Intercept) | 20.663 | 0.6791 | 30.425 | 1.629e-36 |
| log2(dose) | 6.415 | 0.8318 | 7.712 | 2.303e-10 |
| suppVC | -3.700 | 0.9605 | -3.852 | 3.033e-04 |
| log2(dose):suppVC | 2.665 | 1.1763 | 2.266 | 2.737e-02 |

```
d <- expand.grid(supp = c("OJ", "VC"), dose = seq(0.25, 2.25, length = 100))
d$yhat <- predict(m, newdata = d)
```

```
p <- ggplot(ToothGrowth, aes(x = dose, y = len, color = supp)) +
  geom_point(alpha = 0.5) + geom_line(aes(y = yhat), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") + theme_minimal() +
  theme(legend.position = c(0.8, 0.2))
plot(p)
```



In principle, any base of logarithm can be used and it will produce the same model, but with a change in the parameterization that will change the value of β_j in $\beta_j \log_b(x_{ij})$. This is because logs of different bases are *proportional*.¹


```
m <- lm(len ~ log(dose) + supp + log(dose):supp, data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|------------------|----------|------------|---------|-----------|
| (Intercept) | 20.663 | 0.6791 | 30.425 | 1.629e-36 |
| log(dose) | 9.255 | 1.2000 | 7.712 | 2.303e-10 |
| suppVC | -3.700 | 0.9605 | -3.852 | 3.033e-04 |
| log(dose):suppVC | 3.845 | 1.6971 | 2.266 | 2.737e-02 |

```
d <- expand.grid(supp = c("OJ", "VC"), dose = seq(0.25, 2.25, length = 100))
d$yhat <- predict(m, newdata = d)
```

```
p <- ggplot(ToothGrowth, aes(x = dose, y = len, color = supp)) +
  geom_point(alpha = 0.5) + geom_line(aes(y = yhat), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") + theme_minimal() +
  theme(legend.position = c(0.8, 0.2))
plot(p)
```

¹For example, $\log_e(x) = c \log_2(x)$ where $c = 1/\log_2(e)$. In general, $\log_b(x) = \log_a(x)/\log_a(b)$. In R, you can use the function `log(x, b)` to compute the base- b logarithm of x . The default base is $e \approx 2.718$ which produces the *natural* logarithm (sometimes written as \ln). R also understands `log2(x)` and `log10(x)` as the base-2 and base-10 logarithms, respectively.



Note that \log is the “natural” logarithm or base- e logarithm sometimes written as $\ln(x)$. For *interpretation* I find that \log_2 is particularly attractive because we can talk about the effect of *doubling* or *halving* the explanatory variable.

1. Note that

$$\log_2(2x) = \log_2(x) + \log_2(2) = \log_2(x) + 1.$$

So in a linear model with the term $\beta_j \log_2(x_{ij})$, β_j is the change in $E(Y_i)$ if we *double* x_{ij} . For example if we have

$$E(Y) = \beta_0 + \beta_1 \log_2(x),$$

and we *double* x to $2x$ we get

$$E(Y) = \beta_0 + \beta_1 \log_2(2x) = \beta_0 + \beta_1 \log_2(x) + \beta_1.$$

2. Note that

$$\log_2(x/2) = \log_2(x) - \log_2(2) = \log_2(x) - 1.$$

In a linear model with the term $\beta_j \log_2(x_{ij})$, $-\beta_j$ is the change in $E(Y_i)$ if we *halve* x_{ij} . For example if we have

$$E(Y) = \beta_0 + \beta_1 \log_2(x),$$

and we *halve* x to $x/2$ we get

$$E(Y) = \beta_0 + \beta_1 \log_2(x/2) = \beta_0 + \beta_1 \log_2(x) - \beta_1.$$

Note the effect of *doubling* of *halving* dose is constant in this model.

```
contrast(m,
  a = list(dose = 1, supp = c("OJ", "VC")),
  b = list(dose = 0.5, supp = c("OJ", "VC")),
  cnames = c("OJ", "VC"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|----|----------|--------|-------|--------|--------|----|-----------|
| OJ | 6.415 | 0.8318 | 4.749 | 8.081 | 7.712 | 56 | 2.303e-10 |
| VC | 9.080 | 0.8318 | 7.414 | 10.746 | 10.916 | 56 | 1.733e-15 |

```
contrast(m,
  a = list(dose = 2, supp = c("OJ", "VC")),
  b = list(dose = 1, supp = c("OJ", "VC")),
  cnames = c("OJ", "VC"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|----|----------|--------|-------|--------|--------|----|-----------|
| OJ | 6.415 | 0.8318 | 4.749 | 8.081 | 7.712 | 56 | 2.303e-10 |
| VC | 9.080 | 0.8318 | 7.414 | 10.746 | 10.916 | 56 | 1.733e-15 |

Similarly we could use \log_{10} if we wanted to consider the effect of increasing an explanatory variable by a factor of 10, or decreasing it by a factor of 0.1. But regardless of the base, the effect of increasing dose by a *factor* is the same regardless of the dose. For example, suppose we *triple* the dose.

```
contrast(m,
  a = list(dose = 1.5, supp = c("OJ", "VC")),
  b = list(dose = 0.5, supp = c("OJ", "VC")),
  cnames = c("OJ", "VC"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|----|----------|-------|--------|-------|--------|----|-----------|
| OJ | 10.17 | 1.318 | 7.527 | 12.81 | 7.712 | 56 | 2.303e-10 |
| VC | 14.39 | 1.318 | 11.750 | 17.03 | 10.916 | 56 | 1.733e-15 |

```
contrast(m,
  a = list(dose = 3, supp = c("OJ", "VC")),
  b = list(dose = 1, supp = c("OJ", "VC")),
  cnames = c("OJ", "VC"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|----|----------|-------|--------|-------|--------|----|-----------|
| OJ | 10.17 | 1.318 | 7.527 | 12.81 | 7.712 | 56 | 2.303e-10 |
| VC | 14.39 | 1.318 | 11.750 | 17.03 | 10.916 | 56 | 1.733e-15 |

What are the limitations of logarithmic transformations?

1. $\log_b(x)$ is only defined for $x > 0$. We cannot model the expected response for $x \leq 0$.
2. No asymptote. The expected response will continue to increase (or decrease) as x increases.

Exponential Transformations

Consider the linear model

$$E(Y) = \beta_0 + \beta_1 2^{-x/h}$$

where $h > 0$ is some specified value. This applies an *exponential* transformation to x with the following properties.


1. If $x = 0$ then $E(Y) = \beta_0 + \beta_1$, so the “ y -intercept” is $\beta_0 + \beta_1$.
2. As x increases then $E(Y)$ approaches an asymptote of β_0 . This is an *upper* (if $\beta_1 < 0$) or *lower* (if $\beta_1 > 0$) asymptote.²
3. The quantity h can be interpreted as the “half-life” of the curve in the sense that it is the value of x at which the expected responses is half way between the intercept at $\beta_0 + \beta_1$ and its upper/lower asymptote at β_0 because if $x = h$ then

$$E(Y) = \beta_0 + \beta_1 2^{-x/h} = \beta_0 + \beta_1/2,$$

and $\beta_0 + \beta_1/2$ is the midpoint between the “intercept” of $E(Y) = \beta_0 + \beta_1$ and the asymptote of β_0 .

²This can be seen by showing that $\lim_{x \rightarrow \infty} \beta_0 + \beta_1 2^{-x/h} = \beta_0$ if $h > 0$, and by showing that the first derivative of $\beta_0 + \beta_1 2^{-x/h}$ with respect to x is positive if $\beta_1 < 0$ and negative if $\beta_1 > 0$ if $h > 0$.

4. If $\beta_1 < 0$ then $-\beta_1$ is how much $E(Y)$ *increases* from $x = 0$ as it approaches the asymptote, while if $\beta_1 > 0$ then β_1 is how much $E(Y)$ *decreases* from when $x = 0$ as it approaches the asymptote.




Consider again a linear model for the ToothGrowth data with an exponential transformation of dose with $h = 1$.

```
m <- lm(len ~ I(2^(-dose/1)), data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|----------------|----------|------------|---------|-----------|
| (Intercept) | 35.14 | 1.555 | 22.60 | 1.942e-30 |
| I(2^(-dose/1)) | -33.61 | 2.988 | -11.25 | 3.303e-16 |

```
d <- expand.grid(supp = c("OJ", "VC"), dose = seq(0, 6, length = 100))
d$yhat <- predict(m, newdata = d)
```

```
p <- ggplot(ToothGrowth, aes(x = dose, y = len, color = supp)) +
  geom_point(alpha = 0.5) + xlim(0,6) +
  geom_line(aes(y = yhat), color = "black", data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") + theme_minimal() +
  theme(legend.position = c(0.8,0.2))
plot(p)
```

```
lincon(m, a = c(1,1)) # intercept
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|---------|----------|-------|--------|-------|--------|----|--------|
| (1,1),0 | 1.528 | 1.635 | -1.745 | 4.8 | 0.9345 | 58 | 0.3539 |


Now suppose that we let the effect of dose “interact” with supplement type.

```
m <- lm(len ~ I(2^(-dose/1)) + supp + supp:I(2^(-dose/1)), data = ToothGrowth)
summary(m)$coefficients
```

| | Estimate | Std. Error | t value | Pr(> t) |
|-----------------------|----------|------------|---------|-----------|
| (Intercept) | 34.054 | 1.925 | 17.6872 | 1.375e-24 |
| I(2^(-dose/1)) | -27.569 | 3.700 | -7.4519 | 6.199e-10 |
| suppVC | 2.169 | 2.723 | 0.7964 | 4.291e-01 |
| I(2^(-dose/1)):suppVC | -12.083 | 5.232 | -2.3094 | 2.463e-02 |

```
d <- expand.grid(supp = c("OJ", "VC"), dose = seq(0, 6, length = 100))
d$yhat <- predict(m, newdata = d)
```

```
p <- ggplot(ToothGrowth, aes(x = dose, y = len, color = supp)) +
  geom_point(alpha = 0.5) + xlim(0,6) +
  geom_line(aes(y = yhat), data = d) +
  labs(x = "Dose (mg/day)", y = "Odontoblast Length",
       color = "Supplement Type") + theme_minimal() +
  theme(legend.position = c(0.8,0.2))
plot(p)
```



This model can be written as

$$E(Y_i) = \beta_0 + \beta_1 2^{-x_i/h} + \beta_2 d_i + \beta_3 d_i 2^{-x_i/h},$$

where $d_i = 1$ if the supplement type is VC, and $d_i = 0$ otherwise, and $h = 1$. We can also write this model case-wise as

$$E(Y_i) = \begin{cases} \beta_0 + \beta_1 2^{-x_i/h}, & \text{if the supplement type of the } i\text{-th observation is OJ,} \\ \beta_0 + \beta_2 + (\beta_1 + \beta_3) 2^{-x_i/h}, & \text{if the supplement type of the } i\text{-th observation is VC,} \end{cases}$$

or

$$E(Y_i) = \begin{cases} \beta_0 + \beta_1 2^{-x_i/h}, & \text{if the supplement type of the } i\text{-th observation is OJ,} \\ \gamma_0 + \gamma_1 2^{-x_i/h}, & \text{if the supplement type of the } i\text{-th observation is VC,} \end{cases}$$

where $\gamma_0 = \beta_0 + \beta_2$ and $\gamma_1 = \beta_1 + \beta_3$. We can make inferences for the intercepts and asymptotes for *each* supplement type using `lincon`.

```
lincon(m, a = c(1,1,0,0)) # b0 + b1 = intercept for OJ
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-------------|----------|-------|-------|-------|--------|----|----------|
| (1,1,0,0),0 | 6.485 | 2.024 | 2.429 | 10.54 | 3.203 | 56 | 0.002243 |

```
lincon(m, a = c(1,1,1,1)) # g0 + g1 = b0 + b2 + b1 + b3 = intercept for VC
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-------------|----------|-------|--------|--------|--------|----|---------|
| (1,1,1,1),0 | -3.429 | 2.024 | -7.485 | 0.6261 | -1.694 | 56 | 0.09582 |

```
lincon(m, a = c(1,0,1,0)) # g0 = b0 + b2 = asymptote for VC
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|-------------|----------|-------|-------|-------|--------|----|----------|
| (1,0,1,0),0 | 36.22 | 1.925 | 32.37 | 40.08 | 18.81 | 56 | 7.07e-26 |

We can also obtain (approximate) inferences using `contrast`.

```
contrast(m, a = list(dose = 0, supp = c("OJ","VC")),
  cname = c("OJ intercept","VC intercept"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|--------------|----------|-------|--------|---------|--------|----|----------|
| OJ intercept | 6.485 | 2.024 | 2.429 | 10.5401 | 3.203 | 56 | 0.002243 |
| VC intercept | -3.429 | 2.024 | -7.485 | 0.6261 | -1.694 | 56 | 0.095824 |

```
contrast(m, a = list(dose = 100, supp = c("OJ", "VC")),
          cname = c("OJ asymptote", "VC asymptote"))
```

| | estimate | se | lower | upper | tvalue | df | pvalue |
|--------------|----------|-------|-------|-------|--------|----|-----------|
| OJ asymptote | 34.05 | 1.925 | 30.20 | 37.91 | 17.69 | 56 | 1.375e-24 |
| VC asymptote | 36.22 | 1.925 | 32.37 | 40.08 | 18.81 | 56 | 7.070e-26 |

But wouldn't it make sense to have something like the following?

$$E(Y_i) = \begin{cases} \beta_0 + \beta_1 2^{-x_i/h_{OJ}}, & \text{if the supplement type of the } i\text{-th observation is OJ,} \\ \beta_0 + \beta_1 2^{-x_i/h_{VC}}, & \text{if the supplement type of the } i\text{-th observation is VC,} \end{cases}$$

because at $x = 0$ and as $x \rightarrow \infty$ there should be no difference in the supplement type, but there might be a difference in how “fast” the expected response increases with dose. But unless we *know* h_{OJ} and h_{VC} , this model would be *nonlinear* (i.e., the model is not linear if h_{OJ} and h_{VC} are *unknown parameters* as opposed to known values).