

Wednesday, Feb 22

## Iteratively Weighted Least Squares

*Iteratively weighted least squares* can be used when we assume that the variance is proportional to a function of the mean so that

$$\text{Var}(Y_i) \propto h[E(Y_i)],$$

where  $h$  is some specified function, implying that our weights should be

$$w_i = \frac{1}{h[E(Y_i)]}.$$

Because  $E(Y_i)$  is unknown we can use the estimate  $\hat{y}_i$  to obtain weights

$$w_i = \frac{1}{h(\hat{y}_i)}.$$

Because  $\hat{y}_i$  depends on the weights used in the weighted least squares algorithm, and  $w_i$  depends on  $\hat{y}_i$ , we can use the following algorithm known as *iteratively weighted least squares*.

1. Estimate the model using *ordinary least squares* where all  $w_i = 1$ .
2. Compute weights as  $w_i = 1/h(\hat{y}_i)$ .
3. Estimate the model using *weighted least squares* with the weights  $w_i = 1/h(\hat{y}_i)$ .

The second and third steps can be repeated until the estimates and thus the weights stop changing. Typically only a few iterations are necessary.

**Example:** Consider again following data from a study on the effects of fuel reduction on biomass.


```
library(trtools) # for biomass data
```

```
m.ols <- lm(suitable ~ -1 + treatment:total, data = biomass)
summary(m.ols)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t )
treatmentn:total	0.1056	0.04183	2.524	1.31e-02
treatmenty:total	0.1319	0.01121	11.773	7.61e-21

```
d <- expand.grid(treatment = c("n", "y"), total = seq(0, 2767, length = 10))
d$yhat <- predict(m.ols, newdata = d)
```

```
p <- ggplot(biomass, aes(x = total, y = suitable, color = treatment)) +
  geom_point() + geom_line(aes(y = yhat), data = d) + theme_minimal() +
  labs(x = "Total Biomass (kg/ha)", y = "Suitable Biomass (kg/ha)",
       color = "Treatment")
plot(p)
```



```
biomass$yhat <- predict(m.ols)
biomass$rest <- rstudent(m.ols)

p <- ggplot(biomass, aes(x = yhat, y = rest, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Studentized Residual",
       color = "Treatment")
plot(p)
```



Assume that  $\text{Var}(Y_i) \propto E(Y_i)$ , which means the weights should be  $w_i = 1/E(Y_i)$ . We can program the iteratively weighted least squares algorithm as follows.

```
biomass$w <- 1 # initial weights are all equal to one
for (i in 1:5) {
  m.wls <- lm(suitable ~ -1 + treatment:total, weights = w, data = biomass)
  print(coef(m.wls)) # optional
  biomass$w <- 1 / predict(m.wls)
}
```


```
treatmentn:total treatmenty:total
      0.1056      0.1319
treatmentn:total treatmenty:total
      0.1155      0.1578
treatmentn:total treatmenty:total
      0.1155      0.1578
treatmentn:total treatmenty:total
      0.1155      0.1578
treatmentn:total treatmenty:total
      0.1155      0.1578
```

Now let's take a look at the residuals.

```
biomass$yhat <- predict(m.wls)
biomass$rest <- rstudent(m.wls)

p <- ggplot(biomass, aes(x = yhat, y = rest, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Studentized Residual",
```

```
color = "Treatment")
plot(p)
```




That may not be quite enough. Suppose we assume that  $\text{Var}(Y_i) \propto E(Y_i)^p$  where  $p = 2$ .

```
biomass$w <- 1 # initial weights are all equal to one
for (i in 1:5) {
  m.wls <- lm(suitable ~ -1 + treatment:total, weights = w, data = biomass)
  biomass$w <- 1 / predict(m.wls)^2
}
```

Now let's take a look at the residuals.

```
biomass$yhat <- predict(m.wls)
biomass$rest <- rstudent(m.wls)

p <- ggplot(biomass, aes(x = yhat, y = rest, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Studentized Residual",
       color = "Treatment")
plot(p)
```



Better. Maybe too much? We could try  $p = 1.5$  or something like that. The residuals do get a little strange for higher predicted values, but we'll leave it here.

The model is  $E(S_i) = \beta_1 n_i t_i + \beta_2 y_i t_i$ , where  $n_i$  and  $y_i$  are indicator variables for if the  $i$ -th plot was treated or not by fuel reduction. We can also write the model as

$$E(S_i) = \begin{cases} \beta_1 t_i, & \text{if the } i\text{-th plot was not treated by fuel reduction,} \\ \beta_2 t_i, & \text{if the } i\text{-th plot was treated by fuel reduction.} \end{cases}$$

We can use  $\beta_2 - \beta_1$  for inferences about the treatment effect.

```
lincon(m.ols, a = c(-1,1))
```

	estimate	se	lower	upper	tvalue	df	pvalue
(-1,1),0	0.02634	0.0433	-0.05953	0.1122	0.6082	104	0.5444

```
lincon(m.wls, a = c(-1,1))
```

	estimate	se	lower	upper	tvalue	df	pvalue
(-1,1),0	0.06386	0.02359	0.01708	0.1106	2.707	104	0.007937

The `contrast` function from the **trtools** package can also do this. It can make inferences for a *difference of differences*.

```
contrast(m.wls,
  a = list(treatment = "y", total = 1),
  b = list(treatment = "y", total = 0),
  u = list(treatment = "n", total = 1),
  v = list(treatment = "n", total = 0))
```

	estimate	se	lower	upper	tvalue	df	pvalue

```
0.06386 0.02359 0.01708 0.1106 2.707 104 0.007937
```

This estimates  $E(Y_a) - E(Y_b) - [E(Y_u) - E(Y_v)]$ . This can also be done using the `emtrends` function from the `emmeans` package.

```
library(emmeans)
emtrends(m.wls, ~treatment, var = "total") # estimate slopes
```

```
treatment total.trend      SE df lower.CL upper.CL
n              0.125 0.0183 104  0.0888   0.161
y              0.189 0.0149 104  0.1593   0.219
```

Confidence level used: 0.95

```
pairs(emtrends(m.wls, ~ treatment, var = "total")) # estimate difference between slopes
```

```
contrast estimate      SE df t.ratio p.value
n - y      -0.0639 0.0236 104  -2.707  0.0079
```

Recall that both the `emmeans` and `trtools` packages have a `contrast` function. To avoid conflicts or having to use `trtools::contrast` to call the `contrast` function from the `trtools` package later, we can unload the `emmeans` package using `detach`.

```
detach(package:emmeans)
```

Yet another approach to compare the slopes is to change the parameterization. Consider the following model.

```
m.wls <- lm(suitable ~ -1 + total + total:treatment, weights = w, data = biomass)
summary(m.wls)$coefficients
```

```
              Estimate Std. Error t value Pr(>|t|)
total          0.18892    0.01493  12.656 8.836e-23
total:treatment -0.06386    0.02359  -2.707 7.937e-03
```

From `summary` we can see that this model can be written as

$$E(S_i) = \beta_1 t_i + \beta_2 t_i n_i,$$

where  $n_i$  is an indicator variable where  $n_i = 1$  if the treatment was not applied to the  $i$ -th plot, add  $n_i = 0$  otherwise, so we can also write the model as

$$E(S_i) = \begin{cases} (\beta_1 + \beta_2)t_i, & \text{if the } i\text{-th plot was not treated by fuel reduction,} \\ \beta_1 t_i, & \text{if the } i\text{-th plot was treated by fuel reduction.} \end{cases}$$

Note that the meaning of  $\beta_1$  and  $\beta_2$  have changed here. The slopes of the lines with and without treatment are  $\beta_1$  and  $\beta_1 + \beta_2$ , respectively, and the difference between the slopes is  $\beta_1 - (\beta_1 + \beta_2) = -\beta_2$ . So inferences for  $\beta_2$  are for the difference in the slopes (after we reverse the sign). Although not necessary, we can change the reference category to avoid having to reverse the sign.

```
biomass$treatment <- relevel(biomass$treatment, ref = "y")
m.wls <- lm(suitable ~ -1 + total + total:treatment, weights = w, data = biomass)
summary(m.wls)$coefficients
```

```
              Estimate Std. Error t value Pr(>|t|)
total          0.12506    0.01827   6.847 5.428e-10
total:treatmenty 0.06386    0.02359   2.707 7.937e-03
```

Now the model can be written as

$$E(S_i) = \beta_1 t_i + \beta_2 t_i n_i,$$

or


$$E(S_i) = \begin{cases} \beta_1 t_i, & \text{if the } i\text{-th plot was not treated by fuel reduction,} \\ (\beta_1 + \beta_2) t_i, & \text{if the } i\text{-th plot was treated by fuel reduction.} \end{cases}$$

Note: For some reason the reference category (y) is getting an indicator variable here, where normally it does not. I am not sure if this is a bug or intentional, but it appears to be due to the somewhat unusual parameterization I am using.

## Parametric Models for Heteroscedasticity

**Example:** Consider the following data where variability appears to vary by treatment.

```
library(trtools) # for pulse data
p <- ggplot(pulse, aes(x = pulse1, y = pulse2, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Pulse Before", y = "Pulse After", color = "Treatment") +
  theme(legend.position = c(0.85, 0.2))
plot(p)
```



There is one case with missing values on pulse1 and pulse2.

```
subset(pulse, !complete.cases(pulse)) # show observations with missing data
```

```
height weight age gender smokes alcohol exercise treatment pulse1 pulse2 year
76      173    64   20 female      no      yes moderate      sat      NA      NA   97
```

This will cause problems so we are going to remove it.


```
pulse <- subset(pulse, complete.cases(pulse)) # overwrite pulse with only complete cases
```

Let's consider a simple linear model.

```
m <- lm(pulse2 ~ treatment + pulse1 + treatment:pulse1, data = pulse)
summary(m)$coefficients
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	59.41757	10.4467	5.68767	1.171e-07
treatmentsat	-51.25896	15.7451	-3.25554	1.524e-03
pulse1	0.89363	0.1357	6.58544	1.841e-09
treatmentsat:pulse1	-0.01437	0.2049	-0.07011	9.442e-01

```
pulse$yhat <- predict(m)
pulse$rest <- rstudent(m)
p <- ggplot(pulse, aes(x = yhat, y = rest, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Studentized Residual",
       color = "Treatment") +
  theme(legend.position = c(0.8,0.2))
plot(p)
```



Consider that the model assumed by `lm` is

$$E(Y_i) = \beta_0 + \beta_1 t_i + \beta_2 x_i + \beta_3 t_i x_i, \quad (1)$$

$$\text{Var}(Y_i) = \sigma^2, \quad (2)$$

where  $Y_i$  is the second pulse measurement,  $t_i$  is an indicator variable for the treatment (i.e.,  $t_i = 1$  if the  $i$ -th observation was from the sitting treatment condition, and  $t_i = 0$  otherwise), and  $x_i$  is the first pulse



measurement. Maybe it would make sense to have something like

$$\text{Var}(Y_i) = \begin{cases} \sigma_s^2, & \text{if the } i\text{-th observation is from the sitting treatment,} \\ \sigma_r^2, & \text{if the } i\text{-th observation is from the running treatment.} \end{cases}$$

We can estimate such a model using the `glS` function from the **nlme** package.

```
library(nlme) # should come with R
m <- gls(pulse2 ~ treatment + pulse1 + treatment:pulse1, data = pulse,
  method = "ML", weights = varIdent(form = ~ 1|treatment))
summary(m)
```

```
Generalized least squares fit by maximum likelihood
Model: pulse2 ~ treatment + pulse1 + treatment:pulse1
Data: pulse
AIC    BIC logLik
763.1 779.3 -375.6
```

```
Variance function:
Structure: Different standard deviations per stratum
Formula: ~1 | treatment
Parameter estimates:
      sat      ran
1.000 5.723
```

```
Coefficients:
              Value Std.Error t-value p-value
(Intercept)   59.42   15.755   3.771 0.0003
treatmentsat -51.26   16.058  -3.192 0.0019
pulse1         0.89    0.205   4.367 0.0000
treatmentsat:pulse1 -0.01    0.209  -0.069 0.9452
```


```
Correlation:
              (Intr) trtmnt pulse1
treatmentsat -0.981
pulse1       -0.980 0.962
treatmentsat:pulse1 0.962 -0.980 -0.981
```

```
Standardized residuals:
      Min      Q1      Med      Q3      Max
-2.0920 -0.7688 0.1026 0.5886 2.1968
```

```
Residual standard error: 3.634
Degrees of freedom: 109 total; 105 residual
```

Note the different syntax for extracting standardized residuals.

```
pulse$yhat <- predict(m)
pulse$resz <- residuals(m, type = "p") # note different syntax
p <- ggplot(pulse, aes(x = yhat, y = resz, color = treatment)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Standardized Residual",
    color = "Treatment")
plot(p)
```



Here is an example with the CancerSurvival data.

```
library(Stat2Data)
data(CancerSurvival)
m <- gls(Survival ~ Organ, data = CancerSurvival,
  method = "ML", weights = varIdent(form = ~ 1|Organ))
summary(m)
```

```
Generalized least squares fit by maximum likelihood
Model: Survival ~ Organ
Data: CancerSurvival
AIC   BIC logLik
976.8 998.4 -478.4
```

```
Variance function:
Structure: Different standard deviations per stratum
Formula: ~1 | Organ
Parameter estimates:
Stomach Bronchus   Colon   Ovary   Breast
1.0000   0.6119   1.2455   3.0141   3.5504
```


```
Coefficients:
              Value Std.Error t-value p-value
(Intercept)  1395.9    371.0   3.763  0.0004
OrganBronchus -1184.3    374.5  -3.162  0.0025
OrganColon    -938.5    385.5  -2.435  0.0179
OrganOvary    -511.6    565.2  -0.905  0.3691
OrganStomach -1109.9    383.2  -2.896  0.0053
```

```
Correlation:
              (Intr) OrgnBr OrgnC1 OrgnOv
OrganBronchus -0.991
OrganColon    -0.962  0.953
OrganOvary    -0.656  0.650  0.632
OrganStomach  -0.968  0.959  0.932  0.635
```

```
Standardized residuals:
      Min      Q1      Med      Q3      Max
-1.1613 -0.6824 -0.2878  0.1748  3.3435
```

```
Residual standard error: 332.7
Degrees of freedom: 64 total; 59 residual
```

```
CancerSurvival$yhat <- predict(m)
CancerSurvival$resz <- residuals(m, type = "p")
p <- ggplot(CancerSurvival, aes(x = yhat, y = resz, color = Organ)) +
  geom_point() + theme_minimal() +
  labs(x = "Predicted Value", y = "Standardized Residual", color = "Organ")
plot(p)
```



Comments about parametric models for heteroscedasticity.

**Advantages:** Potentially very effective *if* we can specify an accurate model for the variance.

**Disadvantages:** If we do not specify an accurate model for the variance, it may bias estimation of parameters concerning the expected response.

## Heteroscedastic Consistent Standard Errors

The idea is to estimate the model parameters using ordinary least squares, but estimate the standard errors in such a way that we do not assume homoscedasticity. This is sometimes called *heteroscedastic consistent standard errors*, *robust standard errors*, or *sandwich estimators*.

**Example:** Consider again the cancer survival data.

```
m <- lm(Survival ~ Organ, data = CancerSurvival)
```

The **sandwich** package provides resources for using heteroscedastic-consistent standard errors. Technically, what is being estimated is the *covariance matrix* of the parameter estimators.

```
library(sandwich) # for vcovHC used below
vcov(m)           # bad estimate if there is heteroscedasticity
```

	(Intercept)	OrganBronchus	OrganColon	OrganOvary	OrganStomach
(Intercept)	40752	-40752	-40752	-40752	-40752
OrganBronchus	-40752	67121	40752	40752	40752
OrganColon	-40752	40752	67121	40752	40752
OrganOvary	-40752	40752	40752	115464	40752
OrganStomach	-40752	40752	40752	40752	75235

```
vcovHC(m) # better estimate if there is heteroscedasticity
```

	(Intercept)	OrganBronchus	OrganColon	OrganOvary	OrganStomach
(Intercept)	153504	-153504	-153504	-153504	-153504
OrganBronchus	-153504	156256	153504	153504	153504
OrganColon	-153504	153504	164908	153504	153504
OrganOvary	-153504	153504	153504	394879	153504
OrganStomach	-153504	153504	153504	153504	163498

The square root of the diagonal elements are the standard errors.

```
sqrt(diag(vcov(m))) # bad estimates of the standard errors
```

	(Intercept)	OrganBronchus	OrganColon	OrganOvary	OrganStomach
	201.9	259.1	259.1	339.8	274.3

```
sqrt(diag(vcovHC(m))) # better estimates of the standard errors
```

	(Intercept)	OrganBronchus	OrganColon	OrganOvary	OrganStomach
	391.8	395.3	406.1	628.4	404.3

But the usual way to interface with the functions in the **sandwich** package is through other functions.

	Estimate	Std. Error	t value	Pr(> t )	2.5 %	97.5 %
(Intercept)	1395.9	201.9	6.915	3.770e-09	992	1799.9
OrganBronchus	-1184.3	259.1	-4.571	2.530e-05	-1703	-665.9
OrganColon	-938.5	259.1	-3.622	6.083e-04	-1457	-420.1
OrganOvary	-511.6	339.8	-1.506	1.375e-01	-1192	168.4
OrganStomach	-1109.9	274.3	-4.046	1.533e-04	-1659	-561.1

```
confint(m) # bad confidence intervals due to bad standard error estimates
```

	2.5 %	97.5 %
(Intercept)	992	1799.9
OrganBronchus	-1703	-665.9
OrganColon	-1457	-420.1
OrganOvary	-1192	168.4
OrganStomach	-1659	-561.1

```
library(lmtest) # for coeftest and coefci used below
coeftest(m, vcov = vcovHC) # better standard error estimates
```

t test of coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
(Intercept)      1396         392    3.56 0.00073 ***
OrganBronchus   -1184         395   -3.00 0.00400 **
OrganColon       -938         406   -2.31 0.02434 *
OrganOvary       -512         628   -0.81 0.41886
OrganStomach    -1110         404   -2.74 0.00801 **
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coefci(m, vcov = vcovHC) # better confidence intervals
```

```

              2.5 % 97.5 %
(Intercept)    611.9 2179.9
OrganBronchus -1975.3 -393.3
OrganColon     -1751.1 -125.9
OrganOvary     -1769.0  745.8
OrganStomach   -1919.0 -300.8
```

Both `lincon` and `contrast` will accept a `fcov` argument to provide a function to estimate standard errors.

```
lincon(m, fcov = vcovHC)
```

```

      estimate      se  lower upper tvalue df   pvalue
(Intercept)  1395.9 391.8   611.9 2179.9  3.5628 59 0.0007337
OrganBronchus -1184.3 395.3 -1975.3 -393.3 -2.9961 59 0.0039950
OrganColon     -938.5 406.1 -1751.1 -125.9 -2.3111 59 0.0243421
OrganOvary     -511.6 628.4 -1769.0  745.8 -0.8141 59 0.4188611
OrganStomach   -1109.9 404.3 -1919.0 -300.8 -2.7449 59 0.0080080
```

```
organs <- sort(unique(CancerSurvival$Organ)) # sorted organ names
contrast(m, a = list(Organ = organs),
  cnames = organs, fcov = vcovHC)
```

```

      estimate      se  lower upper tvalue df   pvalue
Breast      1395.9 391.80  611.93 2179.9  3.563 59 7.337e-04
Bronchus     211.6  52.46  106.61  316.6  4.033 59 1.604e-04
Colon        457.4 106.79  243.72  671.1  4.283 59 6.884e-05
Ovary        884.3 491.30  -98.75 1867.4  1.800 59 7.698e-02
Stomach       286.0  99.97   85.96  486.0  2.861 59 5.836e-03
```

```
lincon(m, a = c(1,0,0,0,1), fcov = vcovHC)
```

```

      estimate      se  lower upper tvalue df   pvalue
(1,0,0,0,1),0    286 99.97  85.96  486  2.861 59 0.005836
```

You can use a similar approach with the `emmeans` function from the **emmeans** package, but there the argument is `vcov`.

```
library(emmeans)
emmeans(m, ~Organ, vcov = vcovHC)
```

```
Organ    emmean    SE df lower.CL upper.CL
```

Breast	1396	391.8	59	611.9	2180
Bronchus	212	52.5	59	106.6	317
Colon	457	106.8	59	243.7	671
Ovary	884	491.3	59	-98.8	1867
Stomach	286	100.0	59	86.0	486

Confidence level used: 0.95

```
pairs(emmeans(m, ~Organ, vcov = vcovHC), adjust = "none", infer = TRUE)
```

contrast	estimate	SE	df	lower.CL	upper.CL	t.ratio	p.value
Breast - Bronchus	1184.3	395	59	393	1975.3	2.996	0.0040
Breast - Colon	938.5	406	59	126	1751.1	2.311	0.0243
Breast - Ovary	511.6	628	59	-746	1769.0	0.814	0.4189
Breast - Stomach	1109.9	404	59	301	1919.0	2.745	0.0080
Bronchus - Colon	-245.8	119	59	-484	-7.7	-2.066	0.0432
Bronchus - Ovary	-672.7	494	59	-1661	315.9	-1.362	0.1785
Bronchus - Stomach	-74.4	113	59	-300	151.5	-0.659	0.5124
Colon - Ovary	-426.9	503	59	-1433	579.1	-0.849	0.3992
Colon - Stomach	171.4	146	59	-121	464.1	1.172	0.2460
Ovary - Stomach	598.3	501	59	-405	1601.6	1.193	0.2375

Confidence level used: 0.95

Use the function `waldtest` in place of `anova` when using heteroscedastic-consistent standard errors.

```
m.full <- lm(Survival ~ Organ, data = CancerSurvival)
m.null <- lm(Survival ~ 1, data = CancerSurvival)
waldtest(m.null, m.full, vcov = vcovHC)
```

Wald test

```
Model 1: Survival ~ 1
Model 2: Survival ~ Organ
  Res.Df Df    F Pr(>F)
1      63
2      59  4 3.52 0.012 *
```

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Comments about heteroscedastic-consistent standard errors:

**Advantages:** Does not require us to specify a variance structure/function. We let the data inform the estimator.

**Disadvantages:** Highly dependent on the data to help produce better estimates of the standard errors, and tends to work well only if  $n$  is relatively large.

Note: There are a variety of variations of the “sandwich” estimator. Different estimators can be specified through the `type` argument to `vcovHC` so instead of writing `vcov = vcovHC` or `fcov = vcovHC` we write `vcov = function(m) vcovHC(m, type = "HCO")` or `vcov = function(m) vcovHC(m, type = "HCO")` if we wanted to use that particular type of estimator (sometimes called “White’s estimator”).