

Friday, Mar 10

Using the emmeans Package for Poisson and Logistic Regression

The **emmeans** package can be used to produce some of the same inferences that are obtained using **contrast** with respect to estimated expected rates/probabilities as well as rate/odds ratios.

Example: Consider the following Poisson regression model for the **ceriodaphniastrain** data.

```
fleas <- trtools::ceriodaphniastrain
fleas$strain <- factor(fleas$strain, levels = c(1,2), labels = c("a","b"))
m <- glm(count ~ concentration * strain, family = poisson, data = fleas)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	4.4811	0.04350	103.008	0.000e+00
concentration	-1.5979	0.06244	-25.592	1.862e-144
strainb	-0.3367	0.06704	-5.022	5.114e-07
concentration:strainb	0.1253	0.09385	1.336	1.817e-01

We can compute the expected count for a concentration of two for each strain using **contrast**.

```
trtools::contrast(m, tf = exp,
  a = list(strain = c("a","b"), concentration = 2))
```

estimate	lower	upper
3.616	2.970	4.402
3.318	2.671	4.122

And we can do it using **emmeans** if we specify **type = "response"** and use the **at** argument to specify the value of any quantitative explanatory variables.

```
library(emmeans)
emmeans(m, ~ strain, type = "response", at = list(concentration = 2))
```

strain	rate	SE	df	asympt.LCL	asympt.UCL
a	3.62	0.363	Inf	2.97	4.40
b	3.32	0.367	Inf	2.67	4.12

Confidence level used: 0.95

Intervals are back-transformed from the log scale

```
emmeans(m, ~ strain|concentration, type = "response", at = list(concentration = c(1,2,3)))
```

```
concentration = 1:
strain rate SE df asympt.LCL asympt.UCL
a 17.872 0.815 Inf 16.343 19.54
b 14.467 0.725 Inf 13.113 15.96
```

```
concentration = 2:
strain rate SE df asympt.LCL asympt.UCL
a 3.616 0.363 Inf 2.970 4.40
b 3.318 0.367 Inf 2.671 4.12
```

```
concentration = 3:
  strain  rate    SE df asymp.LCL asymp.UCL
a        0.732 0.118 Inf    0.534    1.00
b        0.761 0.136 Inf    0.537    1.08
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

```
emmeans(m, ~ concentration|strain, type = "response", at = list(concentration = c(1,2,3)))
```

```
strain = a:
  concentration  rate    SE df asymp.LCL asymp.UCL
1 17.872 0.815 Inf 16.343 19.54
2 3.616 0.363 Inf 2.970 4.40
3 0.732 0.118 Inf 0.534 1.00
```

```
strain = b:
  concentration  rate    SE df asymp.LCL asymp.UCL
1 14.467 0.725 Inf 13.113 15.96
2 3.318 0.367 Inf 2.671 4.12
3 0.761 0.136 Inf 0.537 1.08
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

```
emmeans(m, ~ concentration*strain, type = "response", at = list(concentration = c(1,2,3)))
```

```
concentration strain  rate    SE df asymp.LCL asymp.UCL
1 a      17.872 0.815 Inf 16.343 19.54
2 a       3.616 0.363 Inf 2.970 4.40
3 a       0.732 0.118 Inf 0.534 1.00
1 b      14.467 0.725 Inf 13.113 15.96
2 b       3.318 0.367 Inf 2.671 4.12
3 b       0.761 0.136 Inf 0.537 1.08
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Note that `emmeans` does produce a valid standard error on the scale of the expected count/rate which `trtools::contrast` does not (by default), and that `trtools::contrast` will show the test statistic and p-value on the log scale if we omit the `tf = exp` argument.

We can compute the rate ratio to compare the two strains at a given concentration.

```
trtools::contrast(m, tf = exp,
  a = list(strain = "a", concentration = 2),
  b = list(strain = "b", concentration = 2))
```

```
estimate lower upper
1.09 0.8132 1.46
```

```
pairs(emmeans(m, ~ strain, type = "response",
  at = list(concentration = 2)), infer = TRUE)
```

```
contrast ratio    SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      1.09 0.163 Inf    0.813    1.46    1    0.576 0.5648
```

Confidence level used: 0.95
 Intervals are back-transformed from the log scale
 Tests are performed on the log scale

```
pairs(emmeans(m, ~ strain*concentration, type = "response",
  at = list(concentration = c(1,2,3))), by = "concentration", infer = TRUE)
```

```
concentration = 1:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      1.235 0.0837 Inf      1.082      1.41      1   3.118  0.0018
```

```
concentration = 2:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      1.090 0.1628 Inf      0.813      1.46      1   0.576  0.5648
```

```
concentration = 3:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      0.961 0.2308 Inf      0.601      1.54      1  -0.164  0.8698
```

Confidence level used: 0.95
 Intervals are back-transformed from the log scale
 Tests are performed on the log scale

```
pairs(emmeans(m, ~ strain|concentration, type = "response",
  at = list(concentration = c(1,2,3))), infer = TRUE)
```

```
concentration = 1:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      1.235 0.0837 Inf      1.082      1.41      1   3.118  0.0018
```

```
concentration = 2:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      1.090 0.1628 Inf      0.813      1.46      1   0.576  0.5648
```

```
concentration = 3:
contrast ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
a / b      0.961 0.2308 Inf      0.601      1.54      1  -0.164  0.8698
```

Confidence level used: 0.95
 Intervals are back-transformed from the log scale
 Tests are performed on the log scale

What about the rate ratio for the effect of concentration?

```
trtools::contrast(m, tf = exp,
  a = list(strain = c("a","b"), concentration = 2),
  b = list(strain = c("a","b"), concentration = 1))
```

```
estimate lower upper
0.2023 0.1790 0.2287
0.2293 0.1999 0.2631
```

```
emmeans(m, ~concentration|strain,
  at = list(concentration = c(2,1)), type = "response")
```

```
strain = a:
concentration rate      SE df asymp.LCL asymp.UCL
```

2	3.62	0.363	Inf	2.97	4.40
1	17.87	0.815	Inf	16.34	19.54

```
strain = b:
  concentration rate SE df asymp.LCL asymp.UCL
2 3.32 0.367 Inf 2.67 4.12
1 14.47 0.725 Inf 13.11 15.96
```

Confidence level used: 0.95
Intervals are back-transformed from the log scale

```
pairs(emmeans(m, ~concentration|strain,
  at = list(concentration = c(2,1)), type = "response"))
```

```
strain = a:
contrast ratio SE df null z.ratio p.value
concentration2 / concentration1 0.202 0.0126 Inf 1 -25.592 <.0001
```

```
strain = b:
contrast ratio SE df null z.ratio p.value
concentration2 / concentration1 0.229 0.0161 Inf 1 -21.015 <.0001
```

Tests are performed on the log scale

```
pairs(emmeans(m, ~concentration*strain,
  at = list(concentration = c(2,1)), type = "response"), by = "strain")
```

```
strain = a:
contrast ratio SE df null z.ratio p.value
concentration2 / concentration1 0.202 0.0126 Inf 1 -25.592 <.0001
```

```
strain = b:
contrast ratio SE df null z.ratio p.value
concentration2 / concentration1 0.229 0.0161 Inf 1 -21.015 <.0001
```

Tests are performed on the log scale

We can also make inferences to *compare* the two rate ratios.

```
trtools::contrast(m, tf = exp,
  a = list(strain = "b", concentration = 2),
  b = list(strain = "b", concentration = 1),
  u = list(strain = "a", concentration = 2),
  v = list(strain = "a", concentration = 1))
```

estimate	lower	upper
1.134	0.9431	1.362

```
pairs(pairs(emmeans(m, ~concentration|strain,
  at = list(concentration = c(2,1)), type = "response")), by = NULL, reverse = TRUE)
```

contrast	ratio	SE	df	null
(concentration2 / concentration1 b) / (concentration2 / concentration1 a)	1.13	0.106	Inf	1
z.ratio	p.value			
1.335	0.1817			

Tests are performed on the log scale

This is a ratio of rate ratios. To just get the test statistic you can leave off `type = "response"`.

```
pairs(pairs(emmeans(m, ~concentration|strain,
  at = list(concentration = c(2,1))), by = NULL, reverse = TRUE)

contrast                                     estimate      SE df
(concentration2 - concentration1 b) - (concentration2 - concentration1 a) 0.125 0.0939 Inf
z.ratio p.value
1.335 0.1817
```

Results are given on the log (not the response) scale.

Example: Consider the following logistic regression model for the insecticide data.

```
m <- glm(cbind(deaths, total-deaths) ~ insecticide * deposit,
  family = binomial, data = trtools::insecticide)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.81091	0.35845	-7.84177	4.442e-15
insecticideboth	1.22575	0.67176	1.82468	6.805e-02
insecticideDDT	-0.03893	0.50722	-0.07676	9.388e-01
deposit	0.62207	0.07786	7.98986	1.351e-15
insecticideboth:deposit	0.37010	0.20897	1.77109	7.655e-02
insecticideDDT:deposit	-0.14143	0.10376	-1.36301	1.729e-01

We can use `trtools::contrast` or `emmeans` to produce estimates of the probability of death for a given insecticide at a given deposit value.

```
trtools::contrast(m, tf = plogis,
  a = list(insecticide = c("g-BHC", "both", "DDT"), deposit = 5),
  cnames = c("g-BHC", "both", "DDT"))
```

	estimate	lower	upper
g-BHC	0.5743	0.5027	0.6429
both	0.9669	0.9212	0.9865
DDT	0.3902	0.3289	0.4550

```
emmeans(m, ~ insecticide, type = "response", at = list(deposit = 5))
```

insecticide	prob	SE	df	asympt.LCL	asympt.UCL
g-BHC	0.574	0.0360	Inf	0.503	0.643
both	0.967	0.0149	Inf	0.921	0.987
DDT	0.390	0.0323	Inf	0.329	0.455

Confidence level used: 0.95

Intervals are back-transformed from the logit scale

Again, `emmeans` produces a valid standard error on the probability scale while `trtools::contrast` does not, and `trtools::contrast` will produce test statistics and p-values on the logit scale when the `tf = plogis` argument is omitted.

We can compute odds ratios to compare the insecticides at a given deposit.

```
pairs(emmeans(m, ~ insecticide, type = "response",
  at = list(deposit = 5)), adjust = "none", infer = TRUE)
```

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.05	0.023	Inf	0.018	0.12	1	-6.275	<.0001
(g-BHC) / DDT	2.11	0.423	Inf	1.424	3.12	1	3.724	0.0002

```
both / DDT          45.71 22.260 Inf    17.600    118.72    1    7.849 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
trtools::contrast(m, tf = exp,
  a = list(insecticide = c("g-BHC", "g-BHC", "both"), deposit = 5),
  b = list(insecticide = c("both", "DDT", "DDT"), deposit = 5),
  cnames = c("g-BHC / both", "g-BHC / DDT", "both / DDT"))
```

	estimate	lower	upper
g-BHC / both	0.04613	0.01765	0.1206
g-BHC / DDT	2.10871	1.42385	3.1230
both / DDT	45.71097	17.59954	118.7243

We can flip/reverse the odds ratios if desired (which can also be done with rate ratios).

```
pairs(emmeans(m, ~ insecticide, type = "response",
  at = list(deposit = 5)), adjust = "none", reverse = TRUE, infer = TRUE)
```

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
both / (g-BHC)	21.677	10.628	Inf	8.293	56.67	1	6.275	<.0001
DDT / (g-BHC)	0.474	0.095	Inf	0.320	0.70	1	-3.724	0.0002
DDT / both	0.022	0.011	Inf	0.008	0.06	1	-7.849	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
trtools::contrast(m, tf = exp,
  a = list(insecticide = c("both", "DDT", "DDT"), deposit = 5),
  b = list(insecticide = c("g-BHC", "g-BHC", "both"), deposit = 5),
  cnames = c("both / g-BHC", "DDT / g-BHC", "DDT / both"))
```

	estimate	lower	upper
both / g-BHC	21.67723	8.292521	56.66581
DDT / g-BHC	0.47422	0.320208	0.70232
DDT / both	0.02188	0.008423	0.05682

We can estimate the odds ratios at several values of deposit.

```
pairs(emmeans(m, ~ insecticide|deposit, type = "response",
  at = list(deposit = c(4,5,6))), adjust = "none", infer = TRUE)
```

deposit = 4:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.07	0.02	Inf	0.035	0.13	1	-8.239	<.0001
(g-BHC) / DDT	1.83	0.37	Inf	1.234	2.72	1	3.004	0.0027
both / DDT	27.41	9.12	Inf	14.274	52.62	1	9.947	<.0001

deposit = 5:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.05	0.02	Inf	0.018	0.12	1	-6.275	<.0001
(g-BHC) / DDT	2.11	0.42	Inf	1.424	3.12	1	3.724	0.0002
both / DDT	45.71	22.26	Inf	17.600	118.72	1	7.849	<.0001

deposit = 6:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.03	0.02	Inf	0.008	0.12	1	-5.080	<.0001
(g-BHC) / DDT	2.43	0.60	Inf	1.495	3.95	1	3.584	0.0003
both / DDT	76.24	51.04	Inf	20.529	283.13	1	6.474	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

```
pairs(emmeans(m, ~ insecticide*deposit, type = "response",
  at = list(deposit = c(4,5,6))), by = "deposit", adjust = "none", infer = TRUE)
```

deposit = 4:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.07	0.02	Inf	0.035	0.13	1	-8.239	<.0001
(g-BHC) / DDT	1.83	0.37	Inf	1.234	2.72	1	3.004	0.0027
both / DDT	27.41	9.12	Inf	14.274	52.62	1	9.947	<.0001

deposit = 5:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.05	0.02	Inf	0.018	0.12	1	-6.275	<.0001
(g-BHC) / DDT	2.11	0.42	Inf	1.424	3.12	1	3.724	0.0002
both / DDT	45.71	22.26	Inf	17.600	118.72	1	7.849	<.0001

deposit = 6:

contrast	odds.ratio	SE	df	asympt.LCL	asympt.UCL	null	z.ratio	p.value
(g-BHC) / both	0.03	0.02	Inf	0.008	0.12	1	-5.080	<.0001
(g-BHC) / DDT	2.43	0.60	Inf	1.495	3.95	1	3.584	0.0003
both / DDT	76.24	51.04	Inf	20.529	283.13	1	6.474	<.0001

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

Here is how we can estimate the odds ratios for the effect of deposit.

```
emmeans(m, ~deposit|insecticide, at = list(deposit = c(2,1)), type = "response") # odds
```

insecticide = g-BHC:

deposit	prob	SE	df	asympt.LCL	asympt.UCL
2	0.1727	0.0318	Inf	0.1190	0.244
1	0.1008	0.0261	Inf	0.0599	0.165

insecticide = both:

deposit	prob	SE	df	asympt.LCL	asympt.UCL
2	0.5985	0.0566	Inf	0.4844	0.703
1	0.3560	0.0892	Inf	0.2049	0.542

insecticide = DDT:

deposit	prob	SE	df	asympt.LCL	asympt.UCL
2	0.1314	0.0271	Inf	0.0867	0.194
1	0.0856	0.0232	Inf	0.0497	0.143

Confidence level used: 0.95

Intervals are back-transformed from the logit scale

```
pairs(emmeans(m, ~deposit|insecticide, at = list(deposit = c(2,1)),
  type = "response"), infer = TRUE) # odds ratios
```

```
insecticide = g-BHC:
contrast      odds.ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
deposit2 / deposit1      1.86 0.145 Inf      1.60      2.17    1   7.990 <.0001
```

```
insecticide = both:
contrast      odds.ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
deposit2 / deposit1      2.70 0.523 Inf      1.84      3.94    1   5.116 <.0001
```

```
insecticide = DDT:
contrast      odds.ratio    SE  df asymp.LCL asymp.UCL null z.ratio p.value
deposit2 / deposit1      1.62 0.111 Inf      1.41      1.85    1   7.007 <.0001
```

Confidence level used: 0.95

Intervals are back-transformed from the log odds ratio scale

Tests are performed on the log odds ratio scale

We can also compare the odds ratios.

```
pairs(pairs(emmeans(m, ~deposit|insecticide, at = list(deposit = c(2,1)))), by = NULL)
```

```
contrast                                     estimate    SE  df z.ratio p.value
(deposit2 - deposit1 g-BHC) - (deposit2 - deposit1 both)   -0.370 0.209 Inf  -1.771 0.1794
(deposit2 - deposit1 g-BHC) - (deposit2 - deposit1 DDT)    0.141 0.104 Inf   1.363 0.3605
(deposit2 - deposit1 both) - (deposit2 - deposit1 DDT)     0.511 0.206 Inf   2.487 0.0344
```

Results are given on the log odds ratio (not the response) scale.

P value adjustment: tukey method for comparing a family of 3 estimates

Here I have left off `type = "response"`. Including it will give ratios of odds ratios, which is a bit confusing, but if all we care about is whether the odds ratios are significantly different this is sufficient. Note that to avoid controlling for family-wise Type I error rate include the option `adjust = "none"`.

Relationship Between Poisson and Logistic Regression

Suppose C_i has a binomial distribution with parameters p_i and m_i so that

$$P(C_i = c) = \binom{m_i}{c} p_i^c (1 - p_i)^{m_i - c}.$$

Define the expected count as $E(C_i) = m_i p_i = \lambda_i$. Then $p_i = \lambda_i / m_i$ so we can write

$$P(C_i = c) = \binom{m_i}{c} \left(\frac{\lambda_i}{m_i} \right)^c \left(1 - \frac{\lambda_i}{m_i} \right)^{m_i - c}.$$

Then it can be shown that

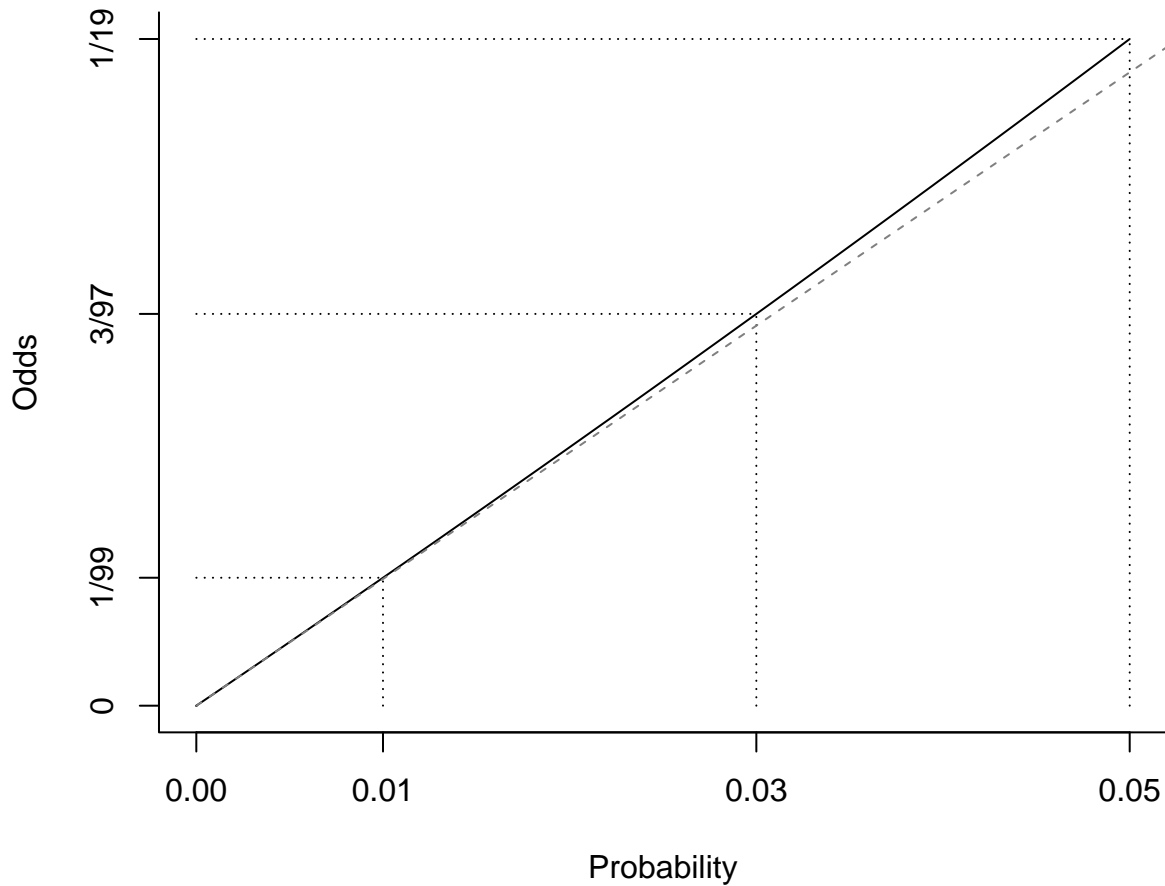
$$\lim_{m_i \rightarrow \infty} \binom{m_i}{c} \left(\frac{\lambda_i}{m_i} \right)^c \left(1 - \frac{\lambda_i}{m_i} \right)^{m_i - c} = \frac{e^{-\lambda_i} \lambda_i^c}{c!},$$

which is the Poisson distribution.

Thus *in practice* if p_i is small relative to m_i we can *approximate a binomial distribution with a Poisson distribution*. Furthermore there is a close relationship between the model parameters. In logistic regression we have

$$O_i = \exp(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik}),$$

where $O_i = p_i / (1 - p_i)$ is the odds of the event. But when p_i is very small then $O_i \approx p_i$.



So then

$$p_i \approx \exp(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik}),$$

and because $E(C_i) = m_i p_i$,

$$E(C_i) \approx \exp(\log m_i + \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \cdots + \beta_k x_{ik}),$$

where $\log m_i$ is used as an offset in a Poisson regression model. That is, we can model a proportion (approximately) as a rate in a Poisson regression model for events that are rare and when m_i (i.e., the denominator of the proportion) is relatively large. This is relatively common in large-scale observational studies.

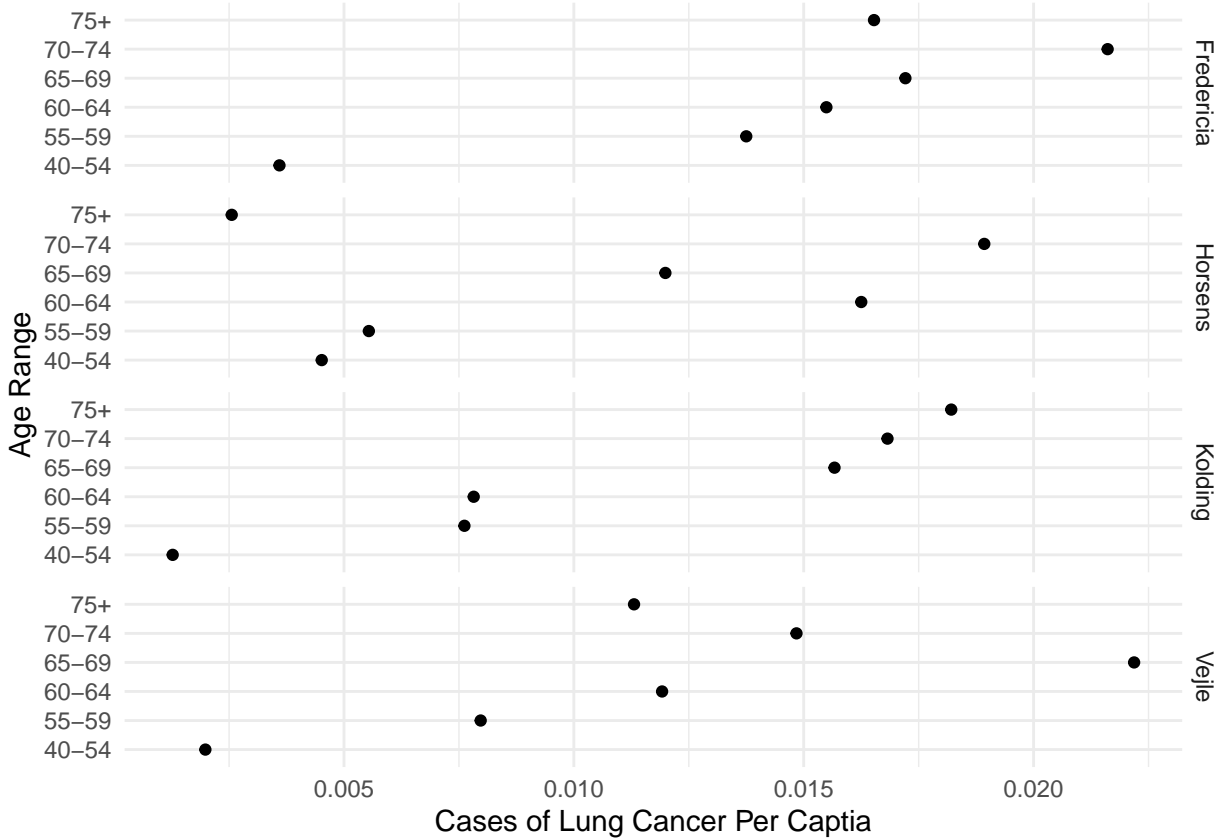
Example: Consider the following data on the incidence of lung cancer in four Danish cities.

```
library(ISwR) # for eba1977 data
head(eba1977)
```

	city	age	pop	cases
1	Fredericia	40-54	3059	11
2	Horsens	40-54	2879	13
3	Kolding	40-54	3142	4
4	Vejle	40-54	2520	5
5	Fredericia	55-59	800	11
6	Horsens	55-59	1083	6

```
p <- ggplot(eba1977, aes(x = age, y = cases/pop)) +
  geom_point() + facet_grid(city ~ .) + coord_flip() +
  labs(x = "Age Range", y = "Cases of Lung Cancer Per Captia") +
```

```
theme_minimal()
plot(p)
```



Consider both a logistic and Poisson regression models to compare the cities while controlling for age.

```
m.b <- glm(cbind(cases, pop-cases) ~ city + age, family = binomial, data = eba1977)
cbind(summary(m.b)$coefficients, confint(m.b))
```

	Estimate	Std. Error	z value	Pr(> z)	2.5 %	97.5 %
(Intercept)	-5.6262	0.2008	-28.021	9.132e-173	-6.0385	-5.249799
cityHorsens	-0.3345	0.1827	-1.830	6.719e-02	-0.6946	0.023561
cityKolding	-0.3764	0.1890	-1.991	4.646e-02	-0.7504	-0.007412
cityVejle	-0.2760	0.1891	-1.459	1.444e-01	-0.6503	0.093162
age55-59	1.1070	0.2490	4.445	8.771e-06	0.6159	1.596828
age60-64	1.5291	0.2325	6.577	4.812e-11	1.0760	1.991225
age65-69	1.7819	0.2305	7.732	1.061e-14	1.3335	2.240675
age70-74	1.8727	0.2365	7.918	2.415e-15	1.4105	2.341695
age75+	1.4289	0.2512	5.688	1.289e-08	0.9328	1.922467

```
m.p <- glm(cases ~ offset(log(pop)) + city + age, family = poisson, data = eba1977)
cbind(summary(m.p)$coefficients, confint(m.p))
```

	Estimate	Std. Error	z value	Pr(> z)	2.5 %	97.5 %
(Intercept)	-5.6321	0.2003	-28.125	4.911e-174	-6.0433	-5.256725
cityHorsens	-0.3301	0.1815	-1.818	6.899e-02	-0.6878	0.025582
cityKolding	-0.3715	0.1878	-1.978	4.789e-02	-0.7432	-0.004967
cityVejle	-0.2723	0.1879	-1.450	1.472e-01	-0.6441	0.094356
age55-59	1.1010	0.2483	4.434	9.230e-06	0.6114	1.589441

age60-64	1.5186	0.2316	6.556	5.528e-11	1.0672	1.979110
age65-69	1.7677	0.2294	7.704	1.314e-14	1.3213	2.224503
age70-74	1.8569	0.2353	7.891	3.005e-15	1.3970	2.323556
age75+	1.4197	0.2503	5.672	1.408e-08	0.9254	1.911381

The expected proportion/rate of cases in Fredericia appears to be the highest. Let's compare that city with the others while controlling for age.

```
trtools::contrast(m.b,
  a = list(city = "Fredericia", age = "40-54"),
  b = list(city = c("Horsens", "Kolding", "Vejle"), age = "40-54"),
  cnames = c("vs Horsens", "vs Kolding", "vs Vejle"), tf = exp)
```

	estimate	lower	upper
vs Horsens	1.397	0.9766	1.999
vs Kolding	1.457	1.0059	2.110
vs Vejle	1.318	0.9097	1.909

```
trtools::contrast(m.p,
  a = list(city = "Fredericia", age = "40-54", pop = 1),
  b = list(city = c("Horsens", "Kolding", "Vejle"), age = "40-54", pop = 1),
  cnames = c("vs Horsens", "vs Kolding", "vs Vejle"), tf = exp)
```

	estimate	lower	upper
vs Horsens	1.391	0.9746	1.985
vs Kolding	1.450	1.0035	2.095
vs Vejle	1.313	0.9086	1.897

Note that since there is no interaction in the model, contrasts for city will not depend on the age group. We can also compute the estimated expected proportion (i.e., probability) or expected rate for each model.

```
trtools::contrast(m.b, a = list(city = levels(eba1977$city), age = "40-54"), tf = plogis)
```

	estimate	lower	upper
	0.003589	0.002424	0.005311
	0.002571	0.001701	0.003885
	0.002466	0.001625	0.003741
	0.002726	0.001787	0.004155

```
trtools::contrast(m.p, a = list(city = levels(eba1977$city), age = "40-54", pop = 1), tf = exp)
```

	estimate	lower	upper
	0.003581	0.002419	0.005303
	0.002574	0.001704	0.003890
	0.002470	0.001628	0.003747
	0.002727	0.001789	0.004158

```
d <- expand.grid(city = levels(eba1977$city), age = levels(eba1977$age))
cbind(d, trtools::glmint(m.b, newdata = d))
```

	city	age	fit	low	upp
1	Fredericia	40-54	0.003589	0.002424	0.005311
2	Horsens	40-54	0.002571	0.001701	0.003885
3	Kolding	40-54	0.002466	0.001625	0.003741
4	Vejle	40-54	0.002726	0.001787	0.004155
5	Fredericia	55-59	0.010780	0.007192	0.016129
6	Horsens	55-59	0.007739	0.005135	0.011648
7	Kolding	55-59	0.007424	0.004884	0.011270
8	Vejle	55-59	0.008201	0.005378	0.012487

```

9 Fredericia 60-64 0.016348 0.011360 0.023473
10 Horsens 60-64 0.011755 0.008104 0.017024
11 Kolding 60-64 0.011278 0.007702 0.016489
12 Vejle 60-64 0.012454 0.008520 0.018170
13 Fredericia 65-69 0.020952 0.014654 0.029876
14 Horsens 65-69 0.015086 0.010513 0.021604
15 Kolding 65-69 0.014476 0.009925 0.021069
16 Vejle 65-69 0.015979 0.010956 0.023252
17 Fredericia 70-74 0.022898 0.015845 0.032986
18 Horsens 70-74 0.016496 0.011299 0.024025
19 Kolding 70-74 0.015830 0.010679 0.023407
20 Vejle 70-74 0.017471 0.011844 0.025703
21 Fredericia 75+ 0.014812 0.009872 0.022169
22 Horsens 75+ 0.010646 0.007042 0.016065
23 Kolding 75+ 0.010214 0.006661 0.015633
24 Vejle 75+ 0.011280 0.007368 0.017232

```

```

d <- expand.grid(city = levels(eba1977$city), age = levels(eba1977$age), pop = 1)
cbind(d, trtools::glmint(m.p, newdata = d))

```

```

      city age pop      fit      low      upp
1 Fredericia 40-54 1 0.003581 0.002419 0.005303
2 Horsens 40-54 1 0.002574 0.001704 0.003890
3 Kolding 40-54 1 0.002470 0.001628 0.003747
4 Vejle 40-54 1 0.002727 0.001789 0.004158
5 Fredericia 55-59 1 0.010769 0.007174 0.016167
6 Horsens 55-59 1 0.007742 0.005133 0.011676
7 Kolding 55-59 1 0.007427 0.004883 0.011297
8 Vejle 55-59 1 0.008202 0.005375 0.012517
9 Fredericia 60-64 1 0.016351 0.011335 0.023587
10 Horsens 60-64 1 0.011755 0.008092 0.017075
11 Kolding 60-64 1 0.011277 0.007690 0.016536
12 Vejle 60-64 1 0.012453 0.008506 0.018231
13 Fredericia 65-69 1 0.020976 0.014623 0.030090
14 Horsens 65-69 1 0.015080 0.010488 0.021681
15 Kolding 65-69 1 0.014467 0.009899 0.021141
16 Vejle 65-69 1 0.015976 0.010929 0.023354
17 Fredericia 70-74 1 0.022932 0.015810 0.033263
18 Horsens 70-74 1 0.016486 0.011266 0.024123
19 Kolding 70-74 1 0.015816 0.010646 0.023497
20 Vejle 70-74 1 0.017466 0.011810 0.025830
21 Fredericia 75+ 1 0.014811 0.009848 0.022273
22 Horsens 75+ 1 0.010647 0.007034 0.016116
23 Kolding 75+ 1 0.010214 0.006654 0.015681
24 Vejle 75+ 1 0.011280 0.007358 0.017292

```

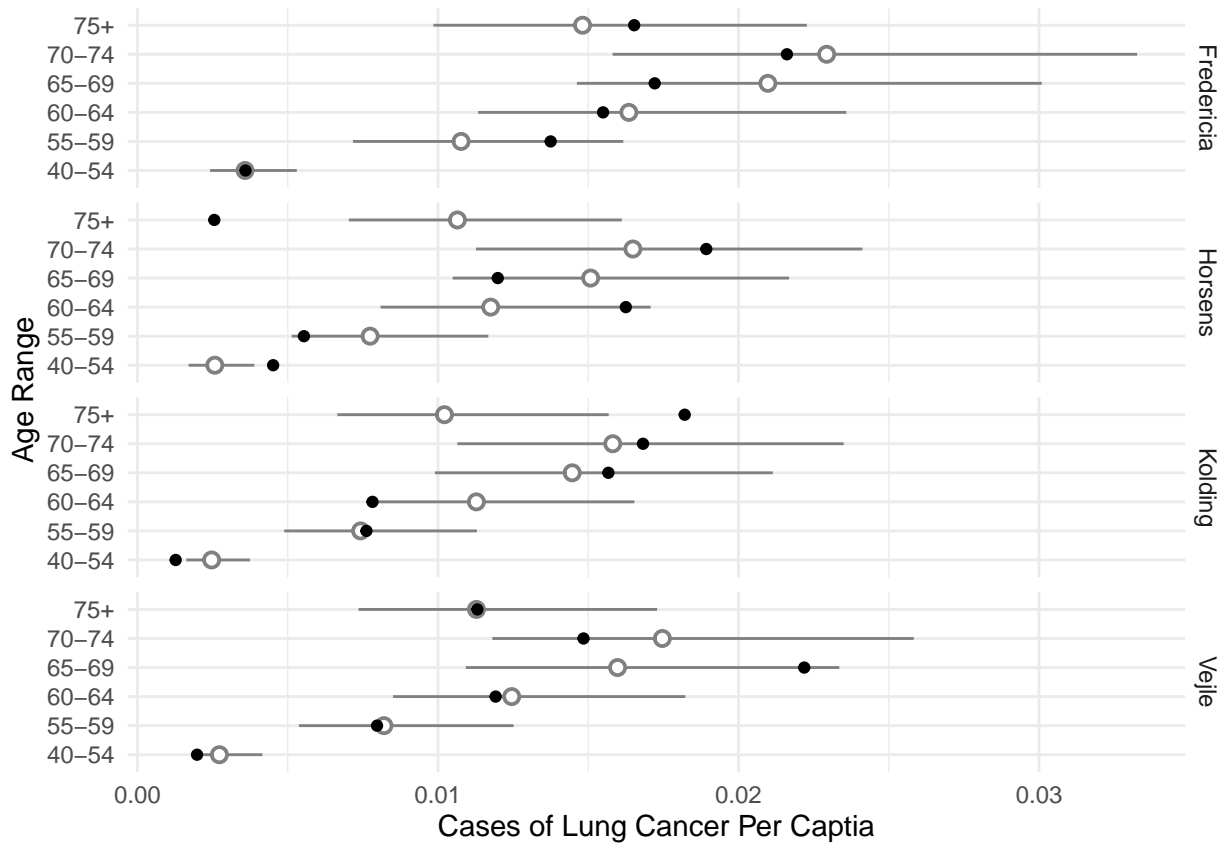
We can use this to make some helpful plots of the estimated rates (or probabilities) of lung cancer.

```

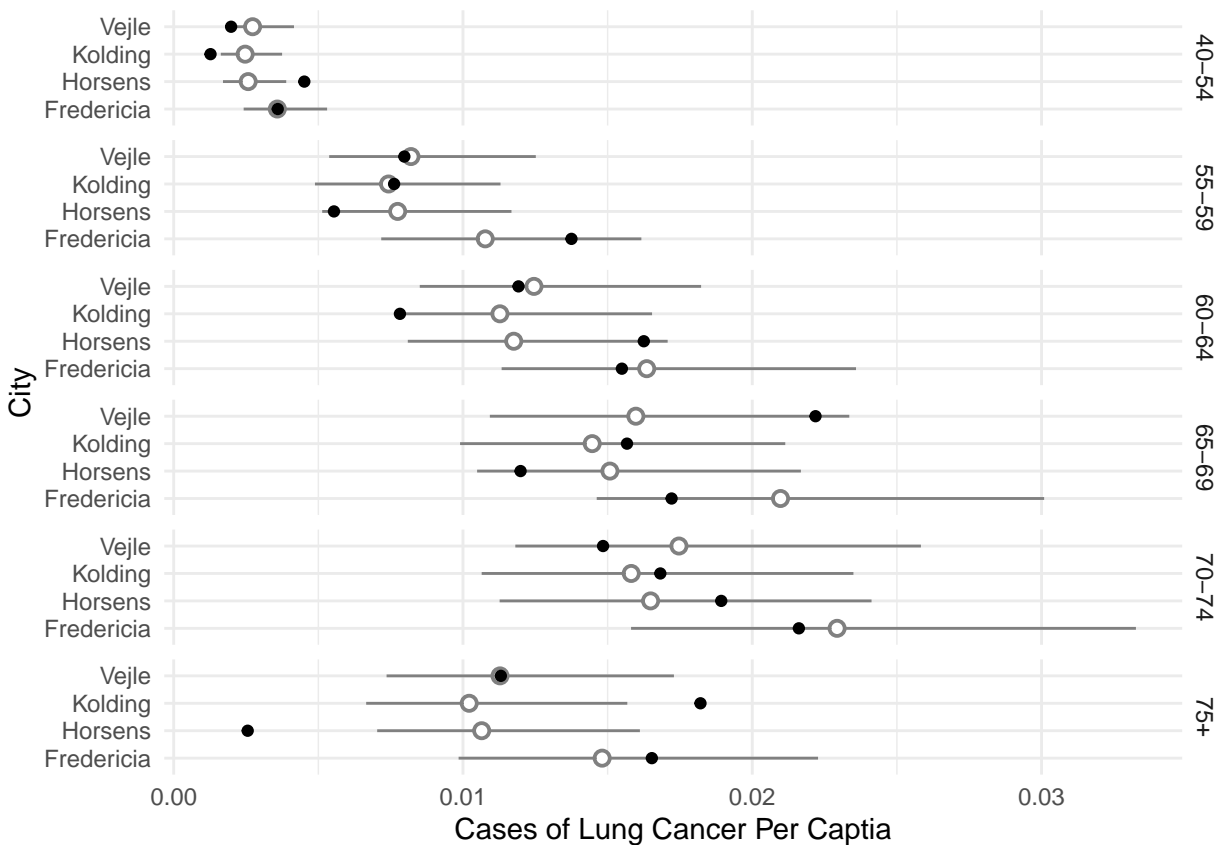
d <- expand.grid(age = levels(eba1977$age), city = levels(eba1977$city), pop = 1)
d <- cbind(d, trtools::glmint(m.p, newdata = d))
p <- ggplot(eba1977, aes(x = age, y = cases/pop)) +
  geom_pointrange(aes(y = fit, ymin = low, ymax = upp),
    shape = 21, fill = "white", data = d, color = grey(0.5)) +
  geom_point() + facet_grid(city ~ .) + coord_flip() +
  labs(x = "Age Range", y = "Cases of Lung Cancer Per Captia") +
  theme_minimal()

```

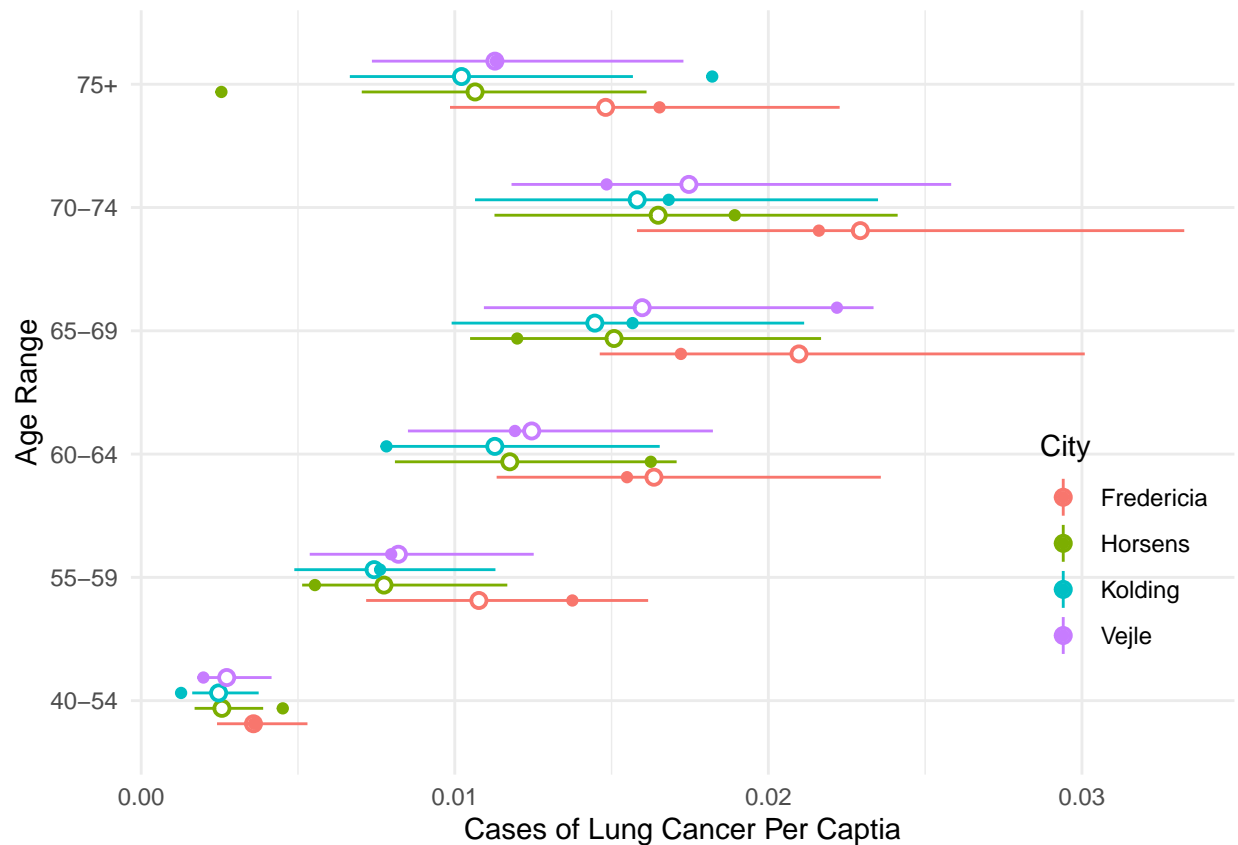
```
plot(p)
```



```
p <- ggplot(eba1977, aes(x = city, y = cases/pop)) +
  geom_pointrange(aes(y = fit, ymin = low, ymax = upp),
    shape = 21, fill = "white", data = d, color = grey(0.5)) +
  geom_point() + facet_grid(age ~ .) + coord_flip() +
  labs(x = "City", y = "Cases of Lung Cancer Per Capita") +
  theme_minimal()
plot(p)
```



```
p <- ggplot(eba1977, aes(x = age, y = cases/pop, color = city)) +
  geom_pointrange(aes(y = fit, ymin = low, ymax = upp),
    shape = 21, fill = "white", data = d,
    position = position_dodge(width = 0.5)) +
  geom_point(position = position_dodge(width = 0.5)) +
  coord_flip() +
  labs(x = "Age Range", y = "Cases of Lung Cancer Per Captia",
    color = "City") +
  theme_minimal() + theme(legend.position = c(0.9,0.3))
plot(p)
```



Separation and Infinite Parameter Estimates

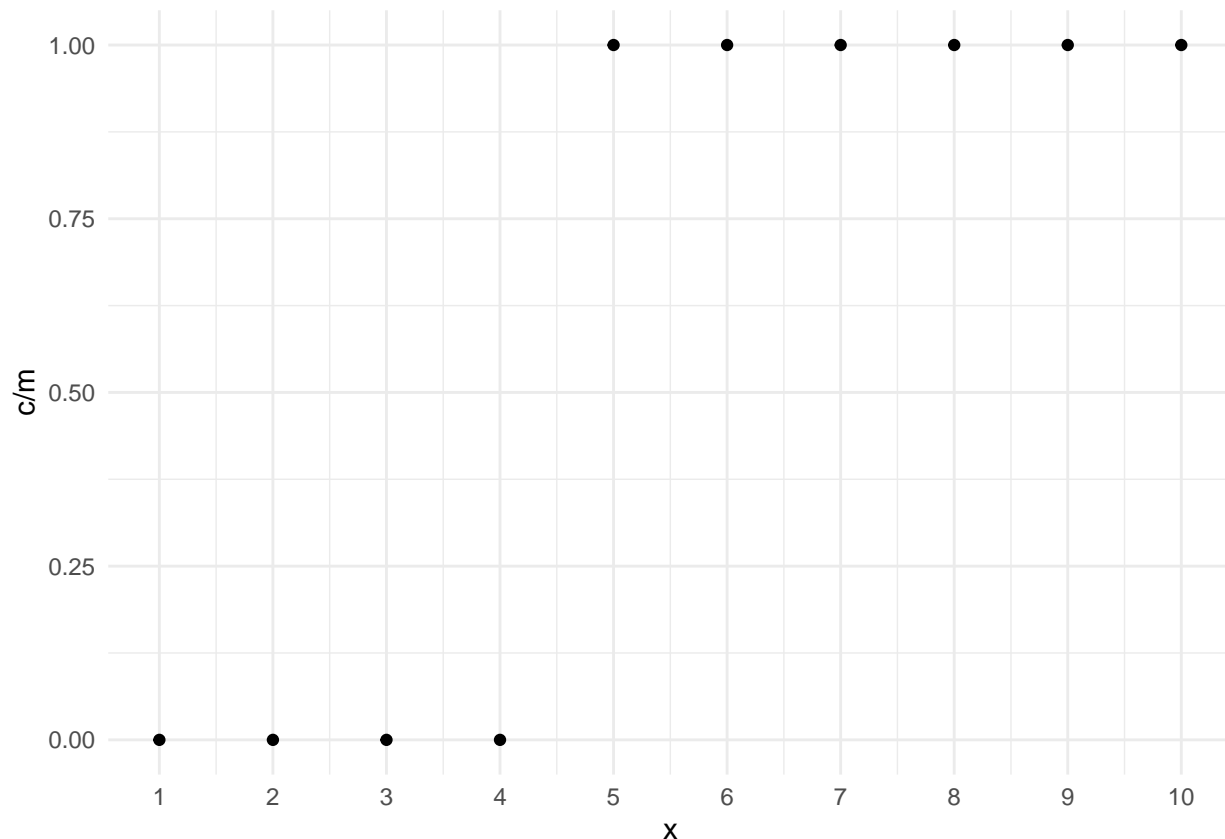
Some GLMs are prone to numerical problems due to (nearly) infinite parameter estimates.

Example: Consider the following data.

```
mydata <- data.frame(m = rep(20, 10), c = rep(c(0,20), c(4,6)), x = 1:10)
mydata
```

	m	c	x
1	20	0	1
2	20	0	2
3	20	0	3
4	20	0	4
5	20	20	5
6	20	20	6
7	20	20	7
8	20	20	8
9	20	20	9
10	20	20	10

```
p <- ggplot(mydata, aes(x = x, y = c/m)) + theme_minimal() +
  geom_point() + scale_x_continuous(breaks = 1:10)
plot(p)
```



If we try to estimate a logistic regression model we get errors and some extreme estimates, standard errors, and confidence intervals.

```
m <- glm(cbind(c,m-c) ~ x, family = binomial, data = mydata)
```

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-212.11	114489	-0.001853	0.9985
x	47.12	25082	0.001879	0.9985

```
confint(m)
```

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: algorithm did not converge

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred


```

                2.5 % 97.5 %
(Intercept) -29559 -28057
x              7969   1966

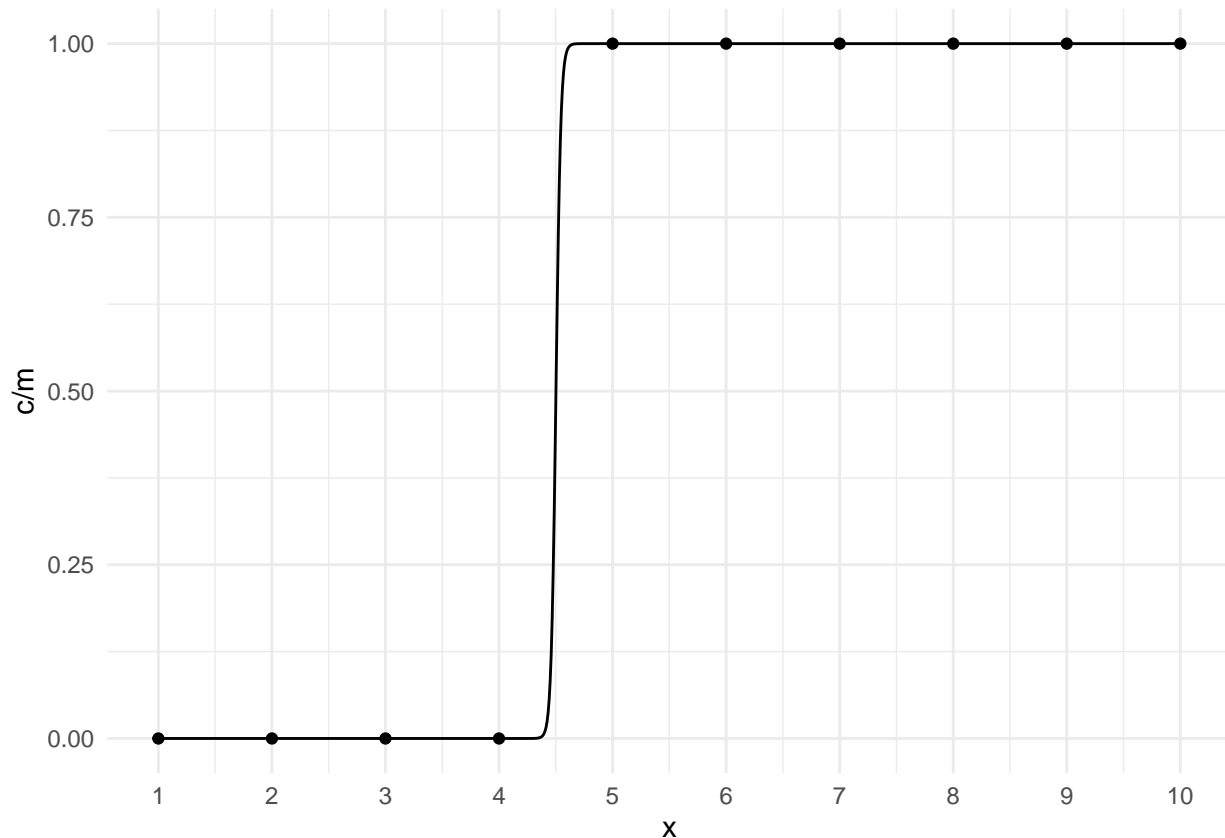
```

But we can still plot the model.

```

d <- data.frame(x = seq(1, 10, length = 1000))
d$yhat <- predict(m, newdata = d, type = "response")
p <- p + geom_line(aes(y = yhat), data = d)
plot(p)

```



The problem is that the estimation procedure “wants” the curve to be a step function, but that only occurs as $\beta_1 \rightarrow \infty$, and the value of x where the estimated expected response is 0.5 equals $-\beta_0/\beta_1$, and for the step function that would be 4.5, so the estimation procedure “wants” the estimate of β_0 to be $-\beta_1 5.5 = -\infty$. This is called *separation*. It is fairly obvious with a single explanatory variable, but much less so with multiple explanatory variables. The example above shows *complete separation* because we can separate the values of y based on the values of x . *Quasi-separation* occurs when this is almost true as in the following example.

```

mydata <- data.frame(m = rep(20, 50), x = seq(1, 10, length = 50),
  c = rep(c(0,20,0,20), c(24,1,1,24)))

```

```

m <- glm(cbind(c,m-c) ~ x, family = binomial, data = mydata)

```

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
summary(m)$coefficients
```

```

                Estimate Std. Error z value Pr(>|z|)
(Intercept)  -39.231      5.542   -7.079 1.448e-12
x              7.133      1.006    7.087 1.371e-12

```

```
confint(m)
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

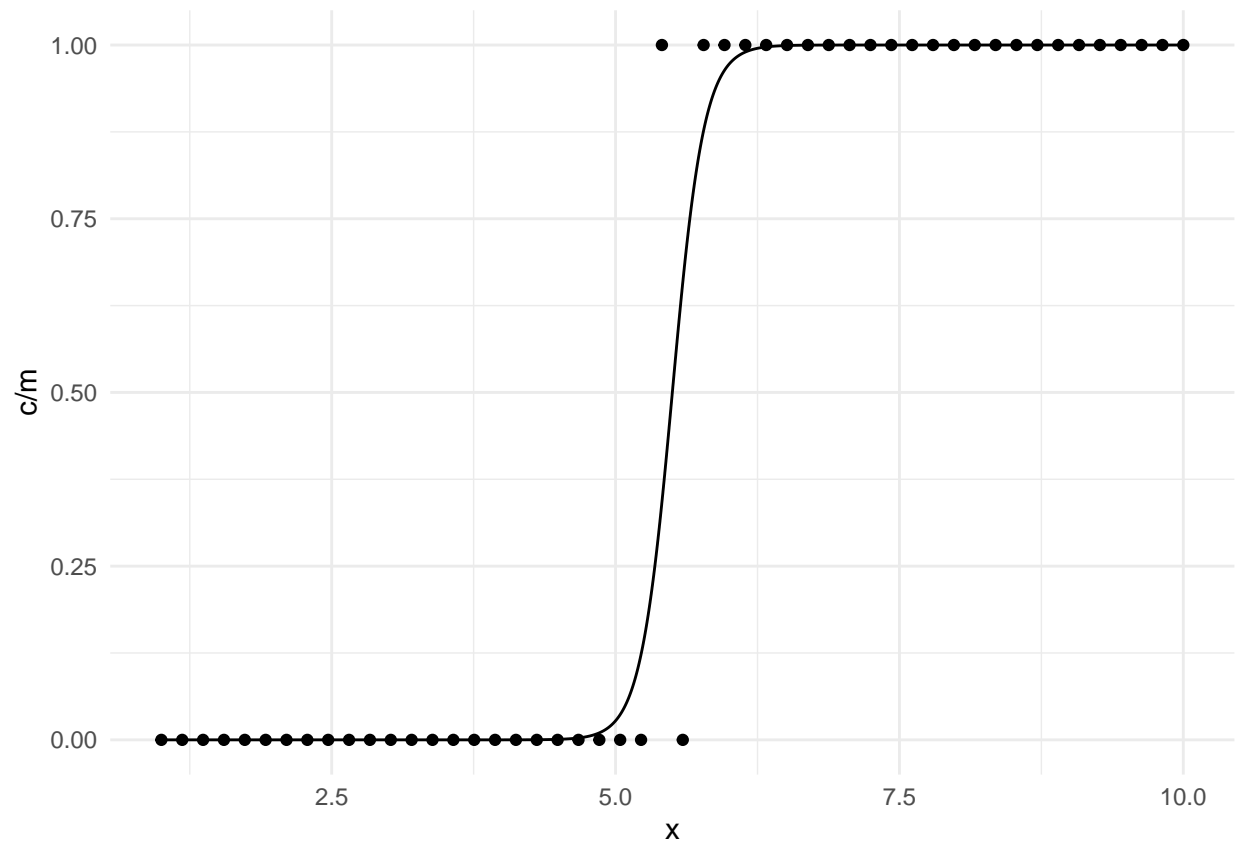
```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
          2.5 %   97.5 %  
(Intercept) -51.696 -29.767  
x             5.414   9.397
```

```
d <- data.frame(x = seq(1, 10, length = 10000))  
d$yhat <- predict(m, newdata = d, type = "response")
```

```
p <- ggplot(mydata, aes(x = x, y = c/m)) + theme_minimal()  
p <- p + geom_point() + geom_line(aes(y = yhat), data = d)  
plot(p)
```



Example: Consider the following data.

```
mydata <- data.frame(m = c(100,100), c = c(25,100), group = c("control","treatment"))
mydata
```

```
   m   c  group
1 100 25 control
2 100 100 treatment
```

```
m <- glm(cbind(c,m-c) ~ group, family = binomial, data = mydata)
summary(m)$coefficients
```

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.099	2.309e-01	-4.7571308	1.964e-06
grouptreatment	28.410	5.169e+04	0.0005496	9.996e-01

```
confint(m)
```

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred

```
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```

                2.5 %      97.5 %
(Intercept)    -1.571    -0.6611
grouptreatment -1849.427 18872.0265

```

A similar problem can happen in Poisson regression where the observed count or rate in a category is zero.

Example: Consider the following data and model.

```
mydata <- data.frame(y = c(20, 10, 50, 15, 0), x = letters[1:5])
mydata
```

```

  y x
1 20 a
2 10 b
3 50 c
4 15 d
5  0 e

```

```
m <- glm(y ~ x, family = poisson, data = mydata)
summary(m)$coefficients
```

```

              Estimate Std. Error   z value Pr(>|z|)
(Intercept)   2.9957    2.236e-01 13.3973220 6.268e-41
xb            -0.6931    3.873e-01 -1.7896983 7.350e-02
xc             0.9163    2.646e-01  3.4632534 5.337e-04
xd            -0.2877    3.416e-01 -0.8422469 3.996e-01
xe            -25.2983   4.225e+04 -0.0005988 9.995e-01

```

```
confint(m)
```

```
Warning: glm.fit: fitted rates numerically 0 occurred
```


2. In some cases with a categorical explanatory variable, we can omit the level(s) where the observed count is zero (in Poisson regression), or the observed proportion is 0 or 1 (in logistic regression). Clearly this precludes inferences concerning that level or its relationship with other levels.
3. For logistic regression (or similar models) a “penalized” or “bias-reduced” estimation method can be used (see the **logistf** and **brglm** packages).