

Wednesday, Apr 27

## Random Effects Approach

The random effects approach conceptualizes the parameters associated with the levels of the many-leveled factor as *random variables*. Another way to think of this is that the levels of that factor are a *sample* of levels from a real or conceptual population of levels.

Note: We sometimes use the term “mixed effects” model for a model where some parameters are modeled as random and some that are not modeled as random (i.e., fixed). Most (but not all) models with random effects also have some fixed effects, and are thus mixed effects models.

**Example:** Consider again the `baserun` data, but a system of subscripts that distinguishes between the *player* and the *observation within each player* so that  $Y_{ij}$  is the  $j$ -th observation of running time for the  $i$ -th player.

```
library(trtools)
head(baserun)
```

```
  round narrow wide
1  5.40    5.50 5.55
2  5.85    5.70 5.75
3  5.20    5.60 5.50
4  5.55    5.50 5.40
5  5.90    5.85 5.70
6  5.45    5.55 5.60
```

If we were to ignore the effect of player we could write a model for these data as

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2},$$

where  $x_{i1}$  and  $x_{i2}$  are indicator variables for two of the three routes.

In the *fixed effects* approach we include an indicator variable for each player, so the model would become

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \beta_3 x_{ij3} + \beta_4 x_{ij4} + \cdots + \beta_{23} x_{ij23},$$

where  $x_{ij3}, x_{ij4}, \dots, x_{ij23}$  are the 21 indicator variables for the 22 players.

In the *random effects* approach we would view  $\beta_3, \beta_4, \dots, \beta_{23}$  as *random variables*. To distinguish the random from the non-random (fixed) parameters I will change the symbols for the indicator variables and the parameters corresponding to the players and write the model as

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \delta_1 z_{ij1} + \delta_2 z_{ij2} + \cdots + \delta_{22} z_{ij22}.$$

Note also that here we have 22 rather than 21 indicator variables (each player has their own parameter). A more compact way to write this model is

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \underbrace{\delta_1 z_{ij1} + \delta_2 z_{ij2} + \cdots + \delta_{22} z_{ij22}}_{\delta_i} = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \delta_i,$$

so that  $\delta_i$  represents the “random effect” of the  $i$ -th player.

Another way to write this model is

$$Y_{ij} = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \delta_i + \epsilon_{ij},$$

where  $\epsilon_{ij}$  is the usual random error term, which is implicitly assumed to be normally-distributed. Thus on the right-hand side of the above expression we have *two* random variables on the right-hand side:  $\delta_i$  and  $\epsilon_{ij}$ .

To complete the model a distribution is needed to be assumed for each  $\delta_i$ . Typically they are assumed to be normally distributed with zero mean and some variance  $\sigma_\delta^2$  so that we write  $\delta_i \sim N(0, \sigma_\delta^2)$ . Because the  $\delta_i$  have a mean of zero they can be viewed as a “deviation” of the effect of the  $i$ -th player from a (conceptual) average player.

The presence of the random  $\delta_i$  parameters fundamentally changes the likelihood function. Specialized inferential methods are (usually) necessary to arrive at correct inferences when random effects are specified. As with other approaches functions to implement these methods require that the data be in “long form” so we reshape the `baserun` data.

```
library(dplyr)
library(tidyr)
baselong <- trtools::baserun %>% mutate(player = factor(letters[1:n()]))) %>%
  pivot_longer(cols = c(round, narrow, wide), names_to = "route", values_to = "time")
head(baselong)
```

```
# A tibble: 6 x 3
  player route   time
  <fct>  <chr>   <dbl>
1 a      round    5.4
2 a      narrow   5.5
3 a      wide     5.55
4 b      round    5.85
5 b      narrow   5.7
6 b      wide     5.75
```

The `lmer` function from the **lme4** package can estimate a *linear mixed effects regression* model with normally-distributed random effects. The model above can be estimated as follows.

```
library(lme4)
m <- lmer(time ~ route + (1 | player), data = baselong)
summary(m)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: time ~ route + (1 | player)
Data: baselong
```

```
REML criterion at convergence: -51.4
```

```
Scaled residuals:
```

	Min	1Q	Median	3Q	Max
	-3.0968	-0.3473	0.0031	0.5001	1.6424

```
Random effects:
```

Groups	Name	Variance	Std.Dev.
player	(Intercept)	0.06448	0.2539
	Residual	0.00745	0.0863

Number of obs: 66, groups: player, 22

```
Fixed effects:
```

	Estimate	Std. Error	t value
(Intercept)	5.53409	0.05718	96.78
routeround	0.00909	0.02603	0.35
routerwide	-0.07500	0.02603	-2.88

Correlation of Fixed Effects:

```
(Intr) rotrnd
routeround -0.228
routewide -0.228 0.500
```

Profile likelihood confidence intervals for  $\sigma_\delta^2$  (the variance of the  $\delta_i$  parameters),  $\sigma^2$  (the variance of  $\epsilon_{ij}$ ), and  $\beta_0$ ,  $\beta_1$ , and  $\beta_2$  can be obtained using `confint`.

```
confint(m)

          2.5 %   97.5 %
.sig01    0.18688 0.34746
.sigma     0.06937 0.10557
(Intercept) 5.42025 5.64793
routeround -0.04186 0.06004
routewide -0.12595 -0.02405
```

Using `lincon` will produce Wald confidence intervals for  $\beta_0$ ,  $\beta_1$ , and  $\beta_2$ .

```
trtools::lincon(m)

          estimate      se    lower    upper  tvalue  df   pvalue
(Intercept) 5.534091 0.05718 5.42202 5.64616 96.7838 Inf 0.000000
routeround   0.009091 0.02603 -0.04192 0.06010 0.3493 Inf 0.726871
routewide   -0.075000 0.02603 -0.12601 -0.02399 -2.8817 Inf 0.003956
```

Other inferences can be made using `trtools::contrast` and the **emmeans** package, but note that player is never specified when using these functions. These tools provide inferences only for the “fixed effects” of the model. We can estimate the expected running time for each route.

```
library(emmeans)

emmeans(m, ~route)

route emmean      SE  df asymp.LCL asymp.UCL
narrow  5.53 0.0572 Inf      5.42      5.65
round   5.54 0.0572 Inf      5.43      5.66
wide    5.46 0.0572 Inf      5.35      5.57
```

Degrees-of-freedom method: asymptotic  
Confidence level used: 0.95

```
trtools::contrast(m, a = list(route = c("narrow", "round", "wide")),
  cnames = c("narrow", "round", "wide"))
```

```
          estimate      se lower upper tvalue  df pvalue
narrow    5.534 0.05718 5.422 5.646 96.78 Inf      0
round     5.543 0.05718 5.431 5.655 96.94 Inf      0
wide      5.459 0.05718 5.347 5.571 95.47 Inf      0
```

Notice that **emmeans** uses the “Kenward-Roger” method of computing approximate degrees of freedom. The issue of degrees of freedom is a difficult problem in models with random effects. Some statisticians suggest just using Wald methods which specify infinite degrees of freedom as an approximation (which is the default in my functions). This can be done using the `lmer.df = "asymptotic"` option.

```
emmeans(m, ~route, lmer.df = "asymptotic")

route emmean      SE  df asymp.LCL asymp.UCL
narrow  5.53 0.0572 Inf      5.42      5.65
```

round	5.54	0.0572	Inf	5.43	5.66
wide	5.46	0.0572	Inf	5.35	5.57

Degrees-of-freedom method: asymptotic  
Confidence level used: 0.95

We can also compare the routes as before.

```
pairs(emmeans(m, ~ route, lmer.df = "asymptotic"), adjust = "none", infer = TRUE)
```

contrast	estimate	SE	df	asympt.LCL	asympt.UCL	z.ratio	p.value
narrow - round	-0.00909	0.026	Inf	-0.0601	0.0419	-0.349	0.7269
narrow - wide	0.07500	0.026	Inf	0.0240	0.1260	2.882	0.0040
round - wide	0.08409	0.026	Inf	0.0331	0.1351	3.231	0.0012

Degrees-of-freedom method: asymptotic  
Confidence level used: 0.95

```
trtools::contrast(m, a = list(route = c("narrow", "round", "wide")),  
  cnames = c("narrow", "round", "wide"))
```

	estimate	se	lower	upper	tvalue	df	pvalue
narrow	5.534	0.05718	5.422	5.646	96.78	Inf	0
round	5.543	0.05718	5.431	5.655	96.94	Inf	0
wide	5.459	0.05718	5.347	5.571	95.47	Inf	0

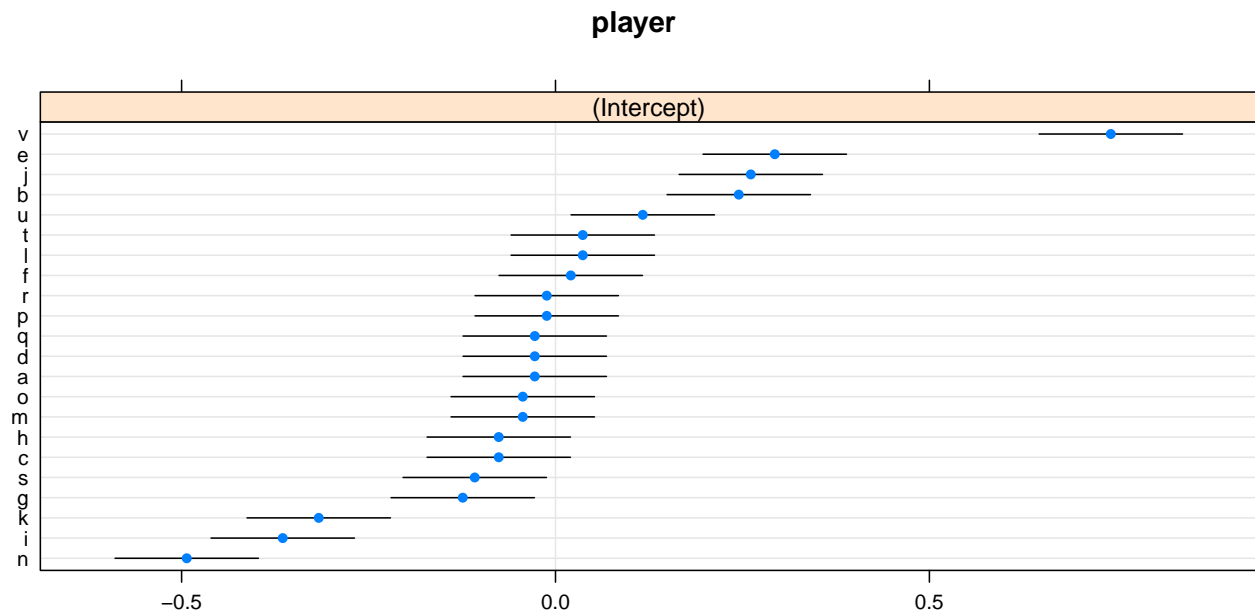
```
trtools::contrast(m,  
  a = list(route = c("narrow", "narrow", "round")),  
  b = list(route = c("round", "wide", "wide")),  
  cnames = c("narrow - round", "narrow - wide", "round - wide"))
```

	estimate	se	lower	upper	tvalue	df	pvalue
narrow - round	-0.009091	0.02603	-0.06010	0.04192	-0.3493	Inf	0.726871
narrow - wide	0.075000	0.02603	0.02399	0.12601	2.8817	Inf	0.003956
round - wide	0.084091	0.02603	0.03308	0.13510	3.2309	Inf	0.001234

Some built-in functions also allow us to plot estimates of the  $\delta_i$  parameters.

```
lattice::dotplot(ranef(m, condVar = TRUE))
```

\$player



Alternatively you can use the `ranef` function to return these estimates and plot them using `ggplot` or something else.

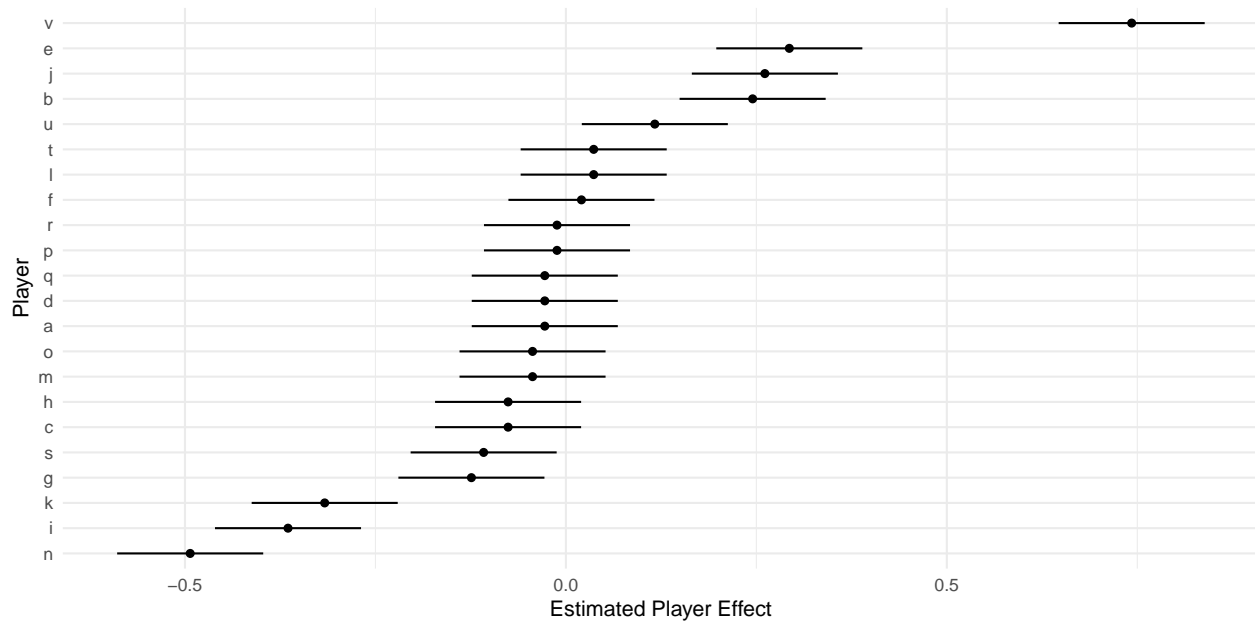
```
d <- as.data.frame(ranef(m))
head(d)
```

	grpvar	term	grp	condval	condsd
1	player	(Intercept)	a	-0.02772	0.0489
2	player	(Intercept)	b	0.24510	0.0489
3	player	(Intercept)	c	-0.07587	0.0489
4	player	(Intercept)	d	-0.02772	0.0489
5	player	(Intercept)	e	0.29325	0.0489
6	player	(Intercept)	f	0.02043	0.0489

```
d <- d %>% mutate(lower = condval - 1.96 * condsd, upper = condval + 1.96 * condsd)
head(d)
```

	grpvar	term	grp	condval	condsd	lower	upper
1	player	(Intercept)	a	-0.02772	0.0489	-0.12357	0.06813
2	player	(Intercept)	b	0.24510	0.0489	0.14925	0.34096
3	player	(Intercept)	c	-0.07587	0.0489	-0.17172	0.01999
4	player	(Intercept)	d	-0.02772	0.0489	-0.12357	0.06813
5	player	(Intercept)	e	0.29325	0.0489	0.19740	0.38910
6	player	(Intercept)	f	0.02043	0.0489	-0.07543	0.11628

```
p <- ggplot(d, aes(x = grp, y = condval)) +
  geom_linerange(aes(ymin = lower, ymax = upper)) +
  geom_point(size = 1.5) +
  theme_minimal() + coord_flip() +
  labs(x = "Player", y = "Estimated Player Effect")
plot(p)
```




**Example:** Now consider again the Sitka data.

```
library(MASS)
head(Sitka, 10)
```

	size	Time	tree	treat	yhat.sub	yhat.avg
1	4.51	152	1	ozone	88.60	67.18
2	4.98	174	1	ozone	164.33	105.05
3	5.41	201	1	ozone	257.27	151.52
4	5.90	227	1	ozone	346.77	196.27
5	6.15	258	1	ozone	453.48	249.63
6	4.24	152	2	ozone	56.53	67.18
7	4.20	174	2	ozone	76.92	105.05
8	4.68	201	2	ozone	101.95	151.52
9	4.92	227	2	ozone	126.05	196.27
10	4.96	258	2	ozone	154.79	249.63

```
p <- ggplot(Sitka, aes(x = Time, y = exp(size))) +
  geom_line(aes(group = tree), alpha = 0.75, size = 0.1) +
  facet_wrap(~ treat) + geom_point(size = 0.5) +
  labs(y = "Size (height times squared diameter)",
       x = "Days Since January 1, 1988") + theme_minimal()
plot(p)
```



First let's consider the model

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \beta_3 x_{ij3} + \delta_i,$$

where  $Y_{ij}$  is the  $j$ -th observation of size for the  $i$ -th tree,  $x_{ij1}$  is an indicator for treatment (ozone),  $x_{ij2}$  is time, and  $x_{ij3} = x_{ij1}x_{ij2}$ .

```
m <- lmer(exp(size) ~ treat * Time + (1 | tree), data = Sitka)
summary(m)
```

```
Linear mixed model fit by REML ['lmerMod']
Formula: exp(size) ~ treat * Time + (1 | tree)
Data: Sitka
```

REML criterion at convergence: 4472

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.811	-0.436	-0.027	0.350	3.620

Random effects:

Groups	Name	Variance	Std.Dev.
tree	(Intercept)	8827	94.0
Residual		2857	53.5

Number of obs: 395, groups: tree, 79

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	-305.123	32.256	-9.46
treatozone	110.675	39.014	2.84
Time	2.509	0.127	19.70
treatozone:Time	-0.788	0.154	-5.12

Correlation of Fixed Effects:

(Intr) tretzn Time

```

treatozone -0.827
Time        -0.799  0.661
treatozone:Time 0.661 -0.799 -0.827


```

```

Sitka$yhat.sub <- predict(m) # for each tree (with deltas)
Sitka$yhat.avg <- predict(m, re.form = NA) # for the "average" tree (deltas = 0)

p <- ggplot(Sitka, aes(x = Time, y = exp(size))) +
  labs(y = "Size (height times squared diameter)",
       x = "Days Since January 1, 1988") +
  theme_minimal() + facet_wrap(~treat) +
  geom_line(aes(y = yhat.sub, group = tree), color = grey(0.75)) +
  geom_line(aes(y = yhat.avg), size = 0.75) +
  geom_point(size = 0.5)
plot(p)

```



This doesn't really capture differences in the growth rates between trees (i.e., an *interaction* between tree and time). Such a model could be written as

$$E(Y_{ij}) = \beta_0 + \beta_1 x_{ij1} + \beta_2 x_{ij2} + \beta_3 x_{ij3} + \delta_i + \gamma_i x_{ij2},$$

where now there are *two* random parameters for each tree:  $\delta_i$  and  $\gamma_i$ . We can also write this model as

$$E(Y_{ij}) = \begin{cases} \beta_0 + \delta_i + (\beta_2 + \gamma_i)t_{ij}, & \text{if the treatment is control,} \\ \beta_0 + \beta_1 + \delta_i + (\beta_2 + \beta_3 + \gamma_i)t_{ij}, & \text{if the treatment is ozone,} \end{cases}$$

where  $t_{ij}$  is time. This means that the linear relationship between time and expected size varies over treatment conditions, but also trees — i.e., each tree has its own intercept and slope (rate).

```

m <- lmer(exp(size) ~ treat * Time + (1 + Time | tree), data = Sitka)

```

```

Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model
failed to converge with max|grad| = 2.20761 (tol = 0.002, component 1)

```

```

Warning in checkConv(attr(opt, "derivs"), opt$par, ctrl = control$checkConv, : Model is nearly unidentifiable:
- Rescale variables?

```



Oh no! Models with random effects are cranky. But let's take the advice of the warning and re-scale time from days to weeks.

```
m <- lmer(exp(size) ~ treat * I(Time/7) + (1 + I(Time/7) | tree), data = Sitka)
```

Warning in checkConv(attr(opt, "derivs"), opt\$par, ctrl = control\$checkConv, : Model failed to converge with max|grad| = 0.0418627 (tol = 0.002, component 1)

That *probably* is not a problem. I suspect it is due to the very high correlation between the random intercept and slope parameters. But changing the optimizer seems to avoid the error.

```
library(optimx)
m <- lmer(exp(size) ~ treat * I(Time/7) + (1 + I(Time/7) | tree), data = Sitka,
  control = lmerControl(optimizer = "optimx", optCtrl = list(method = "nloptn"))
summary(m)
```

Linear mixed model fit by REML ['lmerMod']

Formula: exp(size) ~ treat \* I(Time/7) + (1 + I(Time/7) | tree)

Data: Sitka

Control: lmerControl(optimizer = "optimx", optCtrl = list(method = "nloptn"))

REML criterion at convergence: 3915

Scaled residuals:

Min	1Q	Median	3Q	Max
-2.964	-0.395	-0.049	0.391	4.816

Random effects:

Groups	Name	Variance	Std.Dev.	Corr
tree	(Intercept)	22745.6	150.82	
	I(Time/7)	70.2	8.38	-0.99
Residual		383.2	19.57	

Number of obs: 395, groups: tree, 79

Fixed effects:

	Estimate	Std. Error	t value
(Intercept)	-305.12	31.65	-9.64
treatozone	110.68	38.29	2.89
I(Time/7)	17.56	1.71	10.29
treatozone:I(Time/7)	-5.52	2.06	-2.67

Correlation of Fixed Effects:


	(Intr)	tretzn	I(T/7)
treatozone	-0.827		
I(Time/7)	-0.980	0.810	
trtz:I(T/7)	0.810	-0.980	-0.827

I found that you get more or less the same result even without changing the optimizer. Here's a plot.

```
Sitka$yhat.sub <- predict(m) # for each tree (with deltas)
Sitka$yhat.avg <- predict(m, re.form = NA) # for the "average" tree (deltas = 0)

p <- ggplot(Sitka, aes(x = Time, y = exp(size))) +
  labs(y = "Size (height times squared diameter)",
    x = "Days Since January 1, 1988") +
  theme_minimal() + facet_wrap(~treat) +
  geom_line(aes(y = yhat.sub, group = tree), color = grey(0.75)) +
```

```
geom_line(aes(y = yhat.avg), size = 0.75) +
geom_point(size = 0.5)
plot(p)
```



Now we can estimate and compare the (average) growth rates in the control and ozone conditions (per 100 days with `contrast` and per day with `emmeans`).

```
trtools::contrast(m,
  a = list(Time = 2, treat = c("control", "ozone")),
  b = list(Time = 1, treat = c("control", "ozone")),
  cnames = c("control", "ozone"))
```

	estimate	se	lower	upper	tvalue	df	pvalue
control	2.509	0.2438	2.031	2.987	10.29	Inf	7.718e-25
ozone	1.721	0.1659	1.396	2.046	10.37	Inf	3.227e-25

```
trtools::contrast(m,
  a = list(Time = 2, treat = "control"),
  b = list(Time = 1, treat = "control"),
  u = list(Time = 2, treat = "ozone"),
  v = list(Time = 1, treat = "ozone"))
```

	estimate	se	lower	upper	tvalue	df	pvalue
	0.7881	0.2949	0.21	1.366	2.672	Inf	0.007537

```
emmeans(m, ~ treat, var = "Time", lmer.df = "asymptotic")
```

treat	Time.trend	SE	df	asympt.LCL	asympt.UCL
control	2.51	0.244	Inf	2.03	2.99
ozone	1.72	0.166	Inf	1.40	2.05

Degrees-of-freedom method: asymptotic  
Confidence level used: 0.95

```
pairs(emmeans(m, ~ treat, var = "Time",
  lmer.df = "asymptotic"), infer = TRUE)
```

contrast	estimate	SE	df	asympt.LCL	asympt.UCL	z.ratio	p.value
control - ozone	0.788	0.295	Inf	0.21	1.37	2.672	0.0075

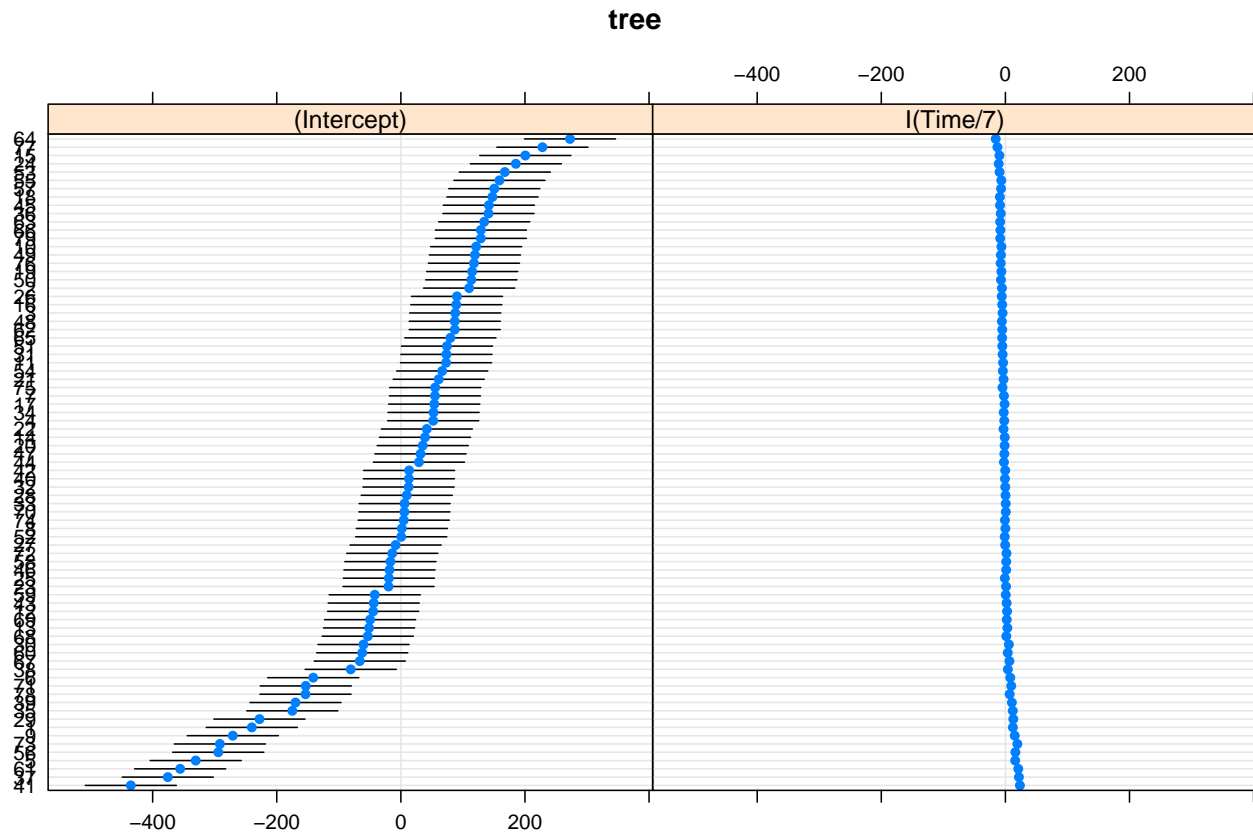
Degrees-of-freedom method: asymptotic

Confidence level used: 0.95

We can plot estimates of the  $\delta_i$  and  $\gamma_i$  parameters for each tree.

```
lattice::dotplot(ranef(m, condVar = TRUE))
```

```
$tree
```



```
d <- as.data.frame(ranef(m))
head(d)
```

	grpvar	term	grp	condval	condsd
1	tree	(Intercept)	1	-240.18	37.45
2	tree	(Intercept)	2	110.08	37.45
3	tree	(Intercept)	3	87.68	37.45
4	tree	(Intercept)	4	52.13	37.45
5	tree	(Intercept)	5	-330.64	37.45
6	tree	(Intercept)	6	-141.21	37.45

```
d <- d %>% mutate(lower = condval - 1.96 * condsd, upper = condval + 1.96 * condsd)
head(d)
```

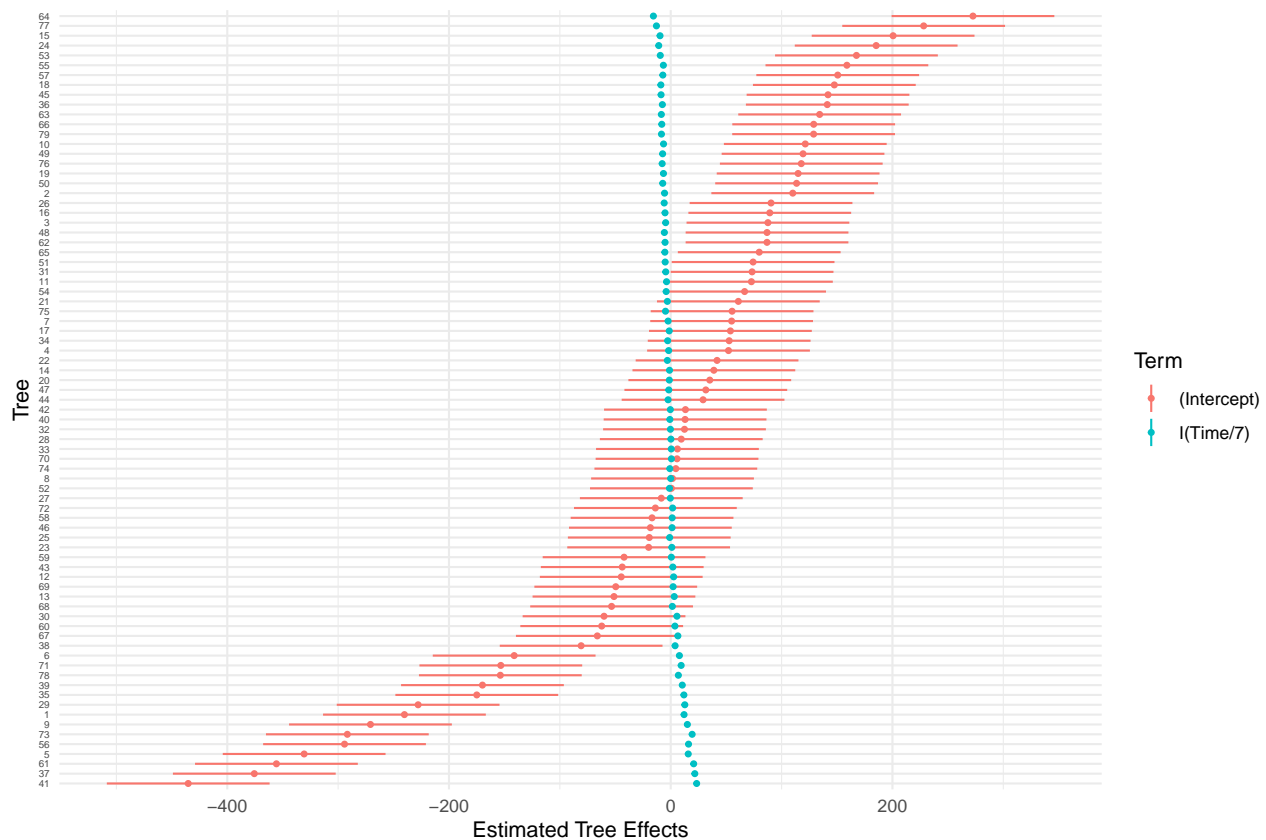
	grpvar	term	grp	condval	condsd	lower	upper
1	tree	(Intercept)	1	-240.18	37.45	-313.59	-166.8
2	tree	(Intercept)	2	110.08	37.45	36.67	183.5

```

3   tree (Intercept)    3   87.68  37.45   14.27  161.1
4   tree (Intercept)    4   52.13  37.45  -21.28  125.5
5   tree (Intercept)    5 -330.64  37.45 -404.05 -257.2
6   tree (Intercept)    6 -141.21  37.45 -214.62  -67.8

p <- ggplot(d, aes(x = grp, y = condval, color = term)) +
  geom_linerange(aes(ymin = lower, ymax = upper)) +
  geom_point(size = 1) +
  theme_minimal() + coord_flip() +
  labs(x = "Tree", y = "Estimated Tree Effects", color = "Term") +
  theme(axis.text.y = element_text(size = 5))
plot(p)

```



**Example:** Consider again the smoking cessation meta analysis data.

```

library(dplyr)
library(tidyr)
quitsmoke <- HSAUR3::smoking
quitsmoke$study <- rownames(quitsmoke)
quitsmoke.quits <- quitsmoke %>% dplyr::select(study, qt, qc) %>%
  rename(gum = qt, control = qc) %>%
  gather(gum, control, key = treatment, value = quit)
quitsmoke.total <- quitsmoke %>% dplyr::select(study, tt, tc) %>%
  rename(gum = tt, control = tc) %>%
  gather(gum, control, key = treatment, value = total)
quitsmoke <- full_join(quitsmoke.quits, quitsmoke.total) %>%
  mutate(study = factor(study)) %>% arrange(study)

```

```
head(quitsmoke)
```

```

      study treatment quit total
1   Blondal89      gum   37    92
2   Blondal89 control   24    90
3  Campbell91      gum   21   107
4  Campbell91 control   21   105
5 Fagerstrom82      gum   30    50
6 Fagerstrom82 control   23    50

```

We can introduce a random “study effect” into a logistic regression model to create a *generalized linear mixed effects regression* model. This would be written as

$$\log \left[ \frac{E(Y_{ij})}{1 - E(Y_{ij})} \right] = \beta_0 + \beta_1 x_{ij} + \delta_i,$$

where  $Y_{ij}$  is the  $j$ -th proportion of people quitting in the  $i$ -th study, and  $x_{ij}$  is an indicator variable for treatment (gum). This model can be estimated as follows.

```

m <- glmer(cbind(quit, total - quit) ~ treatment + (1 | study),
  family = binomial, data = quitsmoke)
summary(m)

```

```
Generalized linear mixed model fit by maximum likelihood (Laplace Approximation) [glmerMod]
```

```

Family: binomial ( logit )
Formula: cbind(quit, total - quit) ~ treatment + (1 | study)
Data: quitsmoke

```

```

      AIC      BIC   logLik deviance df.resid
367.3    373.2   -180.6   361.3      49

```

Scaled residuals:

```

      Min      1Q  Median      3Q      Max
-1.9940 -0.6602 -0.0373  0.4633  2.3042

```

Random effects:

```

Groups Name      Variance Std.Dev.
study (Intercept) 0.412    0.642
Number of obs: 52, groups: study, 26

```

Fixed effects:

```

      Estimate Std. Error z value Pr(>|z|)
(Intercept)  -1.3625    0.1376   -9.90 < 2e-16 ***
treatmentgum  0.5149    0.0655    7.87 3.6e-15 ***
---

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Correlation of Fixed Effects:

```

      (Intr)
treatmentgm -0.281

```

We can estimate the odds ratio for the treatment, which is assumed to be the same for every study in this model.

```
trtools::contrast(m, tf = exp,
  a = list(treatment = "gum"),
  b = list(treatment = "control"))
```

```
estimate lower upper
1.673 1.472 1.902
```

```
pairs(emmeans(m, ~ treatment, type = "response"), reverse = TRUE)
```

```
contrast      odds.ratio    SE  df null z.ratio p.value
gum / control      1.67 0.11 Inf    1   7.867  <.0001
```

Tests are performed on the log odds ratio scale

We can extend the model so that the treatment effect varies over studies (i.e., an interaction between treatment and study).

```
m <- glmer(cbind(quit, total - quit) ~ treatment + (1 + treatment | study),
  family = binomial, data = quitsmoke)
summary(m)
```

Generalized linear mixed model fit by maximum likelihood (Laplace Approximation) [glmerMod]

Family: binomial ( logit )

Formula: cbind(quit, total - quit) ~ treatment + (1 + treatment | study)

Data: quitsmoke

```
      AIC      BIC   logLik deviance df.resid
368.3    378.0  -179.1    358.3        47
```

Scaled residuals:

```
      Min      1Q  Median      3Q      Max
-1.4423 -0.4678  0.0217  0.3796  1.6638
```

Random effects:

```
Groups Name      Variance Std.Dev. Corr
study  (Intercept) 0.4211   0.649
      treatmentgum 0.0508   0.225  -0.12
```

Number of obs: 52, groups: study, 26

Fixed effects:

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  -1.3991     0.1415  -9.89 < 2e-16 ***
treatmentgum   0.5723     0.0887   6.45 1.1e-10 ***
---
```

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

```
(Intr)
treatmentgm -0.340
```

Now our odds ratios are for a “typical” study.

```
trtools::contrast(m, tf = exp,
  a = list(treatment = "gum"),
  b = list(treatment = "control"))
```

```
estimate lower upper
1.772 1.489 2.109
```

```
pairs(emmeans(m, ~ treatment, type = "response"), reverse = TRUE)
```

```
contrast      odds.ratio      SE df null z.ratio p.value
gum / control      1.77 0.157 Inf   1   6.449 <.0001
```

Tests are performed on the log odds ratio scale

Note: In logistic regression, if your response variable is *binary* (i.e., not aggregated counts) use the option `nAGQ = x` where `x` is maybe 21+.

```
m <- glmer(cbind(quit, total - quit) ~ treatment + (1 | study),
  family = binomial, data = quitsmoke, nAGQ = 31)
summary(m)
```

```
Generalized linear mixed model fit by maximum likelihood (Adaptive Gauss-Hermite
  Quadrature, nAGQ = 31) [glmerMod]
Family: binomial ( logit )
Formula: cbind(quit, total - quit) ~ treatment + (1 | study)
Data: quitsmoke
```

```
      AIC      BIC   logLik deviance df.resid
136.2    142.0    -65.1    130.2      49
```

Scaled residuals:

```
      Min      1Q  Median      3Q      Max
-1.9943 -0.6602 -0.0374  0.4631  2.3039
```

Random effects:

```
Groups Name      Variance Std.Dev.
study (Intercept) 0.413    0.642
Number of obs: 52, groups: study, 26
```

Fixed effects:

```
              Estimate Std. Error z value Pr(>|z|)
(Intercept)   -1.3625     0.1378   -9.89   <2e-16 ***
treatmentgum    0.5149     0.0655    7.86    4e-15 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Correlation of Fixed Effects:

```
(Intr)
treatmentgm -0.281
```

It can also be used in other GLMERs. Because of the complexity of the likelihood function in these models, there are many different numerical approaches to estimation.


**Example:** Consider a random effects approach for the `leprosy` data.

```
library(ALA)
head(leprosy)
```

```
      id drug period nBacilli
1     1   A    pre         11
31    1   A   post          6
2     2   B    pre          6
```

```
32 2 B post 0
3 3 C pre 16
33 3 C post 13
```

```
p <- ggplot(leprosy, aes(x = drug, y = nBacilli, fill = period)) +
  geom_dotplot(binaxis = "y", method = "histodot",
    stackdir = "center", binwidth = 1,
    position = position_dodge(width = 0.5)) +
  scale_fill_manual(values = c("white", "black")) +
  labs(x = "Drug", y = "Number of Bacilli", fill = "Period") +
  theme_minimal()
plot(p)
```



```
m <- glmer(nBacilli ~ drug * period + (1 | id),
  family = poisson, data = leprosy)
summary(m)
```

Generalized linear mixed model fit by maximum likelihood (Laplace Approximation) [glmerMod]

Family: poisson ( log )  
 Formula: nBacilli ~ drug \* period + (1 | id)  
 Data: leprosy

AIC	BIC	logLik	deviance	df.resid
363.9	378.6	-175.0	349.9	53

Scaled residuals:

Min	1Q	Median	3Q	Max
-1.8757	-0.5729	0.0637	0.4264	1.9372

Random effects:

Groups Name	Variance	Std.Dev.
id (Intercept)	0.259	0.509

Number of obs: 60, groups: id, 30



Fixed effects:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	2.0936	0.1953	10.72	< 2e-16 ***
drugB	0.0506	0.2737	0.19	0.85320
drugC	0.3836	0.2682	1.43	0.15270
periodpost	-0.5623	0.1704	-3.30	0.00097 ***
drugB:periodpost	0.0680	0.2344	0.29	0.77164
drugC:periodpost	0.5147	0.2114	2.43	0.01490 *

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Correlation of Fixed Effects:

	(Intr)	drugB	drugC	prdpst	drdB:p
drugB	-0.707				
drugC	-0.725	0.515			
periodpost	-0.317	0.226	0.231		
drdB:prdpst	0.230	-0.317	-0.168	-0.727	
drdB:prdpst	0.255	-0.182	-0.321	-0.806	0.586

Estimated ratios for each drug.

```
pairs(emmeans(m, ~ period | drug, type = "response"),
      reverse = TRUE, infer = TRUE)
```

```
drug = A:
contrast    ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
post / pre 0.570 0.0971 Inf    0.408    0.796    1 -3.300 0.0010
```

```
drug = B:
contrast    ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
post / pre 0.610 0.0982 Inf    0.445    0.836    1 -3.071 0.0021
```

```
drug = C:
contrast    ratio      SE df asymp.LCL asymp.UCL null z.ratio p.value
post / pre 0.954 0.1193 Inf    0.746    1.218    1 -0.381 0.7035
```

Confidence level used: 0.95

Intervals are back-transformed from the log scale

Tests are performed on the log scale

```
trtools::contrast(m, tf = exp,
  a = list(period = "post", drug = c("A","B","C")),
  b = list(period = "pre", drug = c("A","B","C")),
  cnames = c("A","B","C"))
```

	estimate	lower	upper
A	0.5699	0.4081	0.7958
B	0.6100	0.4450	0.8362
C	0.9535	0.7461	1.2185

We can also compare the rate ratios.

```
pairs(pairs(emmeans(m, ~ period | drug, type = "response"),
  reverse = TRUE), by = NULL, adjust = "none")
```

contrast	ratio	SE	df	null	z.ratio	p.value
----------	-------	----	----	------	---------	---------

(post / pre A) / (post / pre B)	0.934	0.219	Inf	1	-0.290	0.7716
(post / pre A) / (post / pre C)	0.598	0.126	Inf	1	-2.435	0.0149
(post / pre B) / (post / pre C)	0.640	0.130	Inf	1	-2.191	0.0284

Tests are performed on the log scale

```
trtools::contrast(m, tf = exp,
  a = list(period = "post", drug = "A"),
  b = list(period = "pre", drug = "A"),
  u = list(period = "post", drug = "B"),
  v = list(period = "pre", drug = "B"))
```

estimate	lower	upper
0.9342	0.5901	1.479

```
trtools::contrast(m, tf = exp,
  a = list(period = "post", drug = "A"),
  b = list(period = "pre", drug = "A"),
  u = list(period = "post", drug = "C"),
  v = list(period = "pre", drug = "C"))
```

estimate	lower	upper
0.5977	0.3949	0.9045

```
trtools::contrast(m, tf = exp,
  a = list(period = "post", drug = "B"),
  b = list(period = "pre", drug = "B"),
  u = list(period = "post", drug = "C"),
  v = list(period = "pre", drug = "C"))
```

estimate	lower	upper
0.6397	0.429	0.954